

A JavaScript and jQuery Developer's Guide



Learning

JavaScript Design Patterns

O'REILLY®

Addy Osmani

Learning JavaScript Design Patterns

Addy Osmani

Editor

Mary Treseler

Copyright © 2012 Adnan Osmani

Creative Commons [Attribution-NonCommercial-ShareAlike 3.0](#) unported license. You are free to remix, tweak, and build upon this work non-commercially, as long as you credit Addy Osmani (the copyright holder) and license your new creations under the identical terms. Any of the above conditions can be waived if you get permission from the copyright holder. For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a [link to the license](#).

O'REILLY®

O'Reilly Media

Preface

Design patterns are reusable solutions to commonly occurring problems in software design. They are both exciting and a fascinating topic to explore in any programming language.

One reason for this is that they help us build upon the combined experience of many developers that came before us and ensure we structure our code in an optimized way, meeting the needs of problems we're attempting to solve.

Design patterns also provide us a common vocabulary to describe solutions. This can be significantly simpler than describing syntax and semantics when we're attempting to convey a way of structuring a solution in code form to others.

In this book we will explore applying both classical and modern design patterns to the JavaScript programming language.

Target Audience

This book is targeted at professional developers wishing to improve their knowledge of design patterns and how they can be applied to the JavaScript programming language.

Some of the concepts covered (closures, prototypal inheritance) will assume a level of basic prior knowledge and understanding. If you find yourself needing to read further about these topics, a list of suggested titles is provided for convenience.

If you would like to learn how to write beautiful, structured and organized code, I believe this is the book for you.

Acknowledgments

I will always be grateful for the talented technical reviewers who helped review and improve this book, including those from the community at large. The knowledge and enthusiasm they brought to the project was simply amazing. The official technical reviewers tweets and blogs are also a regular source of both ideas and inspiration and I wholeheartedly recommend checking them out.

- Luke Smith (<http://lucassmith.name>, @ls_n)
- Nicholas Zakas (<http://nczonline.net>, @slicknet)
- Andrée Hansson (<http://andreehansson.se>, @peolanha)
- Alex Sexton (<http://alexsexton.com>, @slexaxton)

I would also like to thank Rebecca Murphey (<http://rebeccamurphey.com>, @rmurphey) for providing the inspiration to write this book and more importantly, continue to make it both available on GitHub and via O'Reilly.

Finally, I would like to thank my wonderful wife Ellie, for all of her support while I was putting together this publication.

Credits

Whilst some of the patterns covered in this book were implemented based on personal experience, many of them have been previously identified by the JavaScript community. This work is as such the production of the combined experience of a number of developers. Similar to Stoyan Stefanov's logical approach to preventing interruption of the narrative with credits (in JavaScript Patterns), I have listed credits and suggested reading for any content covered in the references section.

If any articles or links have been missed in the list of references, please accept my heartfelt apologies. If you contact me I'll be sure to update them to include you on the list.

Reading

Whilst this book is targeted at both beginners and intermediate developers, a basic understanding of JavaScript fundamentals is assumed. Should you wish to learn more about the language, I am happy to recommend the following titles:

- JavaScript: The Definitive Guide by David Flanagan
- Eloquent JavaScript by Marijn Haverbeke
- JavaScript Patterns by Stoyan Stefanov
- Writing Maintainable JavaScript by Nicholas Zakas
- JavaScript: The Good Parts by Douglas Crockford

Chapter 1. Introduction

One of the most important aspects of writing maintainable code is being able to notice the recurring themes in that code and optimize them. This is an area where knowledge of design patterns can prove invaluable.

In the first part of this book, we will explore the history and importance of design patterns which can really be applied to any programming language. If you're already sold on or are familiar with this history, feel free to skip to the chapter ['What is a Pattern?'](#) to continue reading.

Design patterns can be traced back to the early work of a civil engineer named **Christopher Alexander**. He would often write publications about his experience in solving design issues and how they related to buildings and towns. One day, it occurred to Alexander that when used time and time again, certain design constructs lead to a desired optimal effect.

In collaboration with Sara Ishikawa and Murray Silverstein, Alexander produced a pattern language that would help empower anyone wishing to design and build at any scale. This was published back in 1977 in a paper titled 'A Pattern Language', which was later released as a complete hardcover **book**.

Some 30 years ago, software engineers began to incorporate the principles Alexander had written about into the first documentation about design patterns, which was to be a guide for novice developers looking to improve their coding skills. It's important to note that the concepts behind design patterns have actually been around in the programming industry since its inception, albeit in a less formalized form.

One of the first and arguably most iconic formal works published on design patterns in software engineering was a book in 1995 called Design Patterns: Elements Of Reusable Object-Oriented Software. This was written by **Erich Gamma, Richard Helm, Ralph Johnson** and **John Vlissides** - a group that became known as the Gang of Four (or GoF for short).

The GoF's publication is considered quite instrumental to pushing the concept of design patterns further in our field as it describes a number of development

techniques and pitfalls as well as providing twenty-three core Object-Oriented design patterns frequently used around the world today. We will be covering these patterns in more detail in the section ‘Categories of Design Patterns’.

In this book, we will take a look at a number of popular JavaScript design patterns and explore why certain patterns may be more suitable for your projects than others. Remember that patterns can be applied not just to vanilla JavaScript (i.e standard JavaScript code), but also to abstracted libraries such as **jQuery** or **dojo** as well. Before we begin, let’s look at the exact definition of a ‘pattern’ in software design.

Chapter 2. What is a Pattern?

A pattern is a reusable solution that can be applied to commonly occurring problems in software design - in our case - in writing JavaScript-powered applications. Another way of looking at patterns are as templates for how you solve problems - ones which can be used in quite a few different situations.

So, why is it important to understand patterns and be familiar with them?. Design patterns have three main benefits:

1. **Patterns are proven solutions:** They provide solid approaches to solving issues in software development using proven solutions that reflect the experience and insights the developers that helped define and improve them bring to the pattern.
2. **Patterns can be easily reused:** A pattern usually reflects an out of the box solution that can be adapted to suit your own needs. This feature makes them quite robust.
3. **Patterns can be expressive:** When you look at a pattern there's generally a set structure and 'vocabulary' to the solution presented that can help express rather large solutions quite elegantly.

Patterns are **not** an exact solution. It's important that we remember the role of a pattern is merely to provide us with a solution scheme. Patterns don't solve all design problems nor do they replace good software designers, however, they **do** support them. Next we'll take a look at some of the other advantages patterns have to offer.

- **Reusing patterns assists in preventing minor issues that can cause major problems in the application development process.** What this means is when code is built on proven patterns, we can afford to spend less time worrying about the structure of our code and more time focusing on the quality of our overall solution. This is because patterns can encourage us to code in a more structured and organized fashion avoiding the need to refactor it for cleanliness purposes in the future.
- **Patterns can provide generalized solutions which are documented in a**

fashion that doesn't require them to be tied to a specific problem. This generalized approach means that regardless of the application (and in many cases the programming language) you are working with, design patterns can be applied to improve the structure of your code.

- **Certain patterns can actually decrease the overall file-size footprint of your code by avoiding repetition.** By encouraging developers to look more closely at their solutions for areas where instant reductions in repetition can be made, e.g. reducing the number of functions performing similar processes in favor of a single generalized function, the overall size of your codebase can be decreased.
- **Patterns add to a developers vocabulary, which makes communication faster.**
- **Patterns that are frequently used can be improved over time by harnessing the collective experiences other developers using those patterns contribute back to the design pattern community.** In some cases this leads to the creation of entirely new design patterns whilst in others it can lead to the provision of improved guidelines on how specific patterns can be best used. This can ensure that pattern-based solutions continue to become more robust than ad-hoc solutions may be.

We already use patterns everyday

To understand how useful patterns can be, let's review a very simple element selection problem that the jQuery library solves for us everyday.

If we imagine that we have a script where for each DOM element on a page with class "foo" we want to increment a counter, what's the simplest efficient way to query for the list we need?. Well, there are a few different ways this problem could be tackled:

1. Select all of the elements in the page and then store them. Next, filter this list and use regular expressions (or another means) to only store those with the class "foo".
2. Use a modern native browser feature such as `querySelectorAll()` to select all of the elements with the class "foo".

3. Use a native feature such as `getElementsByClassName()` to similarly get back the desired list.

So, which of these is the fastest?. You might be interested to know that it's actually number 3. by a factor of 8-10 times the **alternatives**. In a real-world application however, 3. will not work in versions of Internet Explorer below 9 and thus it's necessary to use 1. where 3. isn't supported.

Developers using jQuery don't have to worry about this problem, as it's luckily abstracted away for us. The library opts for the most optimal approach to selecting elements depending on what your browser supports.

Core internally uses a number of different **design patterns**, the most frequent one being a facade. This provides a simple set of abstracted interfaces (e.g `$el.css()`, `$el.animate()`) to several more complex underlying bodies of code.

We're probably all also familiar with jQuery's `$('selector')`. This is significantly more easy to use for selecting HTML elements on a page versus having to manually handle opt for `getElementById()`, `getElementsByClassName()`, `getElementByTagName` and so on. Although we know that `querySelectorAll()` attempts to solve this problem, compare the effort involved in using jQuery's facade interfaces vs. selecting the most optimal selection paths ourselves. There's no contest! abstractions using patterns can offer real-world value.

We'll be looking at this and more design patterns later on in the book.

Chapter 3. 'Pattern'-ity Testing, Proto-Patterns & The Rule Of Three

Remember that **not every algorithm, best-practice or solution represents** what might be considered a complete pattern. There may be a few key ingredients here that are missing and the pattern community is generally weary of something claiming to be one unless it has been heavily vetted. Even if something is presented to us which **appears** to meet the criteria for a pattern, it should not be considered one until it has undergone suitable periods of scrutiny and testing by others.

Looking back upon the work by Alexander once more, he claims that a pattern should both be a process and a 'thing'. This definition is obtuse on purpose as he follows by saying that it is the process should create the 'thing'. This is a reason why patterns generally focus on addressing a visually identifiable structure i.e you should be able to visually depict (or draw) a picture representing the structure that placing the pattern into practice results in.

In studying design patterns, you may come across the term 'proto-pattern' quite frequently. What is this? Well, a pattern that has not yet been known to pass the 'pattern'-ity tests is usually referred to as a proto-pattern. Proto-patterns may result from the work of someone that has established a particular solution that is worthy of sharing with the community, but may not have yet had the opportunity to have been vetted heavily due to its very young age.

Alternatively, the individual(s) sharing the pattern may not have the time or interest of going through the 'pattern'-ity process and might release a short description of their proto-pattern instead. Brief descriptions or snippets of this type of pattern are known as patlets.

The work involved in fully documenting a qualified pattern can be quite daunting. Looking back at some of the earliest work in the field of design patterns, a pattern may be considered 'good' if it does the following:

- **Solves a particular problem:** Patterns are not supposed to just capture

principles or strategies. They need to capture solutions. This is one of the most essential ingredients for a good pattern.

- **The solution to this problem cannot be obvious:** You can often find that problem-solving techniques attempt to derive from well-known first principles. The best design patterns usually provide solutions to problems indirectly - this is considered a necessary approach for the most challenging problems related to design.
- **The concept described must have been proven:** Design patterns require proof that they function as described and without this proof the design cannot be seriously considered. If a pattern is highly speculative in nature, only the brave may attempt to use it.
- **It must describe a relationship:** In some cases it may appear that a pattern describes a type of module. Although an implementation may appear this way, the official description of the pattern must describe much deeper system structures and mechanisms that explain its relationship to code.

We would be forgiven for thinking that a proto-pattern which fails to meet guidelines isn't worth learning from, however, this is far from the truth. Many proto-patterns are actually quite good. I'm not saying that all proto-patterns are worth looking at, but there are quite a few useful ones in the wild that could assist you with future projects. Use best judgment with the above list in mind and you'll be fine in your selection process.

One of the additional requirements for a pattern to be valid is that they display some recurring phenomenon. This is often something that can be qualified in at least three key areas, referred to as the rule of three. To show recurrence using this rule, one must demonstrate:

1. **Fitness of purpose** - how is the pattern considered successful?
2. **Usefulness**- why is the pattern considered successful?
3. **Applicability** - is the design worthy of being a pattern because it has wider applicability? If so, this needs to be explained. When reviewing or defining a pattern, it is important to keep the above in mind.

Chapter 4. The Structure Of A Design Pattern

You may be curious about how a pattern author might approach outlining structure, implementation and purpose of a new pattern. Traditionally, a pattern is initially be presented in the form of a **rule** that establishes a relationship between:

- A **context**
- A system of **forces** that arises in that context and
- A **configuration** that allows these forces to resolve themselves in context

With this in mind, lets now take a look at a summary of the component elements for a design pattern. A design pattern should have a:

- **Pattern name** and a **description**
- **Context outline** – the contexts in which the pattern is effective in responding to the users needs.
- **Problem statement** – a statement of the problem being addressed so we can understand the intent of the pattern.
- **Solution** – a description of how the user’s problem is being solved in an understandable list of steps and perceptions.
- **Design** – a description of the pattern’s design and in particular, the user’s behavior in interacting with it
- **Implementation**– a guide to how the pattern would be implemented
- **Illustrations** – a visual representation of classes in the pattern (e.g. a diagram))
- **Examples** – an implementation of the pattern in a minimal form
- **Co-requisites** – what other patterns may be needed to support use of the pattern being described?
- **Relations** – what patterns does this pattern resemble? does it closely mimic any others?

- **Known usage** – is the pattern being used in the ‘wild’?. If so, where and how?
- **Discussions** – the team or author’s thoughts on the exciting benefits of the pattern

Design patterns are quite a powerful approach to getting all of the developers in an organization or team on the same page when creating or maintaining solutions. If you or your company ever consider working on your own pattern, remember that although they may have a heavy initial cost in the planning and write-up phases, the value returned from that investment can be quite worth it. Always research thoroughly before working on new patterns however, as you may find it more beneficial to use or build on top of existing proven patterns than starting afresh.

Chapter 5. Writing Design Patterns

Although this book is aimed at those new to design patterns, a fundamental understanding of how a design pattern is written can offer you a number of useful benefits. For starters, you can gain a deeper appreciation for the reasoning behind a pattern being needed. You can also learn how to tell if a pattern (or proto-pattern) is up to scratch when reviewing it for your own needs.

Writing good patterns is a challenging task. Patterns not only need to provide a substantial quantity of reference material for end-users (such as the items found in the structure section above), but they also need to be able to defend why they are necessary. If you've already read the previous section on 'what' a pattern is, you may think that this in itself should help you identify patterns when you see them in the wild. This is actually quite the opposite - you can't always tell if a piece of code you're inspecting follows a pattern.

When looking at a body of code that you think may be using a pattern, you might write down some of the aspects of the code that you believe falls under a particular existing pattern. In many cases of pattern-analysis you'll find that you're just looking at code that follows good principles and design practices that could happen to overlap with the rules for a pattern by accident. Remember - solutions in which neither interactions nor defined rules appear are not patterns.

If you're interested in venturing down the path of writing your own design patterns I recommend learning from others who have already been through the process and done it well. Spend time absorbing the information from a number of different design pattern descriptions and books and take in what's meaningful to you - this will help you accomplish the goals you have laid out for yours. Explore structure and semantics - this can be done by examining the interactions and context of the patterns you are interested in so you can identify the principles that assist in organizing those patterns together in useful configurations.

Once you've exposed yourself to a wealth of information on pattern literature, you may wish to begin your pattern using an existing format and see if you can brainstorm new ideas for improving it or integrating your ideas in there. An example of someone that did this in recent years is Christian Heilmann, who

took the existing module pattern and made some fundamentally useful changes to it to create the revealing module pattern (this is one of the patterns covered later in this book).

If you would like to try your hand at writing a design pattern (even if just for the learning experience of going through the process), the tips I have for doing so would be as follows:

- **Bear in mind practicability:** Ensure that your pattern describes proven solutions to recurring problems rather than just speculative solutions which haven't been qualified.
- **Ensure that you draw upon best practices:** The design decisions you make should be based on principles you derive from an understanding of best-practices.
- **Your design patterns should be transparent to the user:** Design patterns should be entirely transparent to any type of user-experience. They are primarily there to serve the developers using them and should not force changes to behavior in the user-experience that would not be incurred without the use of a pattern.
- **Remember that originality is *not* key in pattern design:** When writing a pattern, you do not need to be the original discoverer of the solutions being documented nor do you have to worry about your design overlapping with minor pieces of other patterns. If your design is strong enough to have broad useful applicability, it has a chance of being recognized as a proper pattern.
- **Know the differences between patterns and design:** A design pattern generally draws from proven best practice and serves as a model for a designer to create a solution. The role of the pattern is to give designers guidance to make the best design choices so they can cater to the needs of their users.
- **Your pattern needs to have a strong set of examples:** A good pattern description needs to be followed by an equally strong set of examples demonstrating the successful application of your pattern. To show broad usage, examples that exhibit good design principles are ideal.

Pattern writing is a careful balance between creating a design that is general, specific and above all, useful. Try to ensure that if writing a pattern you cover the

widest possible areas of application and you should be fine. I hope that this brief introduction to writing patterns has given you some insights that will assist your learning process for the next sections of this book.

Chapter 6. Anti-Patterns

If we consider that a pattern represents a best-practice, an anti-pattern represents a lesson that has been learned. The term anti-patterns was coined in 1995 by Andrew Koenig in the November C++ Report that year, inspired by the GoF's book Design Patterns. In Koenig's report, there are two notions of anti-patterns that are presented. Anti-Patterns:

- Describe a bad solution to a particular problem which resulted in a bad situation occurring
- Describe how to get out of said situation and how to go from there to a good solution

On this topic, Alexander writes about the difficulties in achieving a good balance between good design structure and good context:

“These notes are about the process of design; the process of inventing physical things which display a new physical order, organization, form, in response to function....every design problem begins with an effort to achieve fitness between two entities: the form in question and its context. The form is the solution to the problem; the context defines the problem”.

While it's quite important to be aware of design patterns, it can be equally important to understand anti-patterns. Let us qualify the reason behind this. When creating an application, a project's life-cycle begins with construction however once you've got the initial release done, it needs to be maintained. The quality of a final solution will either be good or bad, depending on the level of skill and time the team have invested in it. Here good and bad are considered in context - a 'perfect' design may qualify as an anti-pattern if applied in the wrong context.

The bigger challenges happen after an application has hit production and is ready to go into maintenance mode. A developer working on such a system who hasn't worked on the application before may introduce a bad design into the project by accident. If said bad practices are created as anti-patterns, they allow developers a means to recognize these in advance so that they can avoid common mistakes that can occur - this is parallel to the way in which design patterns provide us with a

way to recognize common techniques that are useful.

To summarize, an anti-pattern is a bad design that is worthy of documenting.

Examples of anti-patterns in JavaScript are the following:

- Polluting the global namespace by defining a large number of variables in the global context
- Passing strings rather than functions to either `setTimeout` or `setInterval` as this triggers the use of `eval()` internally.
- Modifying the `Object` class prototype (this is a particularly bad anti-pattern)
- Using JavaScript in an inline form as this is inflexible
- The use of `document.write` where native DOM alternatives such as `document.createElement` are more appropriate. `document.write` has been grossly misused over the years and has quite a few disadvantages including that if it's executed after the page has been loaded it can actually overwrite the page you're on, whilst `document.createElement` does not. You can see [here](#) for a live example of this in action. It also doesn't work with XHTML which is another reason opting for more DOM-friendly methods such as `document.createElement` is favorable.

Knowledge of anti-patterns is critical for success. Once you are able to recognize such anti-patterns, you will be able to refactor your code to negate them so that the overall quality of your solutions improves instantly.

Chapter 7. Categories Of Design Pattern

A glossary from the well-known design book, Domain-Driven Terms, rightly states that:

“A design pattern names, abstracts, and identifies the key aspects of a common design structure that make it useful for creating a reusable object-oriented design. The design pattern identifies the participating classes and their instances, their roles and collaborations, and the distribution of responsibilities.

Each design pattern focuses on a particular object-oriented design problem or issue. It describes when it applies, whether or not it can be applied in view of other design constraints, and the consequences and trade-offs of its use. Since we must eventually implement our designs, a design pattern also provides sample ... code to illustrate an implementation.

Although design patterns describe object-oriented designs, they are based on practical solutions that have been implemented in mainstream object-oriented programming languages”

Design patterns can be broken down into a number of different categories. In this section we'll review three of these categories and briefly mention a few examples of the patterns that fall into these categories before exploring specific ones in more detail.

Creational Design Patterns

Creational design patterns focus on handling object creation mechanisms where objects are created in a manner suitable for the situation you are working in. The basic approach to object creation might otherwise lead to added complexity in a project whilst these patterns aim to solve this problem by controlling the creation process.

Some of the patterns that fall under this category are: Constructor, Factory, Abstract, Prototype, Singleton and Builder.

Structural Design Patterns

Structural patterns are concerned with object composition and typically identify simple ways to realize relationships between different objects. They help ensure that when one part of a system changes, the entire structure of the system doesn't need to do the same. They also assist in recasting parts of the system which don't fit a particular purpose into those that do.

Patterns that fall under this category include: Decorator, Facade, Flyweight, Adapter and Proxy.

Behavioral Design Patterns

Behavioral patterns focus on improving or streamlining the communication between disparate objects in a system.

Some behavioral patterns include: Iterator, Mediator, Observer and Visitor.

Chapter 8. Design Pattern Categorization

In my early experiences of learning about design patterns, I personally found the following table a very useful reminder of what a number of patterns has to offer - it covers the 23 Design Patterns mentioned by the GoF. The original table was summarized by Elyse Nielsen back in 2004 and I've modified it where necessary to suit our discussion in this section of the book.

I recommend using this table as reference, but do remember that there are a number of additional patterns that are not mentioned here but will be discussed later in the book.

A brief note on classes

Keep in mind that there will be patterns in this table that reference the concept of 'classes'. JavaScript is a class-less language, however classes can be simulated using functions.

The most common approach to achieving this is by defining a JavaScript function where we then create an object using the `new` keyword. `this` can be used to help define new properties and methods for the object as follows:

```
// A car 'class'
function Car(model) {
  this.model = model;
  this.color = 'silver';
  this.year = '2012';
  this.getInfo = function () {
    return this.model + ' ' + this.year;
  }
}
```

We can then instantiate the object using the `Car` constructor we defined above like this:

```
var myCar = new Car('ford');
myCar.year = '2010';
console.log(myCar.getInfo());
```

For more ways to define 'classes' using JavaScript, see Stoyan Stefanov's useful

post on them.

Let us now proceed to review the table.

Creational	Based on the concept of creating an object.						
Class							
Factory Method	This makes an instance of several derived classes based on interfaced data or events.						
Object							
Abstract Factory	Creates an instance of several families of classes without detailing concrete classes.						
Builder	Separates object construction from its representation, always creates the same type of object.						
Prototype	A fully initialized instance used for copying or cloning.						
Singleton	A class with only a single instance with global access points.						
Structural	Based on the idea of building blocks of objects						
Class							
Adapter	Match interfaces of different classes therefore classes can work together despite incompatible interfaces						
Object							
Adapter	Match interfaces of different classes therefore classes can work together despite incompatible interfaces						
Bridge	Separates an object's interface from its implementation so the two can vary independently						
Composite	A structure of simple and composite objects which makes the total object more than just the sum of its parts.						
Decorator	Dynamically add alternate processing to objects.						
Facade	A single class that hides the complexity of an entire subsystem.						
Flyweight	A fine-grained instance used for efficient sharing of information that is contained elsewhere.						
Proxy	A place holder object representing the true object						
Behavioral	Based on the way objects play and work together.						
Class							
Interpreter	A way to include language elements in an application to match the						

	grammar of the intended language.								
Template Method	Creates the shell of an algorithm in a method, then defer the exact steps to a subclass.								
Object									
Chain of Responsibility	A way of passing a request between a chain of objects to find the object that can handle the request.								
Command	Encapsulate a command request as an object to enable, logging and/or queuing of requests, and provides error-handling for unhandled requests.								
Iterator	Sequentially access the elements of a collection without knowing the inner workings of the collection.								
Mediator	Defines simplified communication between classes to prevent a group of classes from referring explicitly to each other.								
Memento	Capture an object's internal state to be able to restore it later.								
Observer	A way of notifying change to a number of classes to ensure consistency between the classes.								
State	Alter an object's behavior when its state changes								
Strategy	Encapsulates an algorithm inside a class separating the selection from the implementation								
Visitor	Adds a new operation to a class without changing the class								

Chapter 9. JavaScript Design Patterns

We are now going to explore JavaScript implementations of a number of both classical and modern design patterns. This section of the book will cover an introduction to these patterns, whilst the next section will focus on looking at some select patterns in greater detail.

A common question developers regularly ask is what the 'ideal' set of patterns they should be using are. There isn't a singular answer to this question, but with the aid of what you'll learn in this book, you will hopefully be able to use your best judgment to select the right patterns to best suit your project's needs.

The patterns we will be exploring in this section are the:

- [Creational Pattern](#)
- [Constructor Pattern](#)
- [Singleton Pattern](#)
- [Module Pattern](#)
- [Revealing Module Pattern](#)
- [Observer Pattern](#)
- [Mediator Pattern](#)
- [Prototype Pattern](#)
- [Command Pattern](#)
- [Facade Pattern](#)
- [Factory Pattern](#)
- [Mixin Pattern](#)
- [Decorator Pattern](#)
- [Flyweight Pattern](#)

The Creational Pattern

The Creational pattern forms the basis for a number of the other design patterns we'll be reviewing in this section and could be considered the easiest to understand. It deals with the idea of creating new things, specifically new objects. In JavaScript, the three common ways to create new objects are as follows:

```
// Each of the following options will create a new empty object:

var newObject = {}; // or

var newObject = Object.create(null); // or

var newObject = new Object();
```

Where the 'Object' constructor creates an object wrapper for a specific value, or where no value is passed, it will create an empty object and return it.

There are then four ways in which keys and values can then be assigned to an object:

```
// ECMAScript 3 compatible approaches

// 1. Dot syntax
newObject.someKey = 'Hello World'; // Write properties
var key = newObject.someKey; // Access properties

// 2. Square bracket syntax
newObject['someKey'] = 'Hello World'; // Write properties
var key = newObject['someKey']; // Access properties

// ECMAScript 5 only compatible approaches
// For more information see: http://kangax.github.com/es5-compat-table/

// 3. Object.defineProperty
Object.defineProperty(newObject, "someKey", {
  value: "for more control of the property's behavior",
  writable: true,
  enumerable: true,
  configurable: true
});

// If the above feels a little difficult to read, a short-hand could
// be written as follows:

var defineProp = function ( obj, key, value ){
  config.value = value;
  Object.defineProperty(obj, key, config);
}

// Create a new empty object
var man = Object.create(null);
```

```
// Populate the object with properties

defineProp( man, 'car', 'Delorean' );
defineProp( man, 'dob', '1981' );
defineProp( man, 'beard', false );
```

```
// 4. Object.defineProperties

Object.defineProperties(newObject, {
  "someKey": {
    value: "Hello World",
    writable: true
  },
  "anotherKey": {
    value: "Foo bar",
    writable: false
  }
});
```

As we will see a little later in the book, these methods can even be used for inheritance, as follows:

```
var driver = Object.create(man);
defineProp(driver, 'topSpeed', '100mph');
driver.topSpeed // 100mph
```

The Constructor Pattern

The phrase 'constructor' is familiar to most developers, however if you're a beginner it can be useful to review what a constructor is before we get into talking about a pattern dedicated to it.

Constructors are used to create specific types of objects - they both prepare the object for use and can also accept parameters which the constructor uses to set the values of member variables when the object is first created. The idea that a constructor is a paradigm can be found in the majority of programming languages, including JavaScript. You're also able to define custom constructors that define properties and methods for your own types of objects.

Basic Constructors

In JavaScript, constructor functions are generally considered a reasonable way to implement instances. As we saw earlier, JavaScript doesn't support the concept of classes but it does support special constructor functions. By simply prefixing a call to a constructor function with the keyword 'new', you can tell JavaScript you would like function to behave like a constructor and instantiate a new object with the members defined by that function. Inside a constructor, the keyword 'this' references the new object that's being created. Again, a very basic constructor may be:

```
function Car(model, year, miles) {  
  
    this.model = model;  
    this.year = year;  
    this.miles = miles;  
  
    this.toString = function () {  
        return this.model + " has done " + this.miles + " miles";  
    };  
}  
  
var civic = new Car("Honda Civic", 2009, 20000);  
var mondeo = new Car("Ford Mondeo", 2010, 5000);  
  
console.log(civic.toString());  
console.log(mondeo.toString());
```

The above is a simple version of the constructor pattern but it does suffer from some problems. One is that it makes inheritance difficult and the other is that

functions such as `toString()` are redefined for each of the new objects created using the `Car` constructor. This isn't very optimal as the function should ideally be shared between all of the instances of the `Car` type.

Constructors With Prototypes

Functions in JavaScript have a property called a prototype. When you call a JavaScript constructor to create an object, all the properties of the constructor's prototype are then made available to the new object. In this fashion, multiple `Car` objects can be created which access the same prototype. We can thus extend the original example as follows:

```
function Car(model, year, miles) {  
    this.model = model;  
    this.year = year;  
    this.miles = miles;  
}  
  
// Note here that we are using Object.prototype.newMethod rather than  
// Object.prototype so as to avoid redefining the prototype object  
Car.prototype.toString = function () {  
    return this.model + " has done " + this.miles + " miles";  
};  
  
var civic = new Car("Honda Civic", 2009, 20000);  
var mondeo = new Car("Ford Mondeo", 2010, 5000);  
  
console.log(civic.toString());
```

Here, a single instance of `toString()` will now be shared between all of the `Car` objects.

The Singleton Pattern

In conventional software engineering, the singleton pattern can be implemented by creating a class with a method that creates a new instance of the class if one doesn't exist. In the event of an instance already existing, it simply returns a reference to that object.

The singleton pattern is thus known because traditionally, it restricts instantiation of a class to a single object. With JavaScript, singletons serve as a namespace provider which isolate implementation code from the global namespace so-as to provide a single point of access for functions.

The singleton doesn't provide a way for code that doesn't know about a previous reference to the singleton to easily retrieve it - it is not the object or 'class' that's returned by a singleton, it's a structure. Think of how closed variables aren't actually closures - the function scope that provides the closure is the closure.

Singletons in JavaScript can take on a number of different forms and researching this pattern online is likely to result in at least 10 different variations. In its simplest form, a singleton in JS can be an object literal grouped together with its related methods and properties as follows:

```
var mySingleton = {  
  property1: "something",  
  
  property2: "something else",  
  
  method1: function () {  
    console.log('hello world');  
  }  
};
```

If you wished to extend this further, you could add your own private members and methods to the singleton by encapsulating variable and function declarations inside a closure. Exposing only those which you wish to make public is quite straightforward from that point as demonstrated below:

```
var mySingleton = function () {  
  
  // here are our private methods and variables  
  var privateVariable = 'something private';  
  
  function showPrivate() {
```



```

        console.log(privateVariable);
    }

    // public variables and methods (which can access
    // private variables and methods )
    return {

        publicMethod: function () {
            showPrivate();
        },

        publicVar: 'the public can see this!'

    };
};

var single = mySingleton();
single.publicMethod(); // logs 'something private'
console.log(single.publicVar); // logs 'the public can see this!'

```

The above example is great, but let's next consider a situation where you only want to instantiate the singleton when it's needed. To save on resources, you can place the instantiation code inside another constructor function as follows:

```

var Singleton = (function () {
    var instantiated;

    function init() {
        // singleton here
        return {
            publicMethod: function () {
                console.log('hello world');
            },
            publicProperty: 'test'
        };
    }

    return {
        getInstance: function () {
            if (!instantiated) {
                instantiated = init();
            }
            return instantiated;
        }
    };
})();

// calling public methods is then as easy as:
Singleton.getInstance().publicMethod();

```

So, where else is the singleton pattern useful in practice?. Well, it's quite useful when exactly one object is needed to coordinate patterns across the system. Here's one last example of the singleton pattern being used:

```

var SingletonTester = (function () {

    // options: an object containing configuration options for the singleton
    // e.g var options = { name: 'test', pointX: 5};
    function Singleton(options) {

        // set options to the options supplied or an empty object if none provided.
        options = options || {};
        //set the name parameter
        this.name = 'SingletonTester';
        //set the value of pointX
        this.pointX = args.pointX || 6;
        //set the value of pointY
        this.pointY = args.pointY || 10;

    }

    // this is our instance holder
    var instance;

    // this is an emulation of static variables and methods
    var _static = {

        name: 'SingletonTester',

        // This is a method for getting an instance
        // It returns a singleton instance of a singleton object
        getInstance: function (options) {
            if (instance === undefined) {
                instance = new Singleton(options);
            }
            return instance;
        }
    };
    return _static;
})();

var singletonTest = SingletonTester.getInstance({
    pointX: 5
});

console.log(singletonTest.pointX); // outputs 5

```

The Module Pattern

Modules

Modules are an integral piece of any robust application's architecture and typically help in keeping the code for a project organized. In JavaScript, there are several options for implementing modules including both the well-known module pattern as well as object literal notation.

The module pattern is based in part on object literals and so it makes sense to review them first.

Object Literals

In object literal notation, an object is described as a set of comma-separated name/value pairs enclosed in curly braces (`{}`). Names inside the object may be either strings or identifiers that are followed by a colon. There should be no comma used after the final name/value pair in the object as this may result in errors.

```
var myObjectLiteral = {  
  variableKey: variableValue,  
  functionKey: function(){  
    // ...  
  }  
};
```

Object literals don't require instantiation using the `new` operator but shouldn't be used at the start of a statement as the opening `{` may be interpreted as the beginning of a block. Outside of an object, new members may be added to it using assignment as follows `myModule.property = 'someValue'`;

Below we can see a more complete example of a module defined using object literal syntax:

```
var myModule = {  
  myProperty: 'someValue',  
  // object literals can contain properties and methods.  
  // here, another object is defined for configuration  
  // purposes:  
  myConfig: {  
    useCaching: true,  
    language: 'en'  
  },  
  // a very basic method  
  myMethod: function () {  
    console.log('I can haz functionality?');  
  }  
};
```

```

    },
    // output a value based on current configuration
    myMethod2: function () {
        console.log('Caching is:' + (this.myConfig.useCaching) ? 'enabled' :
'disabled');
    },
    // override the current configuration
    myMethod3: function (newConfig) {
        if (typeof newConfig == 'object') {
            this.myConfig = newConfig;
            console.log(this.myConfig.language);
        }
    }
};

myModule.myMethod(); // I can haz functionality
myModule.myMethod2(); // outputs enabled
myModule.myMethod3({
    language: 'fr',
    useCaching: false
}); // fr

```

Using object literals can assist in encapsulating and organizing your code and Rebecca Murphey has previously written about this topic in [depth](#) should you wish to read into object literals further.

That said, if you're opting for this technique, you may be equally as interested in the module pattern. It still uses object literals but only as the return value from a scoping function.

The Module Pattern

The module pattern was originally defined as a way to provide both private and public encapsulation for classes in conventional software engineering.

In JavaScript, the module pattern is used to further emulate the concept of classes in such a way that we're able to include both public/private methods and variables inside a single object, thus shielding particular parts from the global scope. What this results in is a reduction in the likelihood of your function names conflicting with other functions defined in additional scripts on the page.

Privacy

The module pattern encapsulates 'privacy', state and organization using closures. It provides a way of wrapping a mix of public and private methods and variables, protecting pieces from leaking into the global scope and accidentally colliding with

another developer's interface. With this pattern, only a public API is returned, keeping everything else within the closure private.

This gives us a clean solution for shielding logic doing the heavy lifting whilst only exposing an interface you wish other parts of your application to use. The pattern is quite similar to an immediately-invoked functional expression (**IIFE** - see the section on namespacing patterns for more on this) except that an object is returned rather than a function.

It should be noted that there isn't really an explicitly true sense of 'privacy' inside JavaScript because unlike some traditional languages, it doesn't have access modifiers. Variables can't technically be declared as being public nor private and so we use function scope to simulate this concept. Within the module pattern, variables or methods declared are only available inside the module itself thanks to closure. Variables or methods defined within the returning object however are available to everyone.

History

From a historical perspective, the module pattern was originally developed by a number of people including **Richard Cornford** in 2003. It was later popularized by Douglas Crockford in his lectures. Another piece of trivia is that if you've ever played with Yahoo's YUI library, some of its features may appear quite familiar and the reason for this is that the module pattern was a strong influence for YUI when creating their components.

Examples

Let's begin looking at an implementation of the module pattern by creating a module which is self-contained.

```
var testModule = (function () {  
    var counter = 0;  
    return {  
        incrementCounter: function () {  
            return counter++;  
        },  
        resetCounter: function () {  
            console.log('counter value prior to reset:' + counter);  
            counter = 0;  
        }  
    };  
})();
```

```
// test
testModule.incrementCounter();
testModule.resetCounter();
```

Here, other parts of the code are unable to directly read the value of our `incrementCounter()` or `resetCounter()`. The counter variable is actually fully shielded from our global scope so it acts just like a private variable would - its existence is limited to within the module's closure so that the only code able to access its scope are our two functions. Our methods are effectively namespaced so in the test section of our code, we need to prefix any calls with the name of the module (e.g. 'testModule').

When working with the module pattern, you may find it useful to define a simple template that you use for getting started with it. Here's one that covers namespacing, public and private variables:

```
var myNamespace = (function () {

    var myPrivateVar = 0;
    var myPrivateMethod = function (someText) {
        console.log(someText);
    };

    return {

        myPublicVar: "foo",

        myPublicFunction: function (bar) {
            myPrivateVar++;
            myPrivateMethod(bar);
        }
    };

})();
```

Looking at another example, below we can see a shopping basket implemented using the this pattern. The module itself is completely self-contained in a global variable called `basketModule`. The `basket` array in the module is kept private and so other parts of your application are unable to directly read it. It only exists with the module's closure and so the only methods able to access it are those with access to its scope (ie. `addItem()`, `getItem()` etc).

```
var basketModule = (function () {
    var basket = []; //private
    function doSomethingPrivate() {
        //...
    }
})
```

```

function doSomethingElsePrivate() {
  //...
}
return { //exposed to public
  addItem: function (values) {
    basket.push(values);
  },
  getItemCount: function () {
    return basket.length;
  },
  doSomething: doSomethingPrivate(),
  getTotal: function () {
    var q = this.getItemCount(),
        p = 0;
    while (q--) {
      p += basket[q].price;
    }
    return p;
  }
}
})();

```

Inside the module, you'll notice we return an **object**. This gets automatically assigned to **basketModule** so that you can interact with it as follows:

```

// basketModule is an object with properties which can also be methods
basketModule.addItem({
  item: 'bread',
  price: 0.5
});
basketModule.addItem({
  item: 'butter',
  price: 0.3
});

console.log(basketModule.getItemCount());
console.log(basketModule.getTotal());

// however, the following will not work:
console.log(basketModule.basket); // (undefined as not inside the returned object)
console.log(basket); //(only exists within the scope of the closure)

```

The methods above are effectively namespaced inside **basketModule**.

Notice how the scoping function in the above basket module is wrapped around all of our functions, which we then call and immediately store the return value of.

This has a number of advantages including:

- The freedom to have private functions which can only be consumed by our module. As they aren't exposed to the rest of the page (only our exported API is), they're considered truly private.

- Given that functions are declared normally and are named, it can be easier to show call stacks in a debugger when we're attempting to discover what function(s) threw an exception.
- As T.J Crowder has pointed out in the past, it also enables us to return different functions depending on the environment. In the past, I've seen developers use this to perform UA testing in order to provide a code-path in their module specific to IE, but we can easily opt for feature detection these days to achieve a similar goal.

Advantages

We've seen why the singleton pattern can be useful, but why is the module pattern a good choice? For starters, it's a lot cleaner for developers coming from an object-oriented background than the idea of true encapsulation, at least from a JavaScript perspective.

Secondly, it supports private data - so, in the module pattern, public parts of your code are able to touch the private parts, however the outside world is unable to touch the class's private parts (no laughing! Oh, and thanks to David Engfer for the joke).

Disadvantages

The disadvantages of the module pattern are that as you access both public and private members differently, when you wish to change visibility, you actually have to make changes to each place the member was used.

You also can't access private members in methods that are added to the object at a later point. That said, in many cases the module pattern is still quite useful and when used correctly, certainly has the potential to improve the structure of your application.

Other disadvantages include the inability to create automated unit tests for private members and additional complexity when bugs require hot fixes. It's simply not possible to patch privates. Instead, one must override all public methods which interact with the buggy privates. Developers can't easily extend privates either, so it's worth remembering privates are not as flexible as they may initially appear.

Before we dive into how the module pattern may be implemented using different

JavaScript frameworks, here's a simple template for using the pattern:

```
var someModule = (function () {

    // private attributes
    var privateVar = 5;

    // private methods
    var privateMethod =
    function () {
        return 'Private Test';
    };

    return {
        // public attributes
        publicVar: 10,
        // public methods
        publicMethod:
        function () {
            return ' Followed By Public Test ';
        },

        // let's access the private members
        getData:
        function () {
            return privateMethod() + this.publicMethod() + privateVar;
        }
    }
})(); // the parens here cause the anonymous function to execute and return

someModule.getData();
```

How about the module pattern implemented in specific toolkits or frameworks?

Dojo

Dojo provides a convenience method for working with objects called `dojo.setObject()`. This takes as its first argument a dot-separated string such as `myObj.parent.child` which refers to a property called 'child' within an object 'parent' defined inside 'myObj'. Using `setObject()` allows us to set the value of children, creating any of the intermediate objects in the rest of the path passed if they don't already exist.

For example, if we wanted to declare `basket.core` as an object of the `store` namespace, this could be achieved as follows using the traditional way:

```
var store = window.store || {};
```

```
if (!store["basket"]) {
    store.basket = {};
}
if (!store.basket["core"]) {
```

```

    store.basket.core = {};
}

store.basket.core = {
    // ...rest of our logic
}

```

Or as follows using Dojo 1.7 (AMD-compatible version) and above:

```

require(["dojo/_base/customStore"], function(store){

    // using dojo.setObject()
    store.setObject("basket.core", (function() {
        var basket = [];
        function privateMethod() {
            console.log(basket);
        }
        return {
            publicMethod: function(){
                privateMethod();
            }
        };
    }()));

});

```

For more information on `dojo.setObject()`, see the official [documentation](#).

ExtJS

For those using Sencha's ExtJS, you're in for some luck as the official documentation incorporates [examples](#) that do demonstrate how to correctly use the module pattern with the framework.

Below we can see an example of how to define a namespace which can then be populated with a module containing both a private and public API. With the exception of some semantic differences, it's quite close to how the module pattern is implemented in vanilla JavaScript:

```

// create namespace
Ext.namespace('myNameSpace');

// create application
myNameSpace.app = function () {
    // do NOT access DOM from here; elements don't exist yet
    // private variables
    var btn1;
    var privVar1 = 11;

    // private functions
    var btn1Handler = function (button, event) {
        alert('privVar1=' + privVar1);
        alert('this.btn1Text=' + this.btn1Text);
    };
}

```

```

};

// public space
return {
  // public properties, e.g. strings to translate
  btn1Text: 'Button 1',

  // public methods
  init: function () {
    if (Ext.Ext2) {
      btn1 = new Ext.Button({
        renderTo: 'btn1-ct',
        text: this.btn1Text,
        handler: btn1Handler
      });
    } else {
      btn1 = new Ext.Button('btn1-ct', {
        text: this.btn1Text,
        handler: btn1Handler
      });
    }
  }
};
}(); // end of app

```

YUI

Similarly, we can also implement the module pattern when building applications using YUI3. The following example is heavily based on the original YUI module pattern implementation by Eric Miraglia, but again, isn't vastly different from the vanilla JavaScript version:

```

Y.namespace('store.basket') = (function () {

  // private variables:
  var myPrivateVar = "I can be accessed only within Y.store.basket .";

  // private method:
  var myPrivateMethod = function () {
    Y.log("I can be accessed only from within YAHOO.store.basket");
  }

  return {
    myPublicProperty: "I'm a public property.",
    myPublicMethod: function () {
      Y.log("I'm a public method.");

      // Within basket, I can access "private" vars and methods:
      Y.log(myPrivateVar);
      Y.log(myPrivateMethod());

      // The native scope of myPublicMethod is store so we can
      // access public members using "this":
      Y.log(this.myPublicProperty);
    }
  };
})();

```

```
    }  
};  
})();
```

jQuery

There are a number of ways in which jQuery code unspecific to plugins can be wrapped inside the module pattern. Ben Cherry previously suggested an implementation where a function wrapper is used around module definitions in the event of there being a number of commonalities between modules.

In the following example, a `library` function is defined which declares a new library and automatically binds up the `init` function to `document.ready` when new libraries (ie. modules) are created.

```
function library(module) {  
    $(function () {  
        if (module.init) {  
            module.init();  
        }  
    });  
    return module;  
}  
  
var myLibrary = library(function () {  
    return {  
        init: function () {  
            /*implementation*/  
        }  
    };  
})();
```

For further reading on the module pattern, see Ben Cherry's article on it [here](#).

The Revealing Module Pattern

Now that we're a little more familiar with the Module pattern, let's take a look at a slightly improved version - Christian Heilmann's Revealing Module pattern.

The Revealing Module Pattern came about as Heilmann (now at Mozilla) was frustrated with the fact that if you had to repeat the name of the main object when you wanted to call one public method from another or access public variables. He also disliked the Module pattern's requirement for having to switch to object literal notation for the things you wished to make public.

The result of his efforts were an updated pattern where you would simply define all of your functions and variables in the private scope and return an anonymous object at the end of the module along with pointers to both the private variables and functions you wished to reveal as public.

Advantages

Once again, you're probably wondering what the benefits of this approach are. The Revealing Module Pattern allows the syntax of your script to be fairly consistent - it also makes it very clear at the end which of your functions and variables may be accessed publicly, something that is quite useful. In addition, you are also able to reveal private functions with more specific names if you wish.

An example of how to use the revealing module pattern can be found below:

```
var myRevealingModule = (function(){  
  
    var name = 'John Smith';  
    var age = 40;  
  
    function updatePerson(){  
        name = 'John Smith Updated';  
    }  
    function setPerson () {  
        name = 'John Smith Set';  
    }  
    function getPerson () {  
        return name;  
    }  
    return {  
        set: setPerson,  
        get: getPerson  
    };  
})();
```

```
// Sample usage:  
myRevealingModule.get();
```

Disadvantages

A disadvantage of this pattern is that if a private function refers to a public function, that public function can't be overridden if a patch is necessary. This is because the private function will continue to refer to the private implementation and the pattern doesn't apply to public members, only to functions.

It's only members with objects as values that can be used because of pass-by-value rules. However, public object members which refer to private variables are also subject to the no-patch rule notes above. As a result of this, modules created with the revealing module pattern are inherently more fragile than those created with the original module pattern.

The Observer Pattern

The Observer pattern is more popularly known these days as the Publish/Subscribe pattern. It is a design pattern which allows an object (known as a subscriber) to watch another object (the publisher), where we provide a means for the subscriber and publisher form a listen and broadcast relationship.

Subscribers are able to register (subscribe) to receive topic notifications from the publisher when something interesting happens. When the publisher needs to notify observers about interesting topics, it broadcasts (publishes) a notification of these to each observer (which can include specific data related to the topic).

When subscribers are no longer interested in being notified of topics by the publisher they are registered with, they can unregister (or unsubscribe) themselves. The subject will then in turn remove them from the observer collection.

The general idea here is the promotion of loose coupling. Rather than single objects calling on the methods of other objects directly, they instead subscribe to a specific task or activity of another object and are notified when it occurs.

It's often useful to refer back to published definitions of design patterns that are language agnostic to get a broader sense of their usage and advantages over time. The definition of the observer pattern provided in the GoF book, *Design Patterns: Elements of Reusable Object-Oriented Software*, is:

'One or more observers are interested in the state of a subject and register their interest with the subject by attaching themselves. When something changes in our subject that the observer may be interested in, a notify message is sent which calls the update method in each observer. When the observer is no longer interested in the subject's state, they can simply detach themselves.'

Unlimited numbers of objects may observe topics in the subject by registering themselves. Once registered to particular events, the subject will notify all observers when the topic has been fired.

Further motivation behind using the observer pattern is where you need to maintain consistency between related objects without making classes tightly coupled. For example, when an object needs to be able to notify other objects without making assumptions regarding those objects. Another use case is where

abstractions have more than one aspect, where one depends on the other. The encapsulation of these aspects in separate objects allows the variation and re-use of the objects independently.

Advantages

Arguably, the largest benefit of using Pub/Sub is the ability to break down our applications into smaller, more loosely coupled modules, which can also improve general manageability.

It is also a pattern that encourages us to think hard about the relationships between different parts of your application, identifying what layers need to observe or listen for behavior and which need to push notifications regarding behavior occurring to other parts of our apps.

Dynamic relationships may exist between publishers and subscribers when using this pattern. This provides a great deal of flexibility which may not be as easy to implement when disparate parts of your application are tightly coupled.

Whilst it may not always be the best solution to every problem, it remains one of the best tools for designing decoupled systems and should be considered an important tool in any JavaScript developer's utility belt.

Disadvantages

Consequently, some of the issues with this pattern actually stem from its main benefit. By decoupling publishers from subscribers, it can sometimes become difficult to obtain guarantees that particular parts of our applications are functioning as we may expect.

For example, publishers may make an assumption that one or more subscribers are listening to them. Say that we're using such an assumption to log or output errors regarding some application process. If the subscriber performing the logging crashes (or for some reason fails to function), the publisher won't have a way of seeing this due to the decoupled nature of the system.

Another draw-back of the pattern is that observers are quite ignorant to the existence of each other and are blind to the cost of switching in subject. Due to the dynamic relationship between subjects and observers the update dependency can

be difficult to track.

Implementations

Pub/Sub is a design pattern which fits in very well in JavaScript ecosystems, largely because at the core, ECMAScript implementations are event driven. This is particularly true in browser environments as the DOM uses events as it's main interaction API for scripting.

That said, neither ECMAScript nor DOM provide core objects or methods for creating custom events systems in implementation code (with the exception of perhaps the DOM3 CustomEvent, which is bound to the DOM and is thus not generically useful).

Luckily, popular JavaScript libraries such as dojo, jQuery (custom events) and YUI already have utilities that can assist in easily implementing a Pub/Sub system with very little effort.

```
// Publish

// Dojo: dojo.publish("channel", [arg1, arg2, arg3]);
dojo.publish("/login", [{username:"test", userData:"test"}]);

// jQuery: $(obj).trigger("channel", [arg1, arg2, arg3]);
$(el).trigger("/login", [{username:"test", userData:"test"}]);

// YUI: el.publish("channel", [arg1, arg2, arg3]);
el.publish("/login", {username:"test", userData:"test"});


// Subscribe

// Dojo: dojo.subscribe("channel", fn);
var handle = dojo.subscribe("/login", function(data){..});

// jQuery: $(obj).on("channel", [data], fn);
$(el).on("/login", function(event){...});

// YUI: el.on("channel", handler);
el.on("/login", function(data){...});


// Unsubscribe

// Dojo: dojo.unsubscribe(handle);
dojo.unsubscribe(handle);

// jQuery: $(obj).off("channel");
$(el).off("/login");
```

```
// YUI: el.detach("channel");  
el.detach("/login");
```

For those wishing to use the Pub/Sub pattern with vanilla JavaScript (or another library) AmplifyJS includes a clean, library-agnostic implementation of Pub/Sub that can be used with any library or toolkit. You can of course also write your own implementation from scratch or also check out either PubSubJS or OpenAjaxHub, both of which are also library-agnostic.

jQuery developers in particular have quite a few other options for Pub/Sub (in addition to Amplify) and can opt to use one of the many well-developed implementations ranging from Peter Higgins's jQuery plugin to Ben Alman's (optimized) gist on GitHub. Links to just a few of these can be found below.

- Ben Alman's Pub/Sub gist <https://gist.github.com/661855> (recommended)
- Rick Waldron's jQuery-core style take on the above <https://gist.github.com/705311>
- Peter Higgins' plugin <http://github.com/phiggins42/bloody-jquery-plugins/blob/master/pubsub.js>.
- AppendTo's Pub/Sub in AmplifyJS <http://amplifyjs.com>
- Ben Truyman's gist <https://gist.github.com/826794>

So that we are able to get an appreciation for how many of the vanilla JavaScript implementations of the Observer pattern might work, let's take a walk through of a minimalist version of Pub/Sub I released on GitHub under a project called **pubsubz**. This demonstrates the core concepts of subscribe, publish as well as the concept of unsubscribing.

I've opted to base our examples on this code as it sticks closely to both the method signatures and approach of implementation I would expect to see in a JavaScript version of the original observer pattern.

Sample Pub/Sub implementation

```
var pubsub = {};  
  
(function(q) {  
    var topics = {},  
        subUid = -1;
```

```

// Publish or broadcast events of interest
// with a specific topic name and arguments
// such as the data to pass along
q.publish = function( topic, args ) {

    if ( !topics[topic] ) {
        return false;
    }

    var subscribers = topics[topic],
        len = subscribers ? subscribers.length : 0;

    while (len--) {
        subscribers[len].func(topic, args);
    }

    return this;

};

// Subscribe to events of interest
// with a specific topic name and a
// callback function, to be executed
// when the topic/event is observed
q.subscribe = function( topic, func ) {

    if (!topics[topic]) {
        topics[topic] = [];
    }

    var token = (++subUid).toString();
    topics[topic].push({
        token: token,
        func: func
    });
    return token;
};

// Unsubscribe from a specific
// topic, based on a tokenized reference
// to the subscription
q.unsubscribe = function( token ) {
    for ( var m in topics ) {
        if ( topics[m] ) {
            for (var i = 0, j = topics[m].length; i < j; i++) {
                if (topics[m][i].token === token) {
                    topics[m].splice(i, 1);
                    return token;
                }
            }
        }
    }
    return this;
};

```

```
})( pubsub ));
```

Example 1: Basic use of publishers and subscribers

We can now use the implementation to publish and subscribe to events of interest as follows:

```
var testHandler = function (topics, data) {
    console.log(topics + ": " + data);
};

// Subscribers basically "subscribe" (or listen)
// And once they've been "notified" their callback functions are invoked
var testSubscription = pubsub.subscribe('example1', testHandler);

// Publishers are in charge of "publishing" notifications about events
pubsub.publish('example1', 'hello world!');
pubsub.publish('example1', ['test', 'a', 'b', 'c']);
pubsub.publish('example1', [{
    'color': 'blue'
}], {
    'text': 'hello'
});

// Unsubscribe if you no longer wish to be notified
pubsub.unsubscribe(testSubscription);

// This will fail
pubsub.publish('example1', 'hello again! (this will fail)');
```

A jsFiddle version of this example can be found at <http://jsfiddle.net/LxPrq/>

Real-time stock market application

Next, let's imagine we have a web application responsible for displaying real-time stock information.

The application might have a grid for displaying the stock stats and a counter for displaying the last point of update, as well as an underlying data model. When the data model changes, the application will need to update the grid and counter. In this scenario, our subject is the data model and the observers are the grid and counter.

When the observers receive notification that the model itself has changed, they can update themselves accordingly.

Example 2: UI notifications using pub/sub

In the following example, we limit our usage of pub/sub to that of a notification

system. Our subscriber is listening to the topic 'dataUpdated' to find out when new stock information is available. It then triggers 'gridUpdate' which goes on to call hypothetical methods that pull in the latest cached data object and re-render our UI components.

Note: the Mediator pattern is occasionally used to provide a level of communication between UI components without requiring that they communicate with each other directly. For example, rather than tightly coupling our applications, we can have widgets/components publish a topic when something interesting happens. A mediator can then subscribe to that topic and call the relevant methods on other components.

```
var grid = {  
  refreshData: function(){  
    console.log('retrieved latest data from data cache');  
    console.log('updated grid component');  
  },  
  updateCounter: function(){  
    console.log('data last updated at: ' + getCurrentTime());  
  }  
};  
  
// a very basic mediator  
  
var gridUpdate = function(topics, data){  
  grid.refreshData();  
  grid.updateCounter();  
}  
  
var dataSubscription = PubSub.subscribe( 'dataUpdated', gridUpdate );  
PubSub.publish( 'dataUpdated', 'new stock data available!' );  
PubSub.publish( 'dataUpdated', 'new stock data available!' );  
  
function getCurrentTime(){  
  var date = new Date(),  
      m = date.getMonth() + 1,  
      d = date.getDate(),  
      y = date.getFullYear(),  
      t = date.toLocaleTimeString().toLowerCase(),  
      return (m + '/' + d + '/' + y + ' ' + t);  
}
```

Whilst there's nothing terribly wrong with this, there are more optimal ways that we can utilize pub/sub to our advantage.

Example 3: Taking notifications further

Rather than just notifying our subscribers that new data is available, why not actually push the new data through to `gridUpdate` when we publish a new notification from a publisher. In this next example, our publisher will notify subscribers with the actual data that's been updated as well as a timestamp from the data-source of when the new data was added.

In addition to avoiding data having to be read from a cached store, this also avoids client-side calculation of the current time whenever a new data entry gets published.

```
var grid = {
  addEntry: function (data) {
    if (data !== 'undefined') {
      console.log('Entry:'
        + data.title
        + ' Changenet / %'
        + data.changenet
        + '/' + data.percentage + ' % added');
    }
  },
  updateCounter: function (timestamp) {
    console.log('grid last updated at: ' + timestamp);
  }
};

var gridUpdate = function (topics, data) {
  grid.addEntry(data);
  grid.updateCounter(data.timestamp);
}

var gridSubscription = PubSub.subscribe('dataUpdated', gridUpdate);

PubSub.publish('dataUpdated', {
  title: "Microsoft shares",
  changenet: 4,
  percentage: 33,
```

```

        timestamp: '17:34:12'
    });

    PubSub.publish('dataUpdated', {
        title: "Dell shares",
        changenet: 10,
        percentage: 20,
        timestamp: '17:35:16'
    });

```

Example 4: Decoupling applications using Ben Alman's pub/sub implementation

In the following movie ratings example, we'll be using Ben Alman's jQuery implementation of pub/sub to demonstrate how we can decouple a user interface. Notice how submitting a rating only has the effect of publishing the fact that new user and rating data is available.

It's left up to the subscribers to those topics to then delegate what happens with that data. In our case we're pushing that new data into existing arrays and then rendering them using the jQuery.tmpl plugin.

HTML/Templates

```

<script id="userTemplate" type="text/x-jquery-tmpl">
    <li>${user}</li>
</script>

<script id="ratingsTemplate" type="text/x-jquery-tmpl">
    <li><strong>${movie}</strong> was rated ${rating}/5</li>
</script>

<div id="container">

    <div class="sampleForm">
        <p>
            <label for="twitter_handle">Twitter handle:</label>
            <input type="text" id="twitter_handle" />
        </p>
        <p>
            <label for="movie_seen">Name a movie you've seen this year:</label>
            <input type="text" id="movie_seen" />
        </p>
        <p>
            <label for="movie_rating">Rate the movie you saw:</label>
            <select id="movie_rating">
                <option value="1">1</option>
                <option value="2">2</option>
                <option value="3">3</option>
            </select>
        </p>
    </div>

```



```

        <option value="4">4</option>
        <option value="5">5</option>

    </select>
</p>
<p>

    <button id="add">Submit rating</button>

</p>
</div>


<div class="summaryTable">
    <div id="users"><h3>Recent users</h3></div>
    <div id="ratings"><h3>Recent movies rated</h3></div>
</div>

</div>

```

JavaScript

```

(function($) {

    var movieList = [],
        userList = [];

    // subscribers

    $.subscribe( "/new/user", function( e, userName ){

        if(userName.length){
            userList.push({user: userName});
            $( "#userTemplate" ).tmpl( userList[userList.length - 1] ).appendTo(
"#users" );
        }

    });

    $.subscribe( "/new/rating", function( e, movieTitle, userRating ){

        if(movieTitle.length){
            movieList.push({ movie: movieTitle, rating: userRating});
            $( "#ratingsTemplate" ).tmpl( movieList[movieList.length - 1] ).appendTo(
"#ratings" );
        }

    });

    $('#add').on('click', function(){

```

```

    var strUser      = $("#twitter_handle").val(),
        strMovie     = $("#movie_seen").val(),
        strRating    = $("#movie_rating").val();

    // publishers

    $.publish('/new/user', strUser );
    $.publish('/new/rating', [ strMovie, strRating] );

    });

})(jQuery);

```

Example 5: Decoupling an Ajax-based jQuery application

In our final example, we're going to take a practical look at how decoupling our code using pub/sub early on in the development process can save us some potentially painful refactoring later on. This is something Rebecca Murphey touched on in her pub/sub screencast and is another reason why pub/sub is favored by so many developers in the community.

Quite often in Ajax-heavy applications, once we've received a response to a request we want to achieve more than just one unique action. One could simply add all of their post-request logic into a success callback, but there are drawbacks to this approach.

Highly coupled applications sometimes increase the effort required to reuse functionality due to the increased inter-function/code dependency. What this means is that although keeping our post-request logic hardcoded in a callback might be fine if we're just trying to grab a result set once, it's not as appropriate when we want to make further Ajax-calls to the same data source (and different end-behavior) without rewriting parts of the code multiple times. Rather than having to go back through each layer that calls the same data-source and generalizing them later on, we can use pub/sub from the start and save time.

Using pub/sub, we can also easily separate application-wide notifications regarding different events down to whatever level of granularity you're comfortable with, something which can be less elegantly done using other patterns.

Notice how in our sample below, one topic notification is made when a user indicates they want to make a search query and another is made when the request returns and actual data is available for consumption. It's left up to the subscribers to

then decide how to use knowledge of these events (or the data returned). The benefits of this are that, if we wanted, we could have 10 different subscribers utilizing the data returned in different ways but as far as the Ajax-layer is concerned, it doesn't care. Its sole duty is to request and return data then pass it on to whoever wants to use it. This separation of concerns can make the overall design of your code a little cleaner.

HTML/Templates:

```
<form id="flickrSearch">

  <input type="text" name="tag" id="query"/>

  <input type="submit" name="submit" value="submit"/>

</form>

<div id="lastQuery"></div>

<div id="searchResults"></div>

<script id="resultTemplate" type="text/x-jquery-tmpl">
  {{each(i, items) items}}
    <li><p></p></li>
  {{/each}}
</script>
```

JavaScript:

```
(function($) {

  $('#flickrSearch').submit(function( e ){

    e.preventDefault();
    var tags = $(this).find('#query').val();

    if(!tags){return;}
    $.publish('/search/tags', [ $.trim(tags) ]);

  });

  $.subscribe('/search/tags', function(tags){

    $.getJSON('http://api.flickr.com/services/feeds/photos_public.gne?
jsoncallback=?',
      { tags: tags, tagmode: 'any', format: 'json'},
```

```

        function(data){
            if(!data.items.length){ return; }
            $.publish('/search/resultSet', [ data ]);
        });

    });

$.subscribe('/search/tags', function(tags){
    $('#searchResults').html('Searched for:' + tags + '');
});

$.subscribe('/search/resultSet', function(results){

    var holder = $('#searchResults');
    holder.html();
    $('#resultTemplate').tmpl(results).appendTo(holder);

});

});

```

The Observer pattern is useful for decoupling a number of different scenarios in application design and if you haven't been using it, I recommend picking up one of the pre-written implementations mentioned today and just giving it a try out. It's one of the easier design patterns to get started with but also one of the most powerful.

The Mediator Pattern

The dictionary refers to a mediator as a neutral party who assists in negotiations and conflict resolution.

In software engineering, a mediator is a behavioral design pattern that allows us to expose a unified interface through which the different parts of a system may communicate. If it appears a system may have too many direct relationships between modules (colleagues), it may be time to have a central point of control that modules communicate through instead. The Mediator promotes loose coupling by ensuring that instead of modules referring to each other explicitly, their interaction is handled through this central point.

If you would prefer a real-world analogy, consider a typical airport traffic control system. A tower (mediator) handles what planes can take off and land because all communications (notifications being listened out for or broadcast) are done from the planes to the control tower, rather than from plane-to-plane. A centralized controller is key to the success of this system and that's really the role a mediator plays in software design.

In implementation terms, the mediator pattern is essentially a shared subject in the observer pattern. This might assume that a direction Publish/Subscribe relationship between objects or modules in such systems is sacrificed in order to maintain a central point of contact. It may also be considered supplemental - perhaps used for application-level notifications such as a communication between different subsystems that are themselves complex and may desire internal component decoupling through Publish/Subscribe relationships.

Another analogy would be DOM event bubbling and event delegation. If all subscriptions in a system are made against the document rather than individual nodes, the document effectively serves as a mediator. Instead of binding to the events of the individual nodes, a higher level object is given the responsibility of notifying subscribers about interaction events.

A possible implementation of the mediator pattern can be found below:

```
var mediator = (function(){  
    // Storage for our topics/events
```

```

var channels = {};

// Subscribe to an event, supply a callback to be executed
// when that event is broadcast
var subscribe = function(channel, fn){
    if (!channels[channel]) channels[channel] = [];
    channels[channel].push({ context: this, callback: fn });
    return this;
};

// Publish/broadcast an event to the rest of the application
var publish = function(channel){
    if (!channels[channel]) return false;
    var args = Array.prototype.slice.call(arguments, 1);
    for (var i = 0, l = channels[channel].length; i < l; i++) {
        var subscription = channels[channel][i];
        subscription.callback.apply(subscription.context, args);
    }
    return this;
};

return {
    publish: publish,
    subscribe: subscribe,
    installTo: function(obj){
        obj.subscribe = subscribe;
        obj.publish = publish;
    }
};

}());

```

Here is an example that uses the implementation from above. It's effectively centralized Publish/Subscribe where a mediated implementation of the Observer pattern is used:

```

(function( m ){

    // Set a default value for 'person'
    var person = "Luke";

    // Subscribe to a topic/event called 'nameChange' with
    // a callback function which will log the original
    // person's name and (if everything works) the incoming
    // name

    m.subscribe('nameChange', function( arg ){
        console.log( person ); // Luke
        person = arg;
        console.log( person ); // David
    });

    // Publish the 'nameChange' topic/event with the new data
    m.publish( 'nameChange', 'David' );

```

```
})( mediator );
```

Advantages & Disadvantages

The benefits of the Mediator pattern include that broadcasted events can be handled by any number of modules at once. It also can be used for a number of other purposes such as permissions management, given a Mediator can control what messages in a system can be subscribed to and which can be broadcast.

Perhaps the biggest downside of using the Mediator pattern is that it can introduce a single point of failure. Placing a Mediator between modules can also cause a performance hit as they are always communicating indirectly. Because of the nature of loose coupling, it's difficult to establish how a system might react by only looking at the broadcasts. That said, it's useful to remind ourselves that decoupled systems have a number of other benefits - if our modules communicated with each other directly, changes to modules (e.g another module throwing an exception) could easily have a domino effect on the rest of your application. This problem is less of a concern with decoupled systems.

At the end of the day, tight coupling causes all kinds of headaches and this is just another alternative solution, but one which can work very well if implemented correctly.

Mediator Vs. Observer

Developers often wonder what the differences are between the Mediator pattern and the Observer pattern. Admittedly, there is a bit of overlap, but let's refer back to the GoF for an explanation:

"In the Observer pattern, there is no single object that encapsulates a constraint. Instead, the Observer and the Subject must cooperate to maintain the constraint. Communication patterns are determined by the way observers and subjects are interconnected: a single subject usually has many observers, and sometimes the observer of one subject is a subject of another observer."

The Mediator pattern centralizes rather than simply just distributing. It places the responsibility for maintaining a constraint squarely in the mediator.

Mediator Vs. Facade

We will be covering the Facade pattern shortly, but for reference purposes some developers may also wonder whether there are similarities between the Mediator and Facade patterns. They do both abstract the functionality of existing modules, but there are some subtle differences.

The Mediator centralizes communication between modules where it's explicitly referenced by these modules. In a sense this is multidirectional. The Facade however just defines a simpler interface to a module or system but doesn't add any additional functionality. Other modules in the system aren't directly aware of the concept of a facade and could be considered unidirectional.

The Prototype Pattern

The GoF refer to the prototype pattern as one which creates objects based on a template of an existing object through cloning.

We can think of the prototype pattern as being based on prototypal inheritance where we create objects which act as prototypes for other objects. The prototype object itself is effectively used as a blueprint for each object the constructor creates. If the prototype of the constructor function used contains a property called 'name' for example (as per the code sample lower down), then each object created by that same constructor will also have this same property.

Looking at the definitions for the prototype pattern in existing literature non-specific to JavaScript, you **may** find references to concepts outside the scope of the language such as classes. The reality is that prototypal inheritance avoids using classes altogether. There isn't a 'definition' object nor a core object in theory. We're simply creating copies of existing functional objects.

One of the benefits of using the prototype pattern is that we're working with the strengths JavaScript has to offer natively rather than attempting to imitate features of other languages. With other design patterns, this isn't always the case. Not only is the pattern an easy way to implement inheritance, but it can also come with a performance boost as well: when defining a function in an object, they're all created by reference (so all child objects point to the same function) instead of creating their own individual copies.

For those interested, real prototypal inheritance, as defined in the ECMAScript 5 standard, requires the use of `Object.create` (which we've previously looked at briefly). `Object.create` creates an object which has a specified prototype and which optionally contains specified properties (i.e `Object.create(prototype, optionalDescriptorObjects)`). We can also see this being demonstrated in the example below:

```
// No need for capitalization as it's not a constructor
var someCar = {
  drive: function() {},
  name: 'Mazda 3'
};

// Use Object.create to generate a new car
```

```
var anotherCar = Object.create( someCar );
console.log(anotherCar.name); // Now you'll hopefully see that one is a prototype of
the other
```

`Object.create` allows you to easily implement advanced concepts such as differential inheritance where objects are able to directly inherit from other objects. With `Object.create` you're also able to initialise object properties using the second supplied argument. For example:

```
var vehicle = {
  getModel: function () {
    console.log('The model of this vehicle is..' + this.model);
  }
};

var car = Object.create(vehicle, {
  'id': {
    value: MY_GLOBAL.nextId(),
    enumerable: true // writable:false, configurable:false by default
  },
  'model': {
    value: 'Ford',
    enumerable: true
  }
});
```

Here the properties can be initialized on the second argument of `Object.create` using an object literal using the syntax similar to that used by the `Object.defineProperties` and `Object.defineProperty` methods. It allows you to set the property attributes such as `enumerable`, `writable` or `configurable`.

It is worth noting that prototypal relationships can cause trouble when enumerating properties of objects and (as Crockford recommends) wrapping the contents of the loop in a `hasOwnProperty()` check.

If you wish to implement the prototype pattern without directly using `Object.create`, you can simulate the pattern as per the above example as follows:

```
var vehiclePrototype = {
  init: function (carModel) {
    this.model = carModel;
  },
  getModel: function () {
    console.log('The model of this vehicle is..' + this.model);
  }
};

function vehicle(model) {
  function F() {};
```

```

    F.prototype = vehiclePrototype;

    var f = new F();

    f.init(model);
    return f;
}

var car = vehicle('Ford Escort');
car.getModel();

```

Note: This alternative does not allow the user to define read-only properties in the same manner (as the vehiclePrototype may be altered if not careful).

A final alternative implementation of the Prototype pattern is the following:

```

var beget = (function () {
    function F() {}

    return function (proto) {
        F.prototype = proto;
        return new F();
    };
})();

```

One could reference this method from the `vehicle` function. Note, however that `vehicle` here is emulating a constructor, since the prototype pattern does not include any notion of initialization beyond linking an object to a prototype.

The Command Pattern

The Command pattern aims to encapsulate method invocation, requests or operations into a single object and gives you the ability to both parameterize and pass method calls around that can be executed at your discretion. In addition, it enables you to decouple objects invoking the action from the objects which implement them, giving you a greater degree of overall flexibility in swapping out concrete 'classes'.

If you haven't come across concrete classes before, they are best explained in terms of class-based programming languages and are related to the idea of abstract classes. An abstract class defines an interface, but doesn't necessarily provide implementations for all of its member functions. It acts as a base class from which others are derived. A derived class which implements the missing functionality is called a concrete class (you may find these concepts familiar if you've read about the Decorator or Prototype patterns).

The main idea behind the command pattern is that it provides you a means to separate the responsibilities of issuing commands from anything executing commands, delegating this responsibility to different objects instead.

Implementation wise, simple command objects bind together both an action and the object wishing to invoke the action. They consistently include an execution operation (such as `run()` or `execute()`). All Command objects with the same interface can easily be swapped as needed and this is considered one of the larger benefits of the pattern.

To demonstrate the Command pattern we're going to create a simple car purchasing service.

```
(function(){  
    var CarManager = {  
        // request information  
        requestInfo: function( model, id ){  
            return 'The information for ' + model +  
                ' with ID ' + id + ' is foobar';  
        },  
        // purchase the car  
        buyVehicle: function( model, id ){
```

```

        return 'You have successfully purchased Item '
        + id + ', a ' + model;
    },

    // arrange a viewing
    arrangeViewing: function( model, id ){
        return 'You have successfully booked a viewing of '
        + model + ' ( ' + id + ' ) ';
    }
};

})();

```

Taking a look at the above code, it would be trivial to invoke our `CarManager` methods by directly accessing the object. We would all be forgiven for thinking there is nothing wrong with this - technically, it's completely valid JavaScript. There are however scenarios where this may be disadvantageous.

For example, imagine if the core API behind the `CarManager` changed. This would require all objects directly accessing these methods within our application to also be modified. This could be viewed as a layer of coupling which effectively goes against the OOP methodology of loosely coupling objects as much as possible. Instead, we could solve this problem by abstracting the API away further.

Let's now expand on our `CarManager` so that our application of the Command pattern results in the following: accept any named methods that can be performed on the `CarManager` object, passing along any data that might be used such as the Car model and ID.

Here is what we would like to be able to achieve:

```
CarManager.execute("buyVehicle", "Ford Escort", "453543");
```

As per this structure we should now add a definition for the "`CarManager.execute`" method as follows:

```
CarManager.execute = function (name) {
    return CarManager[name] && CarManager[name].apply(CarManager,
    [].slice.call(arguments, 1));
};
```

Our final sample calls would thus look as follows:

```
CarManager.execute("arrangeViewing", "Ferrari", "14523");
CarManager.execute("requestInfo", "Ford Mondeo", "54323");
CarManager.execute("requestInfo", "Ford Escort", "34232");
CarManager.execute("buyVehicle", "Ford Escort", "34232");
```

The Facade Pattern

When we put up a facade, we present an outward appearance to the world which may conceal a very different reality. This was the inspiration for the name behind the next pattern we're going to review - the facade pattern. The facade pattern provides a convenient higher-level interface to a larger body of code, hiding its true underlying complexity. Think of it as simplifying the API being presented to other developers, something which almost always improves usability.

Facades are a structural pattern which can often be seen in JavaScript libraries like jQuery where, although an implementation may support methods with a wide range of behaviors, only a 'facade' or limited abstraction of these methods is presented to the public for use.

This allows us to interact with the facade rather than the subsystem behind the scenes. Whenever you're using jQuery's `$(el).css()` or `$(el).animate()` methods, you're actually using a facade - the simpler public interface that avoids you having to manually call the many internal methods in jQuery core required to get some behavior working. This also avoids the need to manually interact with DOM APIs and maintain state variables.

The jQuery core methods should be considered intermediate abstractions. The more immediate burden to developers is the DOM API and facades are what make the jQuery library so easy to use.

To build on what we've learned, the facade pattern both simplifies the interface of a class and it also decouples the class from the code that utilizes it. This gives us the ability to indirectly interact with subsystems in a way that can sometimes be less prone to error than accessing the subsystem directly. A facade's advantages include ease of use and often a small size-footprint in implementing the pattern.

Let's take a look at the pattern in action. This is an unoptimized code example, but here we're utilizing a facade to simplify an interface for listening to events cross-browser. We do this by creating a common method that can be used in one's code which does the task of checking for the existence of features so that it can provide a safe and cross-browser compatible solution.

```
var addMyEvent = function( el, ev, fn ){  
    if(el.addEventListener){
```

```

        el.addEventListener( ev,fn, false );
    }else if(el.attachEvent){
        el.attachEvent( 'on'+ ev, fn );
    } else{
        el['on' + ev] = fn;
    }
};

```

In a similar manner, we're all familiar with jQuery's `$(document).ready(..)`. Internally, this is actually being powered by a method called `bindReady()`, which is doing this:

```

bindReady: function() {
    ...
    if ( document.addEventListener ) {
        // Use the handy event callback
        document.addEventListener( "DOMContentLoaded", DOMContentLoaded, false );

        // A fallback to window.onload, that will always work
        window.addEventListener( "load", jQuery.ready, false );

        // If IE event model is used
    } else if ( document.attachEvent ) {

        document.attachEvent( "onreadystatechange", DOMContentLoaded );

        // A fallback to window.onload, that will always work
        window.attachEvent( "onload", jQuery.ready );
        ...
    }
}

```

This is another example of a facade, where the rest of the world simply uses the limited interface exposed by `$(document).ready(..)` and the more complex implementation powering it is kept hidden from sight.

Facades don't just have to be used on their own, however. They can also be integrated with other patterns such as the module pattern. As you can see below, our instance of the module patterns contains a number of methods which have been privately defined. A facade is then used to supply a much simpler API to accessing these methods:

```

var module = (function() {
    var _private = {
        i:5,
        get : function() {
            console.log('current value:' + this.i);
        },
        set : function( val ) {
            this.i = val;
        },
        run : function() {
            console.log( 'running' );
        }
    };
}());

```

```

    },
    jump: function(){
        console.log( 'jumping' );
    }
};
return {
    facade : function( args ) {
        _private.set(args.val);
        _private.get();
        if ( args.run ) {
            _private.run();
        }
    }
}
})();

```

```

module.facade({run: true, val:10});
//outputs current value: 10, running

```

In this example, calling `module.facade()` will actually trigger a set of private behavior within the module, but again, the user isn't concerned with this. We've made it much easier for them to consume a feature without needing to worry about implementation-level details.

The Factory Pattern

Similar to other creational patterns, the Factory Pattern deals with the problem of creating objects (which we can think of as ‘factory products’) without the need to specify the exact class of object being created.

Specifically, the Factory Pattern suggests defining an interface for creating an object where you allow the subclasses to decide which class to instantiate. This pattern handles the problem by defining a completely separate method for the creation of objects and which sub-classes are able to override so they can specify the ‘type’ of factory product that will be created.

This is quite useful, in particular if the creation process involved is quite complex, e.g. if it strongly depends on the settings in configuration files.

You can often find factory methods in frameworks where the code for a library may need to create objects of particular types which may be subclassed by scripts using the frameworks.

In our example, let’s take the code used in the original Constructor pattern example and see what this would look like were we to optimize it using the Factory Pattern:

```
function VehicleFactory() {}
VehicleFactory.prototype.vehicleClass = Car;
VehicleFactory.prototype.getVehicle = function (options) {
    return new this.vehicleClass(options);
};

var carFactory = new VehicleFactory();
var car = carFactory.getVehicle({ color: "yellow", turbo: true });
console.log(car instanceof Car); // => true

// approach #1: Modify a VehicleFactory instance to use the Truck class
carFactory.vehicleClass = Truck;

var mover = carFactory.getVehicle({ enclosedCargo: true, length: 26 });
console.log(mover instanceof Truck); // => true

// approach #2: Subclass VehicleFactory to create a factory class that
// builds Trucks
function TruckFactory () {}
TruckFactory.prototype = new VehicleFactory();
TruckFactory.prototype.vehicleClass = Truck;

var truckFactory = new TruckFactory();
var bigfoot = truckFactory.getVehicle({ monster: true, cylinders: 12 });
```

```
console.log(bigfoot instanceof Truck); // => true
```

When To Use The Factory Pattern

The Factory pattern can be especially useful when applied to the following situations:

- When your object's setup requires a high level of complexity
- When you need to generate different instances depending on the environment
- When you're working with many small objects that share the same properties
- When composing classes with instances of other classes that need only satisfy an API contract (aka, duck typing) to work. This is useful for decoupling.

When Not To Use The Factory Pattern

It's generally a good practice to not use the factory pattern in every situation as it can easily add an unnecessarily additional aspect of complexity to your code. It can also make some tests more difficult to run.

It is also useful to be aware of the **Abstract Factory** pattern, which aims to encapsulate a group of individual factories with a common goal. It separates the details of implementation of a set of objects from their general usage.

An Abstract Factory should be used where a system must be independent from the way the objects it creates are generated or it needs to work with multiple types of objects. An example which is both simple and easier to understand is a vehicle factory, which defines ways to get or register vehicles types. The abstract factory can be named `AbstractVehicleFactory`. The abstract factory will allow the definition of types of vehicle like 'car' or 'truck' and concrete factories will implement only classes that fulfill the vehicle contract (e.g `Vehicle.prototype.drive` and `Vehicle.prototype.breakDown`).

```
var AbstractVehicleFactory = (function () {  
  var types = {};  
  
  return {  
    getVehicle: function (type, customizations) {  
      var Vehicle = types[type];  
  
      return (Vehicle) ? return new Vehicle(customizations) : null;  
    },  
  },  
}
```

```

    registerVehicle: function (type, Vehicle) {
        var proto = Vehicle.prototype;

        // only register classes that fulfill the vehicle contract
        if (proto.drive && proto.breakDown) {
            types[type] = Vehicle;
        }

        return AbstractVehicleFactory;
    }
}());

//Usage

AbstractVehicleFactory.registerVehicle("car", Car);
AbstractVehicleFactory.registerVehicle("truck", Truck);

var car = AbstractVehicleFactory.getVehicle("car", { color: "yellow", turbo: true
});
var truck = AbstractVehicleFactory.getVehicle("truck", { monster: true, cylinders:
12 });

```

The Mixin Pattern

In traditional object-oriented programming languages, **mixins** are classes which provide the functionality to be inherited by a subclass. Inheriting from mixins are a means of collecting functionality and classes may inherit functionality from multiple mixins through multiple inheritance.

In the following example, we have a Car defined without any methods. We also have a constructor called 'Mixin'. What we're going to do is augment the Car so it has access to the methods within the Mixin. This code demonstrates how with JavaScript you can augment a constructor to have a particular method without using the typical inheritance methods or duplicating code for each constructor function you have.

```
// Car
var Car = function( settings ){
  this.model = settings.model || 'no model provided';
  this.colour = settings.colour || 'no colour provided';
};

// Mixin
var Mixin = function(){};
Mixin.prototype = {
  driveForward: function(){
    console.log('drive forward');
  },
  driveBackward: function(){
    console.log('drive backward');
  }
};

// Augment existing 'class' with a method from another
function augment( receivingClass, givingClass ) {
  // only provide certain methods
  if ( arguments[2] ) {
    for (var i=2, len=arguments.length; i<len; i++) {
      receivingClass.prototype[arguments[i]] = givingClass.prototype[arguments[i]];
    }
  }
  // provide all methods
  else {
    for ( var methodName in givingClass.prototype ) {
      /* check to make sure the receiving class doesn't
       have a method of the same name as the one currently
       being processed */
      if ( !receivingClass.prototype[methodName] ) {
        receivingClass.prototype[methodName] = givingClass.prototype[methodName];
      }
    }
  }
}
```

```
}  
}  
}
```

```
// Augment the Car have the methods 'driveForward' and 'driveBackward'*/  
augment( Car, Mixin, 'driveForward', 'driveBackward' );  
  
// Create a new Car  
var vehicle = new Car({model: 'Ford Escort', colour: 'blue'});  
  
// Test to make sure we now have access to the methods  
vehicle.driveForward();  
vehicle.driveBackward();
```

It's worth noting that handling mixins that have constructor code is possible, but can be more complicated. Although the mixed in class prototype members are mirrored on the receiving class prototype, instances of the receiving class do not pass instanceof for the mixin class; JS does not support multiple inheritance.

The Decorator Pattern

In this section we're going to explore the **decorator** - a structural design pattern that promotes code reuse and is a flexible alternative to subclassing. This pattern is also useful for modifying existing systems where you may wish to add additional features to objects without the need to change the underlying code that uses them.

Traditionally, the decorator is defined as a design pattern that allows behavior to be added to an existing object dynamically. The idea is that the decoration itself isn't essential to the base functionality of an object otherwise it would be baked into the 'superclass' object itself.

Subclassing

For developers unfamiliar with subclassing, here is a beginner's primer on them before we dive further into decorators: subclassing is a term that refers to inheriting properties for a new object from a base or 'superclass' object.

In traditional OOP, a class B is able to extend another class A. Here we consider A a superclass and B a subclass of A. As such, all instances of B inherit the methods from A. B is however still able to define its own methods, including those that override methods originally defined by A.

Should B need to invoke a method in A that has been overridden, we refer to this as method chaining. Should B need to invoke the constructor A() (the superclass), we call this constructor chaining.

In order to demonstrate subclassing, we first need a base object that can have new instances of itself created. Let's model this around the concept of a person.

```
var Person = function( firstName , lastName ){  
    this.firstName = firstName;  
    this.lastName = lastName;  
    this.gender = 'male'  
};
```

Next, we'll want to specify a new class (object) that's a subclass of the existing Person object. Let's imagine we want to add distinct properties to distinguish a Person from a Superhero whilst inheriting the properties of the Person 'superclass'. As superheroes share many common traits with normal people (e.g. name, gender), this should hopefully illustrate how subclassing works adequately.

```
// a new instance of Person can then easily be created as follows:  
var clark = new Person( "Clark" , "Kent" );  
  
// Define a subclass constructor for for 'Superhero':  
var Superhero = function( firstName, lastName , powers ){  
  
    // Invoke the superclass constructor on the new object  
    // then use .call() to invoke the constructor as a method of  
    // the object to be initialized.  
  
    Person.call(this, firstName, lastName);  
  
    // Finally, store their powers, a new array of traits not found in a normal  
    'Person'  
    this.powers = powers;  
}
```

```
SuperHero.prototype = Object.create(Person.prototype);  
var superman = new Superhero( "Clark" ,"Kent" , ['flight','heat-vision'] );  
console.log(superman); // includes superhero props as well as gender
```

The Superhero definition creates an object which descends from Person. Objects of this type have properties of the objects that are above it in the chain and if we had set default values in the Person object, Superhero is capable of overriding any inherited values with values specific to it's object.

So where do decorators come in?

Decorators

Decorators are used when it's necessary to delegate responsibilities to an object where it doesn't make sense to subclass it. A common reason for this is that the number of features required demand for a very large quantity of subclasses. Can you imagine having to define hundreds or thousands of subclasses for a project? It would likely become unmanageable fairly quickly.

To give you a visual example of where this is an issue, imagine needing to define new kinds of **Superhero**: **SuperheroThatCanFly**, **SuperheroThatCanRunQuickly** and **SuperheroWithXRayVision**.

Now, what if superhero had more than one of these properties?. We'd need to define a subclass called **SuperheroThatCanFlyAndRunQuickly** , **SuperheroThatCanFlyRunQuicklyAndHasXRayVision** etc - effectively, one for each possible combination. As you can see, this isn't very manageable when you factor in different abilities.

The decorator pattern isn't heavily tied to how objects are created but instead focuses on the problem of extending their functionality. Rather than just using inheritance, where we're used to extending objects linearly, we work with a single base object and progressively add decorator objects which provide the additional capabilities. The idea is that rather than subclassing, we add (decorate) properties or methods to a base object so its a little more streamlined.

The extension of objects is something already built into JavaScript and as we know, objects can be extended rather easily with properties being included at any point. With this in mind, a very very simplistic decorator may be implemented as follows:

Example 1: Basic decoration of existing object constructors with new functionality

```
function vehicle( vehicleType ){  
  // properties and defaults  
  this.vehicleType = vehicleType || 'car',  
  this.model = 'default',  
  this.license = '00000-000'  
}
```

```

// Test instance for a basic vehicle
var testInstance = new vehicle('car');
console.log(testInstance);

// vehicle: car, model:default, license: 00000-000

// Lets create a new instance of vehicle, to be decorated*/
var truck = new vehicle('truck');

// New functionality we're decorating vehicle with
truck.setModel = function( modelName ){
    this.model = modelName;
}

truck.setColor = function( color ){
    this.color = color;
}

// Test the value setters and value assignment works correctly
truck.setModel('CAT');
truck.setColor('blue');
console.log(truck);
// vehicle:truck, model:CAT, color: blue

// Demonstrate 'vehicle' is still unaltered
var secondInstance = new vehicle('car');
console.log(secondInstance);

// as before, vehicle: car, model:default, license: 00000-000

```

This type of simplistic implementation is something you're likely familiar with, but it doesn't really demonstrate some of the other strengths of the pattern. For this, we're first going to go through my variation of the Coffee example from an excellent book called Head First Design Patterns by Freeman, Sierra and Bates, which is modelled around a Macbook purchase.

We're then going to look at psuedo-classical decorators.

Example 2: Simply decorate objects with multiple decorators

```

// What we're going to decorate
function MacBook() {
    this.cost = function () { return 997; };
    this.screenSize = function () { return 13.3; };
}

// Decorator 1
function Memory( macbook ) {
    var v = macbook.cost();
    macbook.cost = function() {

```

```

        return v + 75;
    }
}
// Decorator 2
function Engraving( macbook ){
    var v = macbook.cost();
    macbook.cost = function(){
        return v + 200;
    };
}

// Decorator 3
function Insurance( macbook ){
    var v = macbook.cost();
    macbook.cost = function(){
        return v + 250;
    };
}
var mb = new MacBook();
Memory(mb);
Engraving(mb);
Insurance(mb);
console.log(mb.cost()); //1522
console.log(mb.screenSize()); //13.3

```

Here, the decorators are overriding the superclass `.cost()` method to return the current price of the Macbook plus with the cost of the upgrade being specified. It's considered a decoration as the original Macbook object's constructor methods which are not overridden (e.g. `screenSize()`) as well as any other properties which we may define as a part of the Macbook remain unchanged and intact.

As you can probably tell, there isn't really a defined 'interface' in the above example and we're shifting away the responsibility of ensuring an object meets an interface when moving from the creator to the receiver.

Pseudo-classical decorators

We're now going to examine the variation of the decorator presented in 'Pro JavaScript Design Patterns' (PJDP) by Dustin Diaz and Ross Harmes.

Unlike some of the examples from earlier, Diaz and Harmes stick more closely to how decorators are implemented in other programming languages (such as Java or C++) using the concept of an 'interface', which we'll define in more detail shortly.

Note: This particular variation of the decorator pattern is provided for reference purposes. If you find it overly complex for your application's needs, I recommend sticking to one the simpler implementations covered earlier, but I would still read the section. If you haven't yet grasped how decorators are different from subclassing, it may help!.

Interfaces

PJDP describes the decorator as a pattern that is used to transparently wrap objects inside other objects of the same interface. An interface is a way of defining the methods an object **should** have, however, it doesn't actually directly specify how those methods should be implemented.

They can also indicate what parameters the methods take, but this is considered optional.

So, why would you use an interface in JavaScript? The idea is that they're self-documenting and promote reusability. In theory, interfaces also make code more stable by ensuring changes to them must also be made to the classes implementing them.

Below is an example of an implementation of Interfaces in JavaScript using duck-typing - an approach that helps determine whether an object is an instance of constructor/object based on the methods it implements.

```
var TodoList = new Interface('Composite', ['add', 'remove']);
var TodoItem = new Interface('TodoItem', ['save']);
// TodoList class
var myTodoList = function(id, method, action) {
    // implements TodoList, TodoItem
    ...
};
...
```

```

function addTodo( todoInstance ) {
    Interface.ensureImplements(todoInstance, TodoList, TodoItem);
    // This function will throw an error if a required method is not
    implemented,
    // halting execution of the function.
    //...
}

```

where `Interface.ensureImplements` provides strict checking. If you would like to explore interfaces further, I recommend looking at Chapter 2 of Pro JavaScript design patterns. For the Interface class used above, see [here](#).

The biggest problem with interfaces is that, as there isn't built-in support for them in JavaScript, there's a danger of us attempting to emulate the functionality of another language, however, we're going to continue demonstrating their use just to give you a complete view of how the decorator is implemented by other developers.

This variation of decorators and abstract decorators

To demonstrate the structure of this version of the decorator pattern, we're going to imagine we have a superclass that models a macbook once again and a store that allows you to 'decorate' your macbook with a number of enhancements for an additional fee.

Enhancements can include upgrades to 4GB or 8GB Ram, engraving, Parallels or a case. Now if we were to model this using an individual subclass for each combination of enhancement options, it might look something like this:

```

var Macbook = function(){
    //...
}
var MacbookWith4GBRam = function() {},
    MacbookWith8GBRam = function() {},
    MacbookWith4GBRamAndEngraving = function() {},
    MacbookWith8GBRamAndEngraving = function() {},
    MacbookWith4GBRamAndParallels = function() {},
    MacbookWith8GBRamAndParallels = function() {},
    MacbookWith4GBRamAndParallelsAndCase = function() {},
    MacbookWith8GBRamAndParallelsAndCase = function() {},
    MacbookWith4GBRamAndParallelsAndCaseAndInsurance = function() {},
    MacbookWith8GBRamAndParallelsAndCaseAndInsurance = function() {};

```

and so on.

This would be an impractical solution as a new subclass would be required for every possible combination of enhancements that are available. As we'd prefer to keep things simple without maintaining a large set of subclasses, let's look at how decorators may be used to solve this problem better.

Rather than requiring all of the combinations we saw earlier, we should simply have to create five new decorator classes. Methods that are called on these enhancement classes would be passed on to our Macbook class.

In our next example, decorators transparently wrap around their components and can interestingly be interchanged as they use the same interface.

Here's the interface we're going to define for the Macbook:

```
var Macbook = new Interface('Macbook', ['addEngraving', 'addParallels', 'add4GBRam', 'add8GBRam', 'addCase']);
// A Macbook Pro might thus be represented as follows:
var MacbookPro = function(){
    // implements Macbook
}
MacbookPro.prototype = {
    addEngraving: function(){
    },
    addParallels: function(){
    },
    add4GBRam: function(){
    },
    add8GBRam: function(){
    },
    addCase: function(){
    },
    getPrice: function(){
        return 900.00; //base price.
    }
};
```

We're not going to worry about the actual implementation at this point as we'll shortly be passing on all method calls that are made on them.

To make it easier for us to add as many more options as needed later on, an abstract decorator class is defined with default methods required to implement the Macbook interface, which the rest of the options will subclass.

Abstract decorators ensure that we can decorate a base class independently with as many decorators as needed in different combinations (remember the example earlier?) without needing to derive a class for every possible combination.

```
//Macbook decorator abstract decorator class
```

```

var MacbookDecorator = function( macbook ){
  Interface.ensureImplements(macbook, Macbook);
  this.macbook = macbook;
}
MacbookDecorator.prototype = {
  addEngraving: function(){
    return this.macbook.addEngraving();
  },
  addParallels: function(){
    return this.macbook.addParallels();
  },
  add4GBRam: function(){
    return this.macbook.add4GBRam();
  },
  add8GBRam: function(){
    return this.macbook.add8GBRam();
  },
  addCase: function(){
    return this.macbook.addCase();
  },
  getPrice: function(){
    return this.macbook.getPrice();
  }
};

```

What's happening in the above sample is that the Macbook decorator is taking an object to use as the component. It's using the Macbook interface we defined earlier and for each method is just calling the same method on the component. We can now create our option classes just by using the Macbook decorator - simply call the superclass constructor and any methods can be overridden as per necessary.

```

var CaseDecorator = function( macbook ){
  /*call the superclass's constructor next*/
  this.superclass.constructor(macbook);
}

// Let's now extend the superclass
extend( CaseDecorator, MacbookDecorator );

CaseDecorator.prototype.addCase = function(){
  return this.macbook.addCase() + " Adding case to macbook ";
};

CaseDecorator.prototype.getPrice = function(){
  return this.macbook.getPrice() + 45.00;
};

```

As you can see, most of this is relatively easy to implement. What we're doing is overriding the addCase() and getPrice() methods that need to be decorated and we're achieving this by first executing the component's method and then adding to it.

As there's been quite a lot of information presented in this section so far, let's try to bring it all together in a single example that will hopefully highlight what we've learned.

```
// Instantiation of the macbook
var myMacbookPro = new MacbookPro();

// This will return 900.00
console.log(myMacbookPro.getPrice());

// Decorate the macbook
myMacbookPro = new CaseDecorator( myMacbookPro ); /*note*/

// This will return 945.00
console.log(myMacbookPro.getPrice());
```

An important note from PJDP is that in the line denoted **note**, Harmes and Diaz claim that it's important not to create a separate variable to store the instance of your decorators, opting for the same variable instead. The downside to this is that we're unable to access the original macbook object in our example, however we technically shouldn't need to further.

As decorators are able to modify objects dynamically, they're a perfect pattern for changing existing systems. Occasionally, it's just simpler to create decorators around an object versus the trouble of maintaining individual subclasses. This makes maintaining applications of this type significantly more straight-forward.

Implementing decorators with jQuery

As with other patterns I've covered, there are also examples of the decorator pattern that can be implemented with jQuery. `jQuery.extend()` allows you to extend (or merge) two or more objects (and their properties) together into a single object either at run-time or dynamically at a later point.

In this scenario, a target object can be decorated with new functionality without necessarily breaking or overriding existing methods in the source/superclass object (although this can be done).

In the following example, we define three objects: defaults, options and settings. The aim of the task is to decorate the 'defaults' object with additional functionality found in 'options', which we'll make available through 'settings'. We must:

- (a) Leave 'defaults' in an untouched state where we don't lose the ability to access the properties or functions found in it a later point
- (b) Gain the ability to use the decorated properties and functions found in 'options'

```
var decoratorApp = decoratorApp || {};  
// define the objects we're going to use  
decoratorApp = {  
  defaults:{  
    validate: false,  
    limit: 5,  
    name: "foo",  
    welcome: function(){  
      //console.log('welcome!');  
    }  
  },  
  options:{  
    validate: true,  
    name: "bar",  
    helloWorld: function(){  
      //console.log('hello');  
    }  
  },  
  settings:{},  
  printObj: function(obj) {  
    var arr = [];  
    $.each(obj, function(key, val) {  
      var next = key + ": ";  
      next += $.isPlainObject(val) ? printObj(val) : val;  
      arr.push( next );  
    });  
    return "{ " + arr.join(", ") + " }";  
  }  
}
```

```

}
/* merge defaults and options, without modifying defaults */
decoratorApp.settings = $.extend({}, decoratorApp.defaults, decoratorApp.options);
/* what we've done here is decorated defaults in a way that provides access to the
properties and functionality it has to offer (as well as that of the decorator
'options'). defaults itself is left unchanged*/
$('#log').append("<div><b>settings -- </b>" +
decoratorApp.printObj(decoratorApp.settings) + "</div><div><b>options -- </b>" +
decoratorApp.printObj(decoratorApp.options) + "</div><div><b>defaults -- </b>"
+decoratorApp.printObj(decoratorApp.defaults) + "</div>" );
/*
settings -- { validate: true, limit: 5, name: bar, welcome: function (){
console.log('welcome!'); }, helloWorld: function (){ console.log('hello!'); } }
options -- { validate: true, name: bar, helloWorld: function (){
console.log('hello!'); } }
defaults -- { validate: false, limit: 5, name: foo, welcome: function (){
console.log('welcome!'); } }
*/

```

Pros and cons of the pattern

Developers enjoy using this pattern as it can be used transparently and is also fairly flexible - as we've seen, objects can be wrapped or 'decorated' with new behavior and then continue to be used without needing to worry about the base object being modified. In a broader context, this pattern also avoids us needing to rely on large numbers of subclasses to get the same benefits.

There are however drawbacks that you should be aware of when implementing the pattern. If poorly managed, it can significantly complicate your application's architecture as it introduces many small, but similar objects into your namespace. The concern here is that in addition to becoming hard to manage, other developers unfamiliar with the pattern may have a hard time grasping why it's being used.

Sufficient commenting or pattern research should assist with the latter, however as long as you keep a handle on how widespread you use the decorator in your application you should be fine on both counts.

Chapter 10. Flyweight

The Flyweight pattern is considered a useful classical solution for code that's repetitive, slow and inefficient - for example: situations where we might create a large number of similar objects.

It's of particular use in JavaScript where code that's complex in nature may easily use all of the available memory, causing a number of performance issues.

Interestingly, it's been quite underused in recent years. Given how reliant we are on JavaScript for the applications of today, both performance and scalability are often paramount and this pattern (when applied correctly) can assist with improving both.

To give you some quick historical context, the pattern is named after the boxing weight class that includes fighters weighing less than 112lb - Poncho Villa being the most famous fighter in this division. It derives from this weight classification as it refers to the small amount of weight (memory) used.

Flyweights are an approach to taking several similar objects and placing that shared information into a single external object or structure. The general idea is that (in theory) this reduces the resources required to run an overall application. The flyweight is also a structural pattern, meaning that it aims to assist with both the structure of your objects and the relationships between them.

So, how do we apply it to JavaScript?

There are two ways in which the Flyweight pattern can be applied. The first is on the data-layer, where we deal with the concept of large quantities of similar objects stored in memory. The second is on the DOM-layer where the flyweight can be used as a central event-manager to avoid attaching event handlers to every child element in a parent container you wish to have some similar behavior.

As the data-layer is where the flyweight pattern is most used traditionally, we'll take a look at this first.

Flyweight and the data layer

For this application, there are a few more concepts around the classical flyweight

pattern that we need to be aware of. In the Flyweight pattern there's a concept of two states - intrinsic and extrinsic. Intrinsic information may be required by internal methods in your objects which they absolutely can't function without. Extrinsic information can however be removed and stored externally.

Objects with the same intrinsic data can be replaced with a single shared object, created by a factory method, meaning we're able to reduce the overall quantity of objects down significantly. The benefit of this is that we're able to keep an eye on objects that have already been instantiated so that new copies are only ever created should the intrinsic state differ from the object we already have.

We use a manager to handle the extrinsic states. How this is implemented can vary, however as Dustin Diaz correctly points out in Pro JavaScript Design patterns, one approach to this is to have the manager object contain a central database of the extrinsic states and the flyweight objects which they belong to.

Converting code to use the Flyweight pattern

Let's now demonstrate some of these concepts using the idea of a system to manage all of the books in a library. The important meta-data for each book could probably be broken down as follows:

- ID
- Title
- Author
- Genre
- Page count
- Publisher ID
- ISBN

We'll also require the following properties to keep track of which member has checked out a particular book, the date they've checked it out on as well as the expected date of return.

- checkoutDate
- checkoutMember
- dueReturnDate
- availability

Each book would thus be represented as follows, prior to any optimization:

```
var Book = function( id, title, author, genre, pageCount,publisherID, ISBN,
checkoutDate, checkoutMember, dueReturnDate,availability ){
    this.id = id;
    this.title = title;
    this.author = author;
    this.genre = genre;
    this.pageCount = pageCount;
    this.publisherID = publisherID;
    this.ISBN = ISBN;
    this.checkoutDate = checkoutDate;
    this.checkoutMember = checkoutMember;
    this.dueReturnDate = dueReturnDate;
    this.availability = availability;
};
Book.prototype = {
    getTitle:function(){
```

```

        return this.title;
    },
    getAuthor: function(){
        return this.author;
    },
    getISBN: function(){
        return this.ISBN;
    },
    // other getters not shown for brevity
    updateCheckoutStatus: function(bookID, newStatus, checkoutDate, checkoutMember,
    newReturnDate){
        this.id = bookID;
        this.availability = newStatus;
        this.checkoutDate = checkoutDate;
        this.checkoutMember = checkoutMember;
        this.dueReturnDate = newReturnDate;
    },
    extendCheckoutPeriod: function(bookID, newReturnDate){
        this.id = bookID;
        this.dueReturnDate = newReturnDate;
    },
    isPastDue: function(bookID){
        var currentDate = new Date();
        return currentDate.getTime() > Date.parse(this.dueReturnDate);
    }
};

```

This probably works fine initially for small collections of books, however as the library expands to include a larger inventory with multiple versions and copies of each book available, you'll find the management system running slower and slower over time. Using thousands of book objects may overwhelm the available memory, but we can optimize our system using the flyweight pattern to improve this.

We can now separate our data into intrinsic and extrinsic states as follows: data relevant to the book object (title, author etc) is intrinsic whilst the checkout data (checkoutMember, dueReturnDate etc) is considered extrinsic. Effectively this means that only one Book object is required for each combination of book properties. It's still a considerable quantity of objects, but significantly fewer than we had previously.

The following single instance of our book meta-data combinations will be shared among all of the copies of a book with a particular title.

```

// flyweight optimized version
var Book = function (title, author, genre, pageCount, publisherID, ISBN) {
    this.title = title;
    this.author = author;
    this.genre = genre;
    this.pageCount = pageCount;

```

```
    this.publisherID = publisherID;  
    this.ISBN = ISBN;  
};
```

As you can see, the extrinsic states have been removed. Everything to do with library check-outs will be moved to a manager and as the object's data is now segmented, a factory can be used for instantiation.

A Basic Factory

Let's now define a very basic factory. What we're going to have it do is perform a check to see if a book with a particular title has been previously created inside the system. If it has, we'll return it. If not, a new book will be created and stored so that it can be accessed later. This makes sure that we only create a single copy of each unique intrinsic piece of data:

```
// Book Factory singleton
var BookFactory = (function () {
    var existingBooks = {};
    return {
        createBook: function (title, author, genre, pageCount, publisherID, ISBN) {
            // Find out if a particular book meta-data combination has been created before
            var existingBook = existingBooks[ISBN];
            if (existingBook) {
                return existingBook;
            } else {
                // if not, let's create a new instance of it and store it
                var book = new Book(title, author, genre, pageCount, publisherID, ISBN);
                existingBooks[ISBN] = book;
                return book;
            }
        }
    }
})();
```

Managing the extrinsic states

Next, we need to store the states that were removed from the Book objects somewhere - luckily a manager (which we'll be defining as a singleton) can be used to encapsulate them. Combinations of a Book object and the library member that's checked them out will be called Book records. Our manager will be storing both and will also include checkout related logic we stripped out during our flyweight optimization of the Book class.

```
// BookRecordManager singleton
var BookRecordManager = (function () {
  var bookRecordDatabase = {};
  return {
    // add a new book into the library system
    addBookRecord: function (id, title, author, genre, pageCount, publisherID, ISBN,
      checkoutDate, checkoutMember, dueReturnDate, availability) {
      var book = bookFactory.createBook(title, author, genre, pageCount,
        publisherID, ISBN);
      bookRecordDatabase[id] = {
        checkoutMember: checkoutMember,
        checkoutDate: checkoutDate,
        dueReturnDate: dueReturnDate,
        availability: availability,
        book: book;
      };
    },
    updateCheckoutStatus: function (bookID, newStatus, checkoutDate, checkoutMember,
      newReturnDate) {
      var record = bookRecordDatabase[bookID];
      record.availability = newStatus;
      record.checkoutDate = checkoutDate;
      record.checkoutMember = checkoutMember;
      record.dueReturnDate = newReturnDate;
    },
    extendCheckoutPeriod: function (bookID, newReturnDate) {
      bookRecordDatabase[bookID].dueReturnDate = newReturnDate;
    },
    isPastDue: function (bookID) {
      var currentDate = new Date();
      return currentDate.getTime() >
        Date.parse(bookRecordDatabase[bookID].dueReturnDate);
    }
  };
})();
```

The result of these changes is that all of the data that's been extracted from the Book 'class' is now being stored in an attribute of the BookManager singleton (BookDatabase) which is considerable more efficient than the large number of

objects we were previously using. Methods related to book checkouts are also now based here as they deal with data that's extrinsic rather than intrinsic.

This process does add a little complexity to our final solution, however it's a small concern when compared to the performance issues that have been tackled.

Data wise, if we have 30 copies of the same book, we are now only storing it once. Also, every function takes up memory. With the flyweight pattern these functions exist in one place (on the manager) and not on every object, thus saving more memory.

The Flyweight pattern and the DOM

In JavaScript, functions are effectively object descriptors and all functions are also JavaScript objects internally. The goal of the pattern here is thus to make triggering objects have little to no responsibility for the actions they perform and to instead abstract this responsibility up to a global manager. One of the best metaphors for describing the pattern was written by Gary Chisholm and it goes a little like this:

Try to think of the flyweight in terms of a pond. A fish opens its mouth (the event), bubbles raise to the surface (the bubbling) a fly sitting on the top flies away when the bubble reaches the surface (the action). In this example you can easily transpose the fish opening its mouth to a button being clicked, the bubbles as the bubbling effect and the fly flying away to some function being run'.

As jQuery is accepted as one of the best options for DOM-manipulation and selection, we'll be using it for our DOM-related examples.

Example 1: Centralized event handling

For our first practical example, consider scenarios where you may have a number of similar elements or structures on a page that share similar behavior when a user-action is performed against them.

In JavaScript, there's a known bubbling effect in the language so that if an element such as a link or button is clicked, that event is bubbled up to the parent, informing them that something lower down the tree has been clicked. We can use this effect to our advantage.

Normally what you might do when constructing your own accordion component, menu or other list-based widget is bind a click event to each link element in the parent container. Instead of binding the click to multiple elements, we can easily attach a flyweight to the top of our container which can listen for events coming from below. These can then be handled using as simple or as complex logic needed.

The benefit here is that we're converting many independent objects into a few shared ones (potentially saving on memory), similar to what we were doing with our first JavaScript example.

As the types of components mentioned often have the same repeating markup for each section (e.g. each section of an accordion), there's a good chance the behavior of each element that may be clicked is going to be quite similar and relative to similar classes nearby. We'll use this information to construct a very basic accordion using the flyweight below.

A `StateManager` namespace is used here encapsulate our flyweight logic whilst jQuery is used to bind the initial click to a container div. In order to ensure that no other logic on the page is attaching similar handles to the container, an `unbind` event is first applied.

Now to establish exactly what child element in the container is clicked, we make use of a target check which provides a reference to the element that was clicked, regardless of its parent. We then use this information to handle the click event without actually needing to bind the event to specific children when our page loads.

HTML

```
<div id="container">
  <div class="toggle" href="#">More Info (Address)
    <span class="info">
      This is more information
    </span></div>
  <div class="toggle" href="#">Even More Info (Map)
    <span class="info">
      <iframe src="http://www.map-generator.net/extmap.php?
name=London&address=london%2C%20england&width=500...gt;"></iframe>
    </span>
  </div>
</div>
```

JavaScript

```
var stateManager = {
  fly: function () {
    var self = this;
    $('#container').unbind().bind("click", function (e) {
      var target = $(e.originalTarget || e.srcElement);
      if (target.is("div.toggle")) {
        self.handleClick(target);
      }
    });
  },

  handleClick: function (elem) {
    elem.find('span').toggle('slow');
  }
};
```

Example 2: Using the Flyweight for Performance Gains

In our second example, we'll reference some useful performance gains you can get from applying the flyweight pattern to jQuery.

James Padolsey previously wrote a post called '76 bytes for faster jQuery' where he reminds us of an important point: every time jQuery fires off a callback, regardless of type (filter, each, event handler), you're able to access the function's context (the DOM element related to it) via the `this` keyword.

Unfortunately, many of us have become used to the idea of wrapping this in `$()` or `jQuery()`, which means that a new instance of jQuery is constructed every time.

Rather than doing this:

```
$('#div').on('click', function () {
  console.log('You clicked: ' + $(this).attr('id'));
});

// you should avoid using the DOM element to create a
// jQuery object (with the overhead that comes with it)
// and just use the DOM element itself like this:

$('#div').on('click', function () {
  console.log('You clicked: ' + this.id);
});
```

Now with respect to redundant wrapping, where possible with jQuery's utility methods, it's better to use `jQuery.N` as opposed to `jQuery.fn.N` where `N` represents a utility such as `each`. Because not all of jQuery's methods have corresponding single-node functions, Padolsey devised the idea of `jQuery.single`.

The idea here is that a single jQuery object is created and used for each call to `jQuery.single` (effectively meaning only one jQuery object is ever created). The implementation for this can be found below and is a flyweight as we're consolidating multiple possible objects into a more central singular structure.

```
jQuery.single = (function(o){
  var collection = jQuery([1]);
  return function(element) {
    // Give collection the element:
    collection[0] = element;

    // Return the collection:
```

```
        return collection;
    };
});
```

An example of this in action with chaining is:

```
$('#div').on('click', function(){
    var html = jQuery.single(this).next().html();
    console.log(html);
});
```

Note that although we may believe that simply caching our jQuery code may offer just as equivalent performance gains, Padolsey claims that `$.single()` is still worth using and can perform better. That's not to say don't apply any caching at all, just be mindful that this approach can assist. For further details about `$.single`, I recommend reading Padolsey's full post.

Chapter 11. MV* Patterns

In this section, we're going to review two very important architectural patterns - MVC (Model-View-Controller) and MVP (Model-View-Presenter). In the past both of these patterns have been heavily used for structuring desktop and server-side applications, but it's only been in recent years that come to being applied to JavaScript.

As the majority of JavaScript developers currently using these patterns opt to utilize libraries such as Backbone.js for implementing an MVC/MV*-like structure, we will compare how modern solutions such as it differ in their interpretation of MVC compared to classical takes on these patterns.

Let us first now cover the basics.

MVC

MVC is an architectural design pattern that encourages improved application organization through a separation of concerns. It enforces the isolation of business data (Models) from user interfaces (Views), with a third component (Controllers) (traditionally) managing logic, user-input and coordinating both the models and views. The pattern was originally designed by **Trygve Reenskaug** during his time working on Smalltalk-80 (1979) where it was initially called Model-View-Controller-Editor. MVC went on to be described in depth in “**Design Patterns: Elements of Reusable Object-Oriented Software**” (The "GoF" book) in 1994, which played a role in popularizing its use.

Smalltalk-80 MVC

It's important to understand what the original MVC pattern was aiming to solve as it's mutated quite heavily since the days of its origin. Back in the 70's, graphical user-interfaces were far and few between and a concept known as **Separated Presentation** began to be used as a means to make a clear division between domain objects which modeled concepts in the real world (e.g a photo, a person) and the presentation objects which were rendered to the user's screen.

The Smalltalk-80 implementation of MVC took this concept further and had an

objective of separating out the application logic from the user interface. The idea was that decoupling these parts of the application would also allow the reuse of models for other interfaces in the application. There are some interesting points worth noting about Smalltalk-80's MVC architecture:

- A Domain element was known as a Model and were ignorant of the user-interface (Views and Controllers)
- Presentation was taken care of by the View and the Controller, but there wasn't just a single view and controller. A View-Controller pair was required for each element being displayed on the screen and so there was no true separation between them
- The Controller's role in this pair was handling user input (such as key-presses and click events), doing something sensible with them.
- The Observer pattern was relied upon for updating the View whenever the Model changed

Developers are sometimes surprised when they learn that the Observer pattern (nowadays commonly implemented as a Publish/Subscribe system) was included as a part of MVC's architecture many decades ago. In Smalltalk-80's MVC, the View and Controller both observe the Model. As mentioned in the bullet point above, anytime the Model changes, the Views react. A simple example of this is an application backed by stock market data - in order for the application to be useful, any change to the data in our Models should result in the View being refreshed instantly.

Martin Fowler has done an excellent job of writing about the **origins** of MVC over the years and if you are interested in some further historical information about Smalltalk-80's MVC, I recommend reading his work.

MVC For JavaScript Developers

We've reviewed the 70's, but let us now return to the here and now. In modern times, the MVC pattern has been applied to a diverse range of programming languages including of most relevance to us: JavaScript. JavaScript now has a number of frameworks boasting support for MVC (or variations on it, which we refer to as the MV* family), allowing developers to easily add structure to their applications without great effort. You've likely come across at least one of these such frameworks, but they include the likes of Backbone, Ember.js and JavaScriptMVC. Given the importance of avoiding "spaghetti" code, a term which describes code that is very difficult to read or maintain due to its lack of structure, it's imperative that the modern JavaScript developer understand what this pattern provides. This allows us to effectively appreciate what these frameworks enable us to do differently.

We know that MVC is composed of three core components:

Models

Models manage the data for an application. They are concerned with neither the user-interface nor presentation layers but instead represent unique forms of data that an application may require. When a model changes (e.g when it is updated), it will typically notify its observers (e.g views, a concept we will cover shortly) that a change has occurred so that they may react accordingly.

To understand models further, let us imagine we have a JavaScript photo gallery application. In a photo gallery, the concept of a photo would merit its own model as it represents a unique kind of domain-specific data. Such a model may contain related attributes such as a caption, image source and additional meta-data. A specific photo would be stored in an instance of a model and a model may also be reusable. Below we can see an example of a very simplistic model implemented using Backbone.

```
var Photo = Backbone.Model.extend({  
  // Default attributes for the photo  
  defaults: {  
    src: "placeholder.jpg",  
    caption: "A default image",  
  }  
});
```

```

    viewed: false
  },

  // Ensure that each photo created has an `src`.
  initialize: function() {
    this.set({src: this.defaults.src});
  }

});

```

The built-in capabilities of models vary across frameworks, however it is quite common for them to support validation of attributes, where attributes represent the properties of the model, such as a model identifier. When using models in real-world applications we generally also desire model persistence. Persistence allows us to edit and update models with the knowledge that its most recent state will be saved in either: memory, in a user's localStorage data-store or synchronized with a database.

In addition, a model may also have multiple views observing it. If say, our photo model contained meta-data such as its location (longitude and latitude), friends that were present in the a photo (a list of identifiers) and a list of tags, a developer may decide to provide a single view to display each of these three facets.

It is not uncommon for modern MVC/MV* frameworks to provide a means to group models together (e.g. in Backbone, these groups are referred to as "collections"). Managing models in groups allows us to write application logic based on notifications from the group should any model it contains be changed. This avoids the need to manually observe individual model instances.

A sample grouping of models into a simplified Backbone collection can be seen below.

```

var PhotoGallery = Backbone.Collection.extend({

  // Reference to this collection's model.
  model: Photo,

  // Filter down the list of all photos
  // that have been viewed
  viewed: function() {
    return this.filter(function( photo ){
      return photo.get('viewed');
    });
  },

  // Filter down the list to only photos that
  // have not yet been viewed

```

```
unviewed: function() {  
    return this.without.apply( this, this.viewed() );  
}  
  
});
```

Should you read any of the older texts on MVC, you may come across a description of models as also managing application 'state'. In JavaScript applications "state" has a different meaning, typically referring to the current "state" i.e view or sub-view (with specific data) on a users screen at a fixed point. State is a topic which is regularly discussed when looking at Single-page applications, where the concept of state needs to be simulated.

So to summarize, models are primarily concerned with business data.

Views

Views are a visual representation of models that present a filtered view of their current state. A view typically observes a model and is notified when the model changes, allowing the view to update itself accordingly. Design pattern literature commonly refers to views as 'dumb' given that their knowledge of models and controllers in an application is limited.

Users are able to interact with views and this includes the ability to read and edit (i.e get or set the attribute values in) models. As the view is the presentation layer, we generally present the ability to edit and update in a user-friendly fashion. For example, in the former photo gallery application we discussed earlier, model editing could be facilitated through an "edit" view where a user who has selected a specific photo could edit its meta-data.

The actual task of updating the model falls to controllers (which we'll be covering shortly).

Let's explore views a little further using a vanilla JavaScript sample implementation. Below we can see a function that creates a single Photo view, consuming both a model instance and a controller instance.

We define a `render()` utility within our view which is responsible for rendering the contents of the `photoModel` using a JavaScript templating engine (Underscore templating) and updating the contents of our view, referenced by `photoEl`.

The `photoModel` then adds our `render()` callback as one of it's subscribers so

that through the Observer pattern we can trigger the view to update when the model changes.

You may wonder where user-interaction comes into play here. When users click on any elements within the view, it's not the view's responsibility to know what to do next. It relies on a controller to make this decision for it. In our sample implementation, this is achieved by adding an event listener to `photoEl` which will delegate handling the click behavior back to the controller, passing the model information along with it in case it's needed.

The benefit of this architecture is that each component plays its own separate role in making the application function as needed.

```
var buildPhotoView = function( photoModel, photoController ){

    var base        = document.createElement('div'),
        photoEl     = document.createElement('div');

    base.appendChild(photoEl);

    var render= function(){
        // We use a templating library such as Underscore
        // templating which generates the HTML for our
        // photo entry
        photoEl.innerHTML = _.template('photoTemplate',
        {src: photoModel.getSrc()}));
    }

    photoModel.addSubscriber( render );

    photoEl.addEventListener('click', function(){
        photoController.handleEvent('click', photoModel );
    });

    var show = function(){
        photoEl.style.display = '';
    }

    var hide = function(){
        photoEl.style.display = 'none';
    }

    return{
        showView: show,
        hideView: hide
    }

}
```

Templating

In the context of JavaScript frameworks that support MVC/MV*, it is worth briefly discussing JavaScript templating and its relationship to views as we briefly touched upon it in the last section.

It has long been considered (and proven) a performance bad practice to manually create large blocks of HTML markup in-memory through string concatenation. Developers doing so have fallen prey to inperformantly iterating through their data, wrapping it in nested divs and using outdated techniques such as `document.write` to inject the 'template' into the DOM. As this typically means keeping scripted markup inline with your standard markup, it can quickly become both difficult to read and more importantly, maintain such disasters, especially when building non-trivially sized applications.

JavaScript templating solutions (such as Handlebars.js and Mustache) are often used to define templates for views as markup (either stored externally or within script tags with a custom type - e.g `text/template`) containing template variables. Variables may be delimited using a variable syntax (e.g `{{name}}`) and frameworks are typically smart enough to accept data in a JSON form (of which model instances can be converted to) such that we only need be concerned with maintaining clean models and clean templates. Most of the grunt work to do with population is taken care of by the framework itself. This has a large number of benefits, particularly when opting to store templates externally as this can give way to templates being dynamically loaded on an as-needed basis when it comes to building larger applications.

Below we can see two examples of HTML templates. One implemented using the popular Handlebars.js framework and another using Underscore's templates.

Handlebars.js:

```
<li class="photo">
  <h2>{{caption}}</h2>
  
  <div class="meta-data">
    {{metadata}}
  </div>
</li>
```

Underscore.js Microtemplates:

```
<li class="photo">
  <h2><%= caption %></h2>
  
```

```
<div class="meta-data">  
  <%= metadata %>  
</div>  
</li>
```

It is also worth noting that in classical web development, navigating between independent views required the use of a page refresh. In Single-page JavaScript applications however, once data is fetched from a server via Ajax, it can simply be dynamically rendered in a new view within the same page without any such refresh being necessary. The role of navigation thus falls to a "router", which assists in managing application state (e.g allowing users to bookmark a particular view they have navigated to). As routers are however neither a part of MVC nor present in every MVC-like framework, I will not be going into them in greater detail in this section.

To summarize, views are a visual representation of our application data.

Controllers

Controllers are an intermediary between models and views which are classically responsible for two tasks: they both update the view when the model changes and update the model when the user manipulates the view.

In our photo gallery application, a controller would be responsible for handling changes the user made to the edit view for a particular photo, updating a specific photo model when a user has finished editing.

In terms of where most JavaScript MVC frameworks detract from what is conventionally considered "MVC" however, it is with controllers. The reasons for this vary, but in my honest opinion it is that framework authors initially look at the server-side interpretation of MVC, realize that it doesn't translate 1:1 on the client-side and re-interpret the C in MVC to mean something they feel makes more sense. The issue with this however is that it is subjective, increases the complexity in both understanding the classical MVC pattern and of course the role of controllers in modern frameworks.

As an example, let's briefly review the architecture of the popular architectural framework Backbone.js. Backbone contains models and views (somewhat similar to what we reviewed earlier), however it doesn't actually have true controllers. Its views and routers act a little similar to a controller, but neither are actually

controllers on their own.

In this respect, contrary to what might be mentioned in the official documentation or in blog posts, Backbone is neither a truly MVC/MVP nor MVVM framework. It's in fact better to consider it a member of the MV* family which approaches architecture in its own way. There is of course nothing wrong with this, but it is important to distinguish between classical MVC and MV* should you be relying on advice from classical literature on the former to help with the latter.

Controllers in another library (Spine.js) vs Backbone.js

Spine.js

We now know that controllers are traditionally responsible for updating the view when the model changes (and similarly the model when the user updates the view). As the framework we'll be discussing in this book (Backbone) doesn't have its **own** explicit controllers, it can be useful for us to review the controller from another MVC framework to appreciate the difference in implementations. For this, let's take a look at a sample controller from Spine.js:

In this example, we're going to have a controller called `PhotosController` which will be in charge of individual photos in the application. It will ensure that when the view updates (e.g a user editd the photo meta-data) the corresponding model does too.

Note: We won't be delving heavily into Spine.js at all, but will just take a ten-foot view of what its controllers can do:

```
// Controllers in Spine are created by inheriting from Spine.Controller

var PhotosController = Spine.Controller.sub({
  init: function(){
    this.item.bind("update", this.proxy(this.render));
    this.item.bind("destroy", this.proxy(this.remove));
  },

  render: function(){
    // Handle templating
    this.replace($("#photoTemplate").tpl(this.item));
    return this;
  },

  remove: function(){
```



```
        this.el.remove();  
        this.release();  
    }  
});
```

In Spine, controllers are considered the glue for an application, adding and responding to DOM events, rendering templates and ensuring that views and models are kept in sync (which makes sense in the context of what we know to be a controller).

What we're doing in the above example is setting up listeners in the `update` and `destroy` events using `render()` and `remove()`. When a photo entry gets updated, we re-render the view to reflect the changes to the meta-data. Similarly, if the photo gets deleted from the gallery, we remove it from the view. In case you were wondering about the `tpl()` function in the code snippet: in the `render()` function, we're using this to render a JavaScript template called `#photoTemplate` which simply returns a HTML string used to replace the controller's current element.

What this provides us with is a very lightweight, simple way to manage changes between the model and the view.

Backbone.js

Later on in this section we're going to revisit the differences between Backbone and traditional MVC, but for now let's focus on controllers.

In Backbone, one shares the responsibility of a controller with both the `Backbone.View` and `Backbone.Router`. Some time ago Backbone did once come with its own `Backbone.Controller`, but as the naming for this component didn't make sense for the context in which it was being used, it was later renamed to `Router`.

Routers handle a little more of the controller responsibility as it's possible to bind the events there for models and have your view respond to DOM events and rendering. As Tim Branyen (another Bocoup-based Backbone contributor) has also previously pointed out, it's possible to get away with not needing `Backbone.Router` at all for this, so a way to think about it using the Router paradigm is probably:

```
var PhotoRouter = Backbone.Router.extend({  
  routes: { "photos/:id": "route" },
```

```
route: function(id) {  
  var item = photoCollection.get(id);  
  var view = new PhotoView({ model: item });  
  
  something.html( view.render().el );  
}  
}):
```

To summarize, the takeaway from this section is that controllers manage the logic and coordination between models and views in an application.

What does MVC give us?

This separation of concerns in MVC facilitates simpler modularization of an application's functionality and enables:

- Easier overall maintenance. When updates need to be made to the application it is very clear whether the changes are data-centric, meaning changes to models and possibly controllers, or merely visual, meaning changes to views.
- Decoupling models and views means that it is significantly more straightforward to write unit tests for business logic
- Duplication of low-level model and controller code (i.e what you may have been using instead) is eliminated across the application
- Depending on the size of the application and separation of roles, this modularity allows developers responsible for core logic and developers working on the user-interfaces to work simultaneously

Delving deeper

Right now, you likely have a basic understanding of what the MVC pattern provides, but for the curious, we can explore it a little further.

The GoF (Gang of Four) do not refer to MVC as a design pattern, but rather consider it a "set of classes to build a user interface". In their view, it's actually a variation of three other classical design patterns: the Observer (Pub/Sub), Strategy and Composite patterns. Depending on how MVC has been implemented in a framework, it may also use the Factory and Decorator patterns.

As we've discussed, models represent application data whilst views are what the user is presented on screen. As such, MVC relies on Pub/Sub for some of its core communication (something that surprisingly isn't covered in many articles about the MVC pattern). When a model is changed it notifies the rest of the application it has been updated. The controller then updates the view accordingly. The observer nature of this relationship is what facilitates multiple views being attached to the same model.

For developers interested in knowing more about the decoupled nature of MVC

(once again, depending on the implementation), one of the goals of the pattern is to help define one-to-many relationships between a topic and its observers. When a topic changes, its observers are updated. Views and controllers have a slightly different relationship. Controllers facilitate views to respond to different user input and are an example of the Strategy pattern.

Summary

Having reviewed the classical MVC pattern, we should now understand how it allows us to cleanly separate concerns in an application. We should also now appreciate how JavaScript MVC frameworks may differ in their interpretation of the MVC pattern, which although quite open to variation, still shares some of the fundamental concepts the original pattern has to offer.

When reviewing a new JavaScript MVC/MV* framework, remember - it can be useful to step back and review how it's opted to approach architecture (specifically, how it supports implementing models, views, controllers or other alternatives) as this can better help you grok how the framework expects to be used.

MVP

Model-view-presenter (MVP) is a derivative of the MVC design pattern which focuses on improving presentation logic. It originated at a company named **Taligent** in the early 1990s while they were working on a model for a C++ CommonPoint environment. Whilst both MVC and MVP target the separation of concerns across multiple components, there are some fundamental differences between them.

For the purposes of this summary we will focus on the version of MVP most suitable for web-based architectures.

Models, Views & Presenters

The P in MVP stands for presenter. It's a component which contains the user-interface business logic for the view. Unlike MVC, invocations from the view are delegated to the presenter, which are decoupled from the view and instead talk to it through an interface. This allows for all kinds of useful things such as being able to mock views in unit tests.

The most common implementation of MVP is one which uses a Passive View (a view which is for all intents and purposes "dumb"), containing little to no logic. MVP models are almost identical to MVC models and handle application data. The presenter acts as a mediator which talks to both the view and model, however both of these are isolated from each other. They effectively bind models to views, a responsibility which was previously held by controllers in MVC. Presenters are at the heart of the MVP pattern and as you can guess, incorporate the presentation logic behind views.

Solicited by a view, presenters perform any work to do with user requests and pass data back to them. In this respect, they retrieve data, manipulate it and determine how the data should be displayed in the view. In some implementations, the presenter also interacts with a service layer to persist data (models). Models may trigger events but it's the presenters role to subscribe to them so that it can update the view. In this passive architecture, we have no concept of direct data binding. Views expose setters which presenters can use to set data.

The benefit of this change from MVC is that it increases the testability of your

application and provides a more clean separation between the view and the model. This isn't however without its costs as the lack of data binding support in the pattern can often mean having to take care of this task separately.

Although a common implementation of a **Passive View** is for the view to implement an interface, there are variations on it, including the use of events which can decouple the View from the Presenter a little more. As we don't have the interface construct in JavaScript, we're using more a protocol than an explicit interface here. It's technically still an API and it's probably fair for us to refer to it as an interface from that perspective.

There is also a **Supervising Controller** variation of MVP, which is closer to the MVC and **MVVM** patterns as it provides data-binding from the Model directly from the View. Key-value observing (KVO) plugins (such as Derick Bailey's Backbone.ModelBinding plugin) tend to bring Backbone out of the Passive View and more into the Supervising Controller or MVVM variations.

MVP or MVC?

MVP is generally used most often in enterprise-level applications where it's necessary to reuse as much presentation logic as possible. Applications with very complex views and a great deal of user interaction may find that MVC doesn't quite fit the bill here as solving this problem may mean heavily relying on multiple controllers. In MVP, all of this complex logic can be encapsulated in a presenter, which can simplify maintenance greatly.

As MVP views are defined through an interface and the interface is technically the only point of contact between the system and the view (other than a presenter), this pattern also allows developers to write presentation logic without needing to wait for designers to produce layouts and graphics for the application.

Depending on the implementation, MVP may be more easy to automatically unit test than MVC. The reason often cited for this is that the presenter can be used as a complete mock of the user-interface and so it can be unit tested independent of other components. In my experience this really depends on the languages you are implementing MVP in (there's quite a difference between opting for MVP for a JavaScript project over one for say, ASP.net).

At the end of the day, the underlying concerns you may have with MVC will likely hold true for MVP given that the differences between them are mainly semantic. As long as you are cleanly separating concerns into models, views and controllers (or presenters) you should be achieving most of the same benefits regardless of the pattern you opt for.

MVC, MVP and Backbone.js

There are very few, if any architectural JavaScript frameworks that claim to implement the MVC or MVP patterns in their classical form as many JavaScript developers don't view MVC and MVP as being mutually exclusive (we are actually more likely to see MVP strictly implemented when looking at web frameworks such as ASP.net or GWT). This is because it's possible to have additional presenter/view logic in your application and yet still consider it a flavor of MVC.

Backbone contributor **Irene Ros** (of Boston-based Bocoup) subscribes to this way of thinking as when she separates views out into their own distinct components, she needs something to actually assemble them for her. This could either be a controller route (such as a `Backbone.Router`, covered later in the book) or a callback in response to data being fetched.

That said, some developers do however feel that Backbone.js better fits the description of MVP than it does MVC. Their view is that:

- The presenter in MVP better describes the `Backbone.View` (the layer between View templates and the data bound to it) than a controller does
- The model fits `Backbone.Model` (it isn't greatly different to the models in MVC at all)
- The views best represent templates (e.g Handlebars/Mustache markup templates)

A response to this could be that the view can also just be a View (as per MVC) because Backbone is flexible enough to let it be used for multiple purposes. The V in MVC and the P in MVP can both be accomplished by `Backbone.View` because they're able to achieve two purposes: both rendering atomic components and assembling those components rendered by other views.

We've also seen that in Backbone the responsibility of a controller is shared with

both the Backbone.View and Backbone.Router and in the following example we can actually see that aspects of that are certainly true.

Our Backbone PhotoView uses the Observer pattern to 'subscribe' to changes to a View's model in the line `this.model.bind('change', ...)`. It also handles templating in the `render()` method, but unlike some other implementations, user interaction is also handled in the View (see `events`).

```
var PhotoView = Backbone.View.extend({

  //... is a list tag.
  tagName: "li",

  // Pass the contents of the photo template through a templating
  // function, cache it for a single photo
  template: _.template($('#photo-template').html()),

  // The DOM events specific to an item.
  events: {
    "click img" : "toggleViewed"
  },

  // The PhotoView listens for changes to
  // its model, re-rendering. Since there's
  // a one-to-one correspondence between a
  // **Photo** and a **PhotoView** in this
  // app, we set a direct reference on the model for convenience.

  initialize: function() {
    _.bindAll(this, 'render');
    this.model.bind('change', this.render);
    this.model.bind('destroy', this.remove);
  },

  // Re-render the photo entry
  render: function() {
    $(this.el).html(this.template(this.model.toJSON()));
    return this;
  },

  // Toggle the `"viewed"` state of the model.
  toggleViewed: function() {
    this.model.viewed();
  }

});
```

Another (quite different) opinion is that Backbone more closely resembles **Smalltalk-80 MVC**, which we went through earlier.

As regular Backbone user Derick Bailey has **previously** put it, it's ultimately best

not to force Backbone to fit any specific design patterns. Design patterns should be considered flexible guides to how applications may be structured and in this respect, Backbone fits neither MVC nor MVP. Instead, it borrows some of the best concepts from multiple architectural patterns and creates a flexible framework that just works well.

It is however worth understanding where and why these concepts originated, so I hope that my explanations of MVC and MVP have been of help. Call it **the Backbone way**, MV* or whatever helps reference its flavor of application architecture. Most structural JavaScript frameworks will adopt their own take on classical patterns, either intentionally or by accident, but the important thing is that they help us develop applications which are organized, clean and can be easily maintained.

MVVM

MVVM (Model View ViewModel) is an architectural pattern based on MVC and MVP, which attempts to more clearly separate the development of user-interfaces (UI) from that of the business logic and behavior in an application. To this end, many implementations of this pattern make use of declarative data bindings to allow a separation of work on Views from other layers.

This facilitates UI and development work occurring almost simultaneously within the same codebase. UI developers write bindings to the ViewModel within their document markup (HTML), where the Model and ViewModel are maintained by developers working on the logic for the application.

History

MVVM (by name) was originally defined by Microsoft for use with Windows Presentation Foundation (**WPF**) and **Silverlight**, having been officially announced in 2005 by **John Grossman** in a blog post about Avalon (the codename for WPF). It also found some popularity in the Adobe Flex community as an alternative to simply using MVC.

Prior to Microsoft adopting the MVVM name, there was however a movement in the community to go from MVP to MVPM: Model View **PresentationModel**. Martin Fowler wrote an **article** on PresentationModels back in 2004 for those interested in reading more about it. The idea of a **PresentationModel** had been around much longer than this article, however it was considered the big break in the idea and greatly helped popularize it.

There was quite a lot of uproar in the "alt.net" circles after Microsoft announced MVVM as an alternative to MVPM. Many claimed the company's dominance in the GUI world was giving them the opportunity to take over the community as a whole, renaming existing concepts as they pleased for marketing purposes. A progressive crowd recognized that whilst MVVM and MVPM were effectively the same idea, they came in slightly different packages.

In recent years, MVVM has been implemented in JavaScript in the form of structural frameworks such as **KnockoutJS**, **Kendo MVVM** and **Knockback.js**,

with an overall positive response from the community.

Let's now review the three components that compose MVVM.

Model

As with other members of the MV* family, the Model in MVVM represents domain-specific data or information that our application will be working with. A typical example of domain-specific data might be a user account (e.g name, avatar, e-mail) or a music track (e.g title, year, album).

Models hold information, but typically don't handle behavior. They don't format information or influence how data appears in the browser as this isn't their responsibility. Instead, formatting of data is handled by the View, whilst behavior is considered business logic that should be encapsulated in another layer that interacts with the Model - the ViewModel.

The only exception to this rule tends to be validation and it's considered acceptable for Models to validate data being used to define or update existing models (e.g does an e-mail address being input meet the requirements of a particular Regular expression?).

In KnockoutJS, Models fall under the above definition, but often make Ajax calls to a server-side service to both read and write Model data.

If we were constructing a simple Todo application, a KnockoutJS Model representing a single Todo item could look as follows:

```
var Todo = function (content, done) {  
    this.content = ko.observable(content);  
    this.done = ko.observable(done);  
    this.editing = ko.observable(false);  
};
```

Note: You may notice in the above snippet that we are calling a method `observables()` on the KnockoutJS namespace `ko`. In KnockoutJS, observables are special JavaScript objects that can notify subscribers about changes and automatically detect dependencies. This allows us to synchronize Models and ViewModels when the value of a Model attribute is modified.

View

As with MVC, the View is the only part of the application of users actually interact with. They are an interactive UI that represent the state of a ViewModel. In this sense, MVVM View is considered active rather than passive, but what does this mean?.

A passive View has no real knowledge of the models in our application and is manipulated by a controller. MVVM's active View contains the data-bindings, events and behaviors which require an understanding of the Model and ViewModel. Although these behaviors can be mapped to properties, the View is still responsible for handling events to the ViewModel.

It's important to remember the View isn't responsible here for handling state - it keeps this in sync with the ViewModel.

A KnockoutJS View is simply a HTML document with declarative bindings to link it to the ViewModel. KnockoutJS Views display information from the ViewModel, pass commands to it (e.g a user clicking on an element) and update as the state of the ViewModel changes. Templates generating markup using data from the ViewModel can however also be used for this purpose.

To give a brief initial example, we can look to the JavaScript MVVM framework KnockoutJS for how it allows the definition of a ViewModel and it's related bindings in markup:

ViewModel:

```
var aViewModel = {  
  contactName: ko.observable('John');  
};
```

View:

```
<input id="source" data-bind="value: contactName, valueUpdate: 'keyup'" /></p>  
  
<div data-bind="visible: contactName().length > 10">  
  You have a really long name!  
</div>
```

Our input text-box (source) obtains it's initial value from `contactName`, automatically updating this value whenever `contactName` changes. As the data binding is two-way, typing into the text-box will update `contactName` accordingly so the values are always in sync.

Although implementation specific to KnockoutJS, the `<div>` containing the 'You

have a really long name! text also contains simple validation (once again in the form of data bindings). If the input exceeds 10 characters, it will display, otherwise it will remain hidden.

Moving on to a more advanced example, we can return to our Todo application. A trimmed down KnockoutJS View for this, including all the necessary data-bindings may look as follows.

```
<div id="todoapp">
  <header>
    <h1>Todos</h1>
    <input id="new-todo" type="text" data-bind="value: current, valueUpdate:
'afterkeydown', enterKey: add"
      placeholder="What needs to be done?"/>
  </header>
  <section id="main" data-bind="block: todos().length">

    <input id="toggle-all" type="checkbox" data-bind="checked: allCompleted">
    <label for="toggle-all">Mark all as complete</label>

    <ul id="todo-list" data-bind="foreach: todos">

      <!-- item -->
      <li data-bind="css: { done: done, editing: editing }">
        <div class="view" data-bind="event: { dblclick: $root.editItem }">
          <input class="toggle" type="checkbox" data-bind="checked: done">
          <label data-bind="text: content"></label>
          <a class="destroy" href="#" data-bind="click: $root.remove"></a>
        </div>
        <input class="edit" type="text"
          data-bind="value: content, valueUpdate: 'afterkeydown',
enterKey: $root.stopEditing, selectAndFocus: editing, event: { blur:
$root.stopEditing }"/>
        </li>

      </ul>

    </section>
  </div>
```

Note that the basic layout of the mark-up is relatively straight-forward, containing an input textbox (new-todo) for adding new items, togglers for marking items as complete and a list (todo-list) with a template for a Todo item in the form of an li.

The data bindings in the above markup can be broken down as follows:

- The input textbox new-todo has a data-binding for the current property, which is where the value of the current item being added is stored. Our

ViewModel (shown shortly) observes the `current` property and also has a binding against the `add` event. When the enter key is pressed, the `add` event is triggered and our ViewModel can then trim the value of `current` and add it to the Todo list as needed

- The input checkbox `toggle-all` can mark all of the current items as completed if clicked. If checked, it triggers the `allCompleted` event, which can be seen in our ViewModel
- The item `li` has the class `done`. When a task is marked as done, the CSS class `editing` is marked accordingly. If double-clicking on the item, the `$root.editItem` callback will be executed
- The checkbox with the class `toggle` shows the state of the `done` property
- A label contains the text value of the Todo item (`content`)
- There is also a remove button that will call the `$root.remove` callback when clicked.
- An input textbox used for editing mode also holds the value of the Todo item content. The `enterKey` event will set the `editing` property to true or false

ViewModel

The ViewModel can be considered a specialized Controller that acts as a data converter. It changes Model information into View information, passing commands from the View to the Model.

For example, let us imagine that we have a model containing a date attribute in unix format (e.g 1333832407). Rather than our models being aware of a user's view of the date (e.g 04/07/2012 @ 5:00pm), where it would be necessary to convert the address to it's display format, our model simply holds the raw format of the data. Our View contains the formatted date and our ViewModel acts as a middle-man between the two.

In this sense, the ViewModel might be looked upon as more of a Model than a View but it does handle most of the View's display logic. The ViewModel may also expose methods for helping to maintain the View's state, update the model based on the action's on a View and trigger events on the View.

In summary, the ViewModel sits behind our UI layer. It exposes data needed by a View (from a Model) and can be viewed as the source our Views go to for both data and actions.

KnockoutJS interprets the ViewModel as the representation of data and operations that can be performed on a UI. This isn't the UI itself nor the data model that persists, but rather a layer that can also hold the yet to be saved data a user is working with. Knockout's ViewModels are implemented JavaScript objects with no knowledge of HTML markup. This abstract approach to their implementation allows them to stay simple, meaning more complex behavior can be more easily managed on-top as needed.

A partial KnockoutJS ViewModel for our Todo application could thus look as follows:

```
// our main ViewModel
var ViewModel = function (todos) {
    var self = this;

    // map array of passed in todos to an observableArray of Todo objects
    self.todos = ko.observableArray(ko.utils.arrayMap(todos, function (todo) {
        return new Todo(todo.content, todo.done);
    }));

    // store the new todo value being entered
    self.current = ko.observable();

    // add a new todo, when enter key is pressed
    self.add = function (data, event) {
        var newTodo, current = self.current().trim();
        if (current) {
            newTodo = new Todo(current);
            self.todos.push(newTodo);
            self.current("");
        }
    };

    // remove a single todo
    self.remove = function (todo) {
        self.todos.remove(todo);
    };

    // remove all completed todos
    self.removeCompleted = function () {
        self.todos.remove(function (todo) {
            return todo.done();
        });
    };

    // writeable computed observable to handle marking all complete/incomplete
```

```

        self.allCompleted = ko.computed({
            //always return true/false based on the done flag of all todos
            read:function () {
                return !self.remainingCount();
            },
            //set all todos to the written value (true/false)
            write:function (newValue) {
                ko.utils.arrayForEach(self.todos(), function (todo) {
                    //set even if value is the same, as subscribers are not notified in
that case
                    todo.done(newValue);
                });
            }
        });

        // edit an item
        self.editItem = function(item) {
            item.editing(true);
        };
        ..

```

Above we are basically providing the methods needed to add, edit or remove items as well as the logic to mark all remaining items as having been completed Note: The only real difference worth noting from previous examples in our ViewModel are observable arrays. In KnockoutJS, if we wish to detect and respond to changes on a single object, we would use **observables**. If however we wish to detect and respond to changes of a collection of things, we can use an **observableArray** instead. A simpler example of how to use observables arrays may look as follows:

```

// Define an initially an empty array
var myObservableArray = ko.observableArray();

// Add a value to the array and notify our observers
myObservableArray.push('A new todo item');

```

Note: The complete Knockout.js Todo application we reviewed above can be grabbed from [TodoMVC](#) if interested.

Recap: The View and the ViewModel

Views and ViewModels communicate using data-bindings and events. As we saw in our initial ViewModel example, the ViewModel doesn't just expose Model attributes but also access to other methods and features such as validation.

Our Views handle their own user-interface events, mapping them to the ViewModel as necessary. Models and attributes on the ViewModel are synchronized and updated via two-way data-binding.

Triggers (data-triggers) also allow us to further react to changes in the state of our Model attributes.

Recap: The ViewModel and the Model

Whilst it may appear the ViewModel is completely responsible for the Model in MVVM, there are some subtleties with this relationship worth noting. The ViewModel can expose a Model or Model attributes for the purposes of data-binding and can also contain interfaces for fetching and manipulating properties exposed in the view.

Pros and Cons

You now hopefully have a better appreciation for what MVVM is and how it works. Let's now review the advantages and disadvantages of employing the pattern:

Advantages

- MVVM Facilitates easier parallel development of a UI and the building blocks that power it
- Abstracts the View and thus reduces the quantity of business logic (or glue) required in the code behind it
- The ViewModel can be easier to unit test than event-driven code
- The ViewModel (being more Model than View) can be tested without concerns of UI automation and interaction

Disadvantages

- For simpler UIs, MVVM can be overkill
- Whilst data-bindings can be declarative and nice to work with, they can be harder to debug than imperative code where we simply set breakpoints
- Data-bindings in non-trivial applications can create a lot of book-keeping. You also don't want to end up in a situation where bindings are heavier than the objects being bound to
- In larger applications, it can be more difficult to design the ViewModel up front to get the necessary amount of generalization

MVVM With Looser Data-Bindings

It's not uncommon for JavaScript developers from an MVC or MVP background to review MVVM and complain about its true separation of concerns. Namely, the quantity of inline data-bindings maintained in the HTML markup of a View.

I must admit that when I first reviewed implementations of MVVM (e.g KnockoutJS, Knockback), I was surprised that any developer would want to return to the days of old where we mixed logic (JavaScript) with our markup and found it quickly unmaintainable. The reality however is that MVVM does this for a number of good reasons (which we've covered), including facilitating designers to more easily bind to logic from their markup.

For the purists among us, you'll be happy to know that we can now also greatly reduce how reliant we are on data-bindings thanks to a feature known as custom binding providers, introduced in KnockoutJS 1.3 and available in all versions since.

KnockoutJS by default has a data-binding provider which searches for any elements with `data-bind` attributes on them such as in the below example.

```
<input id="new-todo" type="text" data-bind="value: current, valueUpdate:
  'afterkeydown', enterKey: add" placeholder="What needs to be done?"/>
```

When the provider locates an element with this attribute, it parses it and turns it into a binding object using the current data context. This is the way KnockoutJS has always worked, allowing you to declaratively add bindings to elements which KnockoutJS binds to the data at that layer.

Once you start building Views that are no longer trivial, you may end up with a large number of elements and attributes whose bindings in markup can become difficult to manage. With custom binding providers however, this is no longer a problem.

A binding provider is primarily interested in two things:

- When given a DOM node, does it contain any data-bindings?
- If the node passed this first question, what does the binding object look like in the current data context?.

Binding providers implement two functions:

- `nodeHasBindings`: this takes in a DOM node which doesn't necessarily have to be an element
- `getBindings`: returns an object representing the bindings as applied to the current data context

A skeleton binding provider might thus look as follows:

```
var ourBindingProvider = {
  nodeHasBindings: function(node) {
    // returns true/false
  },

  getBindings: function(node, bindingContext) {
    // returns a binding object
  }
};
```

Before we get to fleshing out this provider, let's briefly discuss logic in data-bind attributes.

If when using Knockout's MVVM you find yourself dissatisfied with the idea of application logic being overly tied into your View, you can change this. We could implement something a little like CSS classes to assign bindings by name to elements. Ryan Niemeyer (of knockmeout.net) has previously suggested using `data-class` for this to avoid confusing presentation classes with data classes, so let's get our `nodeHasBindings` function supporting this:

```
// does an element have any bindings?
function nodeHasBindings(node) {
  return node.getAttribute ? node.getAttribute("data-class") : false;
};
```

Next, we need a sensible `getBindings()` function. As we're sticking with the idea of CSS classes, why not also consider supporting space-separated classes to allow us to share binding specs between different elements?.

Let's first review what our bindings will look like. We create an object to hold them where our property names need to match the keys we wish to use in our data-classes.

Note: There isn't a great deal of work required to convert a KnockoutJS application from using traditional data-bindings over to unobstrusive bindings with custom binding providers. We simply pull out all of our data-bind attributes, replace them with data-class attributes and place our bindings in a binding object as per below:

```

var viewModel = new ViewModel(todos || []);
var bindings = {

    newTodo: {
        value: viewModel.current,
        valueUpdate: 'afterkeydown',
        enterKey: viewModel.add
    },
    taskTooltip : { visible: viewModel.showTooltip },
    checkAllContainer : {visible: viewModel.todos().length },
    checkAll: {checked: viewModel.allCompleted },

    todos: {foreach: viewModel.todos },
    todoListItem: function() { return { css: { editing: this.editing } }; },
    todoListItemWrapper: function() { return { css: { done: this.done } }; },
    todoCheckBox: function() {return { checked: this.done }; },
    todoContent: function() { return { text: this.content, event: { dblclick:
this.edit } };},
    todoDestroy: function() {return { click: viewModel.remove };},

    todoEdit: function() { return {
        value: this.content,
        valueUpdate: 'afterkeydown',
        enterKey: this.stopEditing,
        event: { blur: this.stopEditing } }; },

    todoCount: {visible: viewModel.remainingCount},
    remainingCount: { text: viewModel.remainingCount },
    remainingCountWord: function() { return { text:
viewModel.getLabel(viewModel.remainingCount) };},

    todoClear: {visible: viewModel.completedCount},
    todoClearAll: {click: viewModel.removeCompleted},
    completedCount: { text: viewModel.completedCount },
    completedCountWord: function() { return { text:
viewModel.getLabel(viewModel.completedCount) }; },

    todoInstructions: {visible: viewModel.todos().length}
};

....

```

There are however two lines missing from the above snippet - we still need our `getBindings` function, which will loop through each of the keys in our data-class attributes and build up the resulting object from each of them. If we detect that the binding object is a function, we call it with our current data using the context `this`. Our complete custom binding provider would look as follows:

```

// We can now create a bindingProvider that uses
// something different than data-bind attributes
ko.customBindingProvider = function(bindingObject) {
    this.bindingObject = bindingObject;

```

```

        //determine if an element has any bindings
        this.nodeHasBindings = function(node) {
            return node.getAttribute ? node.getAttribute("data-class") : false;
        };

    // return the bindings given a node and the bindingContext
    this.getBindings = function(node, bindingContext) {
        var result = {};
        var classes = node.getAttribute("data-class");
        if (classes) {
            classes = classes.split(' ');
            //evaluate each class, build a single object to return
            for (var i = 0, j = classes.length; i < j; i++) {
                var bindingAccessor = this.bindingObject[classes[i]];
                if (bindingAccessor) {
                    var binding = typeof bindingAccessor == "function" ?
bindingAccessor.call(bindingContext.$data) : bindingAccessor;
                    ko.utils.extend(result, binding);
                }
            }
        }
        return result;
    };
};

```

Thus, the final few lines of our bindings object can be defined as follows:

```

    // set ko's current bindingProvider equal to our new binding provider
    ko.bindingProvider.instance = new ko.customBindingProvider(bindings);

    // bind a new instance of our ViewModel to the page
    ko.applyBindings(viewModel);
})();

```

What we're doing here is effectively defining constructor for our binding handler which accepts an object (bindings) which we use to lookup our bindings. We could then re-write the markup for our application View using data-classes as follows:

```

<div id="create-todo">
    <input id="new-todo" data-class="newTodo" placeholder="What needs to
be done?" />
    <span class="ui-tooltip-top" data-class="taskTooltip"
style="display: none;">Press Enter to save this task</span>
</div>
<div id="todos">
    <div data-class="checkAllContainer" >
        <input id="check-all" class="check" type="checkbox" data-
class="checkAll" />
        <label for="check-all">Mark all as complete</label>
    </div>
    <ul id="todo-list" data-class="todos" >
        <li data-class="todoListItem" >
            <div class="todo" data-class="todoListItemWrapper" >

```

```

        <div class="display">
            <input class="check" type="checkbox" data-
class="todoCheckBox" />
            <div class="todo-content" data-class="todoContent"
style="cursor: pointer;"></div>
            <span class="todo-destroy" data-class="todoDestroy">
</span>
        </div>
        <div class="edit">
            <input class="todo-input" data-class="todoEdit"/>
        </div>
    </div>
</li>
</ul>
</div>

```

Neil Kerkin has put together a complete TodoMVC demo app using the above, which can be accessed and played around with [here](#).

Whilst it may look like quite a lot of work in the explanation above, now that you have a generic `getBindings` method written, it's a lot more trivial to simply re-use it and use data-classes rather than strict data-bindings for writing your KnockoutJS applications instead. The net result is hopefully cleaner markup with your data bindings being shifted from the View to a bindings object instead.

MVC Vs. MVP Vs. MVVM

Both MVP and MVVM are derivatives of MVC. The key difference between it and its derivatives is the dependency each layer has on other layers as well as how tightly bound they are to each other.

In MVC, the View sits on top of our architecture with the controller laying below this. Models sit below the controller and so our Views know about our controllers and controllers know about Models. Here, our Views have direct access to Models. Exposing the complete Model to the View however may have security and performance costs, depending on the complexity of our application. MVVM attempts to avoid these issues.

In MVP, the role of the controller is replaced with a Presenter. Presenters sit at the same level as views, listening to events from both the View and model and mediating the actions between them. Unlike MVVM, there isn't a mechanism for binding Views to ViewModels, so we instead rely on each View implementing an interface allowing the Presenter to interact with the View.

MVVM consequently allows us to create View-specific subsets of a Model which can contain state and logic information, avoiding the need to expose the entire Model to a View. Unlike MVP's Presenter, a ViewModel is not required to reference a View. The View can bind to properties on the ViewModel which in turn expose data contained in Models to the View. As we've mentioned, the abstraction of the View means there is less logic required in the code behind it.

One of the downsides to this however is that a level of interpretation is needed between the ViewModel and the View and this can have performance costs. The complexity of this interpretation can also vary - it can be as simple as copying data or as complex as manipulating them to a form we would like the View to see. MVC doesn't have this problem as the whole Model is readily available and such manipulation can be avoided.

Backbone.js Vs. KnockoutJS

Understanding the subtle differences between MVC, MVP and MVVM are important but developers ultimately will ask whether they should consider using KnockoutJS over Backbone based in what we've learned. The following notes may be of help here:

- Both libraries are designed with different goals in mind and its often not as simple as just choosing MVC or MVVM
- If data-binding and two-way communication are are your main concerns, KnockoutJS is definitely the way to go. Practically any attribute or value stored in DOM nodes can be mapped to JavaScript objects with this approach.
- Backbone excels with its ease of integration with RESTful services, whilst KnockoutJS Models are simply JavaScript objects and code needed for updating the Model must be written by the developer.
- KnockoutJS has a focus on automating UI bindings, which requires significantly more verbose custom code if attempting to do this with Backbone. This isn't a problem with Backbone itself par se as it purposefully attempts to stay out of the UI. Knockback does however attempt to assist with this problem.
- With KnockoutJS, we can bind our own functions to ViewModel observables, which are executed anytime the observable changes. This allows us the same level of flexibility as can be found in Backbone
- Backbone has a solid routing solution built-in, whilst KnockoutJS offers no routing options out of the box. One can however easily fill this behavior in if needed using Ben Alman's [BBQ plugin](#) or a standalone routing system like Miller Medeiros's excellent [Crossroads](#).

To conclude, I personally find KnockoutJS more suitable for smaller applications whilst Backbone's feature set really shines when building anything non-trivial. That said, many developers have used both frameworks to write applications of varying complexity and I recommend trying out both at a smaller scale before making a decision on which might work best for your project.

If you wish to read more about MVVM or Knockout, further recommended reading can be found below:

- The Advantages Of MVVM
- SO: What are the problems with MVVM?
- MVVM Explained
- How does MVVM compare to MVC?
- Custom bindings in KnockoutJS
- Exploring Knockout with TodoMVC

Namespacing Patterns

In this section, I'll be discussing both intermediate and advanced conventions for namespacing in JavaScript. We're going to begin with the latter, however if you're new to namespacing with the language and would like to learn more about some of the fundamentals, please feel free to skip to the section titled '[namespacing fundamentals](#)' to continue reading.

What is namespacing?

In many programming languages, namespacing is a technique employed to avoid **collisions** with other objects or variables in the global namespace. They're also extremely useful for helping organize blocks of functionality in your application into easily manageable groups that can be uniquely identified.

In JavaScript, namespacing at an enterprise level is critical as it's important to safeguard your code from breaking in the event of another script on the page using the **same** variable or method names as you are. With the number of **third-party** tags regularly injected into pages these days, this can be a common problem we all need to tackle at some point in our careers. As a well-behaved 'citizen' of the global namespace, it's also imperative that you do your best to similarly not prevent other developer's scripts executing due to the same issues.

Whilst JavaScript doesn't really have built-in support for namespaces like other languages, it does have objects and closures which can be used to achieve a similar effect.

Advanced namespacing patterns

In this section, I'll be exploring some advanced patterns and utility techniques that have helped me when working on larger projects requiring a re-think of how application namespacing is approached. I should state that I'm not advocating any of these as **the** way to do things, but rather just ways that I've found work in practice.

Automating nested namespacing

As you're probably aware, a nested namespace provides an organized hierarchy of structures in an application and an example of such a namespace could be the following: application.utilities.drawing.canvas.2d. In JavaScript the equivalent of this definition using the object literal pattern would be:

```
var application = {  
  utilities:{  
    drawing:{  
      canvas:{  
        2d:{  
          //...  
        }  
      }  
    }  
  }  
};
```

Wow, that could be done better.

One of the obvious challenges with this pattern is that each additional depth you wish to create requires yet another object to be defined as a child of some parent in your top-level namespace. This can become particularly laborious when multiple depths are required as your application increases in complexity.

How can this problem be better solved? In [JavaScript Patterns](#), [Stoyan Stefanov](#) presents a very-clever approach for automatically defining nested namespaces under an existing global variable using a convenience method that takes a single string argument for a nest, parses this and automatically populates your base namespace with the objects required.

The method he suggests using is the following, which I've updated it to be a generic function for easier re-use with multiple namespaces:

```

// top-level namespace being assigned an object literal
var myApp = myApp || {};

// a convenience function for parsing string namespaces and
// automatically generating nested namespaces
function extend( ns, ns_string ) {
    var parts = ns_string.split('.'),
        parent = ns,
        pl, i;

    pl = parts.length;
    for (i = 0; i < pl; i++) {
        // create a property if it doesnt exist
        if (typeof parent[parts[i]] == 'undefined') {
            parent[parts[i]] = {};
        }

        parent = parent[parts[i]];
    }

    return parent;
}

// sample usage:
// extend myApp with a deeply nested namespace
var mod = extend(myApp, 'myApp.modules.module2');
// the correct object with nested depths is output
console.log(mod);
// minor test to check the instance of mod can also
// be used outside of the myApp namespace as a clone
// that includes the extensions
console.log(mod == myApp.modules.module2); //true
// further demonstration of easier nested namespace
// assignment using extend
extend(myApp, 'moduleA.moduleB.moduleC.moduleD');
extend(myApp, 'longer.version.looks.like.this');
console.log(myApp);

```

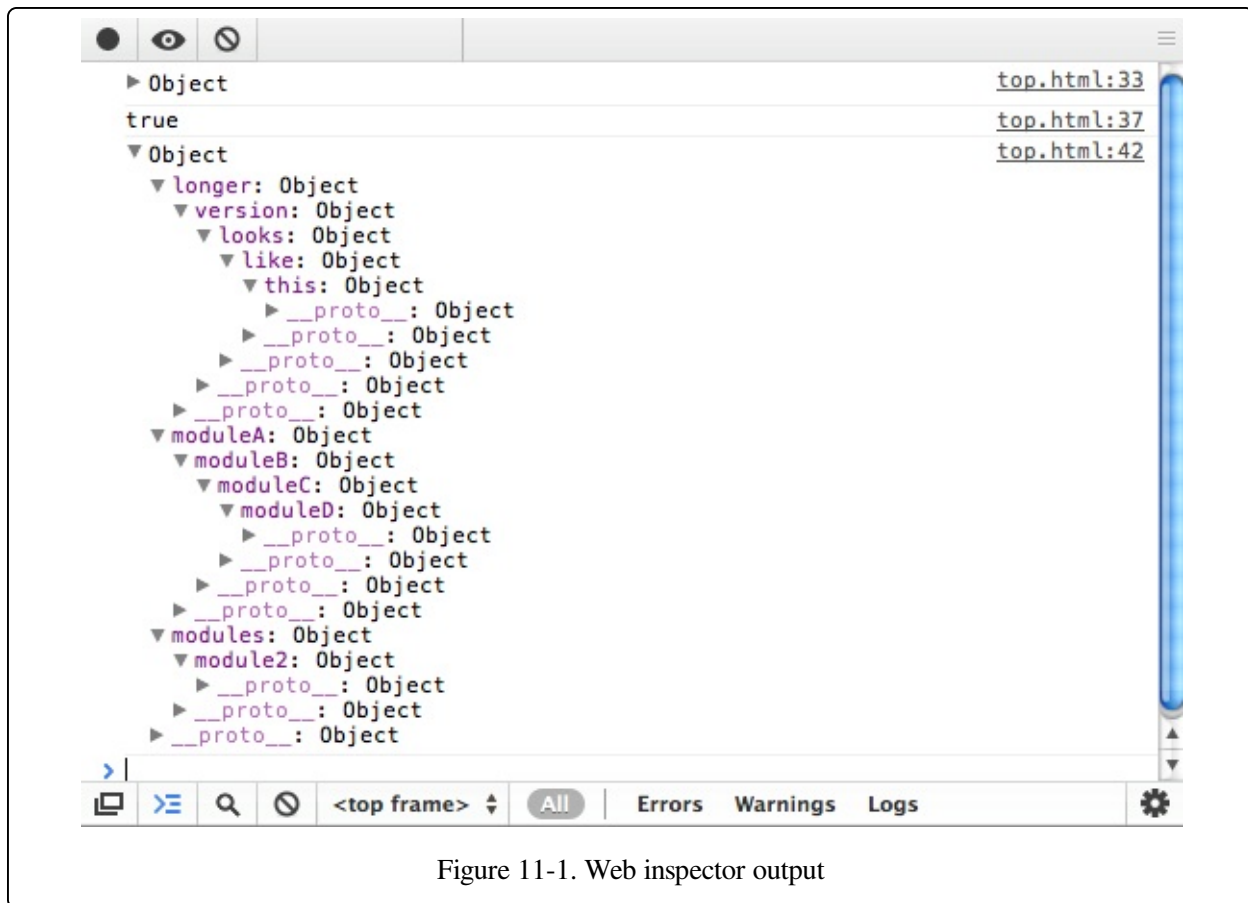


Figure 11-1. Web inspector output

Note how where one would previously have had to explicitly declare the various nests for their namespace as objects, this can now be easily achieved using a single, cleaner line of code. This works exceedingly well when defining purely namespaces alone, but can seem a little less flexible when you want to define both functions and properties at the same time as declaring your namespaces. Regardless, it is still incredibly powerful and I regularly use a similar approach in some of my projects.

Dependency declaration pattern

In this section we're going to take a look at a minor augmentation to the nested namespacing pattern you may be used to seeing in some applications. We all know that local references to objects can decrease overall lookup times, but let's apply this to namespacing to see how it might look in practice:

```
// common approach to accessing nested namespaces
myApp.utilities.math.fibonacci(25);
myApp.utilities.math.sin(56);
myApp.utilities.drawing.plot(98,50,60);
```

```
// with local/cached references
var utils = myApp.utilities,
    maths = utils.math,
    drawing = utils.drawing;

// easier to access the namespace
maths.fibonacci(25);
maths.sin(56);
drawing.plot(98, 50, 60);

// note that the above is particularly performant when
// compared to hundreds or thousands of calls to nested
// namespaces vs. a local reference to the namespace
```

Working with a local variable here is almost always faster than working with a top-level global (e.g. `myApp`). It's also both more convenient and more performant than accessing nested properties/sub-namespaces on every subsequent line and can improve readability in more complex applications.

Stoyan recommends declaring localized namespaces required by a function or module at the top of your function scope (using the single-variable pattern) and calls this a dependency declaration pattern. One of the benefits this offers is a decrease in locating dependencies and resolving them, should you have an extendable architecture that dynamically loads modules into your namespace when required.

In my opinion this pattern works best when working at a modular level, localizing a namespace to be used by a group of methods. Localizing namespaces on a per-function level, especially where there is significant overlap between namespace dependencies would be something I would recommend avoiding where possible. Instead, define it further up and just have them all access the same reference.

Deep object extension

An alternative approach to automatic namespacing is deep object extension. Namespaces defined using object literal notation may be easily extended (or merged) with other objects (or namespaces) such that the properties and functions of both namespaces can be accessible under the same namespace post-merge.

This is something that's been made fairly easy to accomplish with modern JavaScript frameworks (e.g. see jQuery's `$.extend`), however, if you're looking to extend object (namespaces) using vanilla JS, the following routine may be of

assistance.

```
// extend.js
// written by andrew dupont, optimized by addy osmani
function extend(destination, source) {
  var toString = Object.prototype.toString,
      objTest = toString.call({});
  for (var property in source) {
    if (source[property] && objTest == toString.call(source[property])) {
      destination[property] = destination[property] || {};
      extend(destination[property], source[property]);
    } else {
      destination[property] = source[property];
    }
  }
  return destination;
};

console.group("objExtend namespacing tests");

// define a top-level namespace for usage
var myNS = myNS || {};

// 1. extend namespace with a 'utils' object
extend(myNS, {
  utils:{
  }
});

console.log('test 1', myNS);
//myNS.utils now exists

// 2. extend with multiple depths (namespace.hello.world.wave)
extend(myNS, {
  hello:{
    world:{
      wave:{
        test: function(){
          //...
        }
      }
    }
  }
});

// test direct assignment works as expected
myNS.hello.test1 = 'this is a test';
myNS.hello.world.test2 = 'this is another test';
console.log('test 2', myNS);

// 3. what if myNS already contains the namespace being added
// (e.g. 'library')? we want to ensure no namespaces are being
// overwritten during extension
myNS.library = {
```

```

        foo:function(){}
    };

    extend(myNS, {
        library:{
            bar:function(){
                //...
            }
        }
    });

    // confirmed that extend is operating safely (as expected)
    // myNS now also contains library.foo, library.bar
    console.log('test 3', myNS);

    // 4. what if we wanted easier access to a specific namespace without having
    // to type the whole namespace out each time?.

    var shorterAccess1 = myNS.hello.world;
    shorterAccess1.test3 = "hello again";
    console.log('test 4', myNS);
    //success, myApp.hello.world.test3 is now 'hello again'

    console.groupEnd();

```

Note: The above implementation is not cross-browser compatible for all objects and should be considered a proof-of-concept only. You may find the Underscore.js `extend` method a simpler, more cross-browser implementation to start with <http://documentcloud.github.com/underscore/docs/underscore.html#section-67>. Alternatively, a version of the jQuery `$.extend()` method extracted from core can be found here: <https://gist.github.com/1123784>.

If you do however happen to be using jQuery in your application, you can achieve the exact same object namespace extensibility with `$.extend` as follows:

```

// top-level namespace
var myApp = myApp || {};

// directly assign a nested namespace
myApp.library = {
    foo:function(){ /*...*/ }
};

// deep extend/merge this namespace with another
// to make things interesting, let's say it's a namespace
// with the same name but with a different function
// signature: $.extend(deep, target, object1, object2)
$.extend(true, myApp, {
    library:{
        bar:function(){
            //...
        }
    }
});

```

```
    }  
  });  
  
  console.log('test', myApp);  
  // myApp now contains both library.foo() and library.bar() methods  
  // nothing has been overwritten which is what we're hoping for.
```

For the sake of thoroughness, please see [here](#) for jQuery \$.extend equivalents to the rest of the namespacing experiments found in this section.

Namespacing Fundamentals

Namespaces can be found in almost any serious JavaScript application. Unless you're working with a code-snippet, it's imperative that you do your best to ensure that you're implementing namespacing correctly as it's not just simple to pick-up, it'll also avoid third party code clobbering your own. The patterns we'll be examining in this section are:

1. Single global variables
2. Object literal notation
3. Nested namespacing
4. Immediately-invoked Function Expressions
5. Namespace injection

1. Single global variables

One popular pattern for namespacing in JavaScript is opting for a single global variable as your primary object of reference. A skeleton implementation of this where we return an object with functions and properties can be found below:

```
var myApplication = (function(){  
    function(){  
        //...  
    },  
    return{  
        //...  
    }  
})();
```

Although this works for certain situations, the biggest challenge with the single global variable pattern is ensuring that no one else has used the same global variable name as you have in the page.

2. Prefix namespacing

One solution to the above problem, as mentioned by [Peter Michaux](#), is to use prefix namespacing. It's a simple concept at heart, but the idea is you select a unique prefix namespace you wish to use (in this example, "myApplication_") and then define any methods, variables or other objects after the prefix as follows:

```
var myApplication_propertyA = {};
var myApplication_propertyB = {};
function myApplication_myMethod(){ /*...*/ }
```

This is effective from the perspective of trying to lower the chances of a particular variable existing in the global scope, but remember that a uniquely named object can have the same effect. This aside, the biggest issue with the pattern is that it can result in a large number of global objects once your application starts to grow. There is also quite a heavy reliance on your prefix not being used by any other developers in the global namespace, so be careful if opting to use this.

For more on Peter's views about the single global variable pattern, read his excellent post on them [here](#).

3. Object literal notation

Object literal notation (which we also cover in the module pattern section of the book) can be thought of as an object containing a collection of key:value pairs with a colon separating each pair of keys and values where keys can also represent new namespaces.

```
var myApplication = {
  getInfo:function(){ /**/ },

  // we can also populate object literal to support
  // further object namespaces containing anything really:
  models : {},
  views : {
    pages : {}
  },
  collections : {}
};
```

One can also opt for adding properties directly to the namespace:

```
myApplication.foo = function(){
  return "bar";
}
myApplication.utils = {
  toString:function(){
    //...
  },
  export: function(){
    //...
  }
}
```

Object literals have the advantage of not polluting the global namespace but assist

in organizing code and parameters logically. They're beneficial if you wish to create easily-readable structures that can be expanded to support deep nesting. Unlike simple global variables, object literals often also take into account tests for the existence of a variable by the same name so the chances of collision occurring are significantly reduced.

The code at the very top of the next sample demonstrates the different ways in which you can check to see if a variable (object namespace) already exists before defining it. You'll commonly see developers using Option 1, however Options 3 and 5 may be considered more thorough and Option 4 is considered a good best-practice.

```
// This doesn't check for existence of 'myApplication' in
// the global namespace. Bad practice as you can easily
// clobber an existing variable/namespace with the same name
var myApplication = {};

// The following options *do* check for variable/namespace existence.
// If already defined, we use that instance, otherwise we assign a new
// object literal to myApplication.
//
// Option 1: var myApplication = myApplication || {};
// Option 2: if(!MyApplication) MyApplication = {};
// Option 3: var myApplication = myApplication = myApplication || {}
// Option 4: myApplication || (myApplication = {});
// Option 5: var myApplication = myApplication === undefined ? {} : myApplication;
//
```

There is of course a huge amount of variance in how and where object literals are used for organizing and structuring code. For smaller applications wishing to expose a nested API for a particular self-enclosed module, you may just find yourself using the Revealing Module Pattern, which we covered earlier in the book:

```
var namespace = (function () {

    // defined within the local scope
    var privateMethod1 = function () { /* ... */ }
    var privateMethod2 = function () { /* ... */ }
    var privateProperty1 = 'foobar';

    return {
        // the object literal returned here can have as many
        // nested depths as you wish, however as mentioned,
        // this way of doing things works best for smaller,
        // limited-scope applications in my personal opinion
        publicMethod1: privateMethod1,

        // nested namespace with public properties
    }
})();
```

```

        properties:{
            publicProperty1: privateProperty1
        },

        // another tested namespace
        utils:{
            publicMethod2: privateMethod2
        }
        ...
    }
    })();

```

The benefit of object literals is that they offer us a very elegant key/value syntax to work with; one where we're able to easily encapsulate any distinct logic or functionality for our application in a way that clearly separates it from others and provides a solid foundation for extending your code.

A possible downside however is that object literals have the potential to grow into long syntactic constructs. Opting to take advantage of the nested namespace pattern (which also uses the same pattern as its base)

This pattern has a number of other useful applications too. In addition to namespacing, it's often of benefit to decouple the default configuration for your application into a single area that can be easily modified without the need to search through your entire codebase just to alter them - object literals work great for this purpose. Here's an example of a hypothetical object literal for configuration:

```

var myConfig = {
    language: 'english',
    defaults: {
        enableGeolocation: true,
        enableSharing: false,
        maxPhotos: 20
    },
    theme: {
        skin: 'a',
        toolbars: {
            index: 'ui-navigation-toolbar',
            pages: 'ui-custom-toolbar'
        }
    }
}

```

Note that there are really only minor syntactical differences between the object literal pattern and a standard JSON data set. If for any reason you wish to use JSON for storing your configurations instead (e.g. for simpler storage when sending to the back-end), feel free to. For more on the object literal pattern, I

recommend reading Rebecca Murphey's excellent [article](#) on the topic.

4. Nested namespacing

An extension of the object literal pattern is nested namespacing. It's another common pattern used that offers a lower risk of collision due to the fact that even if a namespace already exists, it's unlikely the same nested children do.

Does this look familiar?

```
YAHOO.util.Dom.getElementsByClassName('test');
```

Older versions of Yahoo!'s YUI library use the nested object namespacing pattern regularly. At AOL we also used this pattern in many of our larger applications. A sample implementation of nested namespacing may look like this:

```
var myApp = myApp || {};  
  
// perform a similar existence check when defining nested  
// children  
myApp.routers = myApp.routers || {};  
myApp.model = myApp.model || {};  
myApp.model.special = myApp.model.special || {};  
  
// nested namespaces can be as complex as required:  
// myApp.utilities.charting.html5.plotGraph(/*..*/);  
// myApp.modules.financePlanner.getSummary();  
// myApp.services.social.facebook.realtimeStream.getLatest();
```

Note: The above differs from how YUI3 approaches namespacing as modules there use a sandboxed API host object with far less and far shallower namespacing.

You can also opt to declare new nested namespaces/properties as indexed properties as follows:

```
myApp["routers"] = myApp["routers"] || {};  
myApp["models"] = myApp["models"] || {};  
myApp["controllers"] = myApp["controllers"] || {};
```

Both options are readable, organized and offer a relatively safe way of namespacing your application in a similar fashion to what you may be used to in other languages. The only real caveat however is that it requires your browser's JavaScript engine first locating the myApp object and then digging down until it gets to the function you actually wish to use.

This can mean an increased amount of work to perform lookups, however

developers such as [Juriy Zaytsev](#) have previously tested and found the performance differences between single object namespacing vs the 'nested' approach to be quite negligible.

5. Immediately-invoked Function Expressions (IIFE)s

An **IIFE** is effectively an unnamed function which is immediately invoked after it's been defined. In JavaScript, because both variables and functions explicitly defined within such a context may only be accessed inside of it, function invocation provides an easy means to achieving privacy.

This is one of the many reasons why IIFEs are a popular approach to encapsulating application logic to protect it from the global namespace. You've probably come across this pattern before under the name of a self-executing (or self-invoked) anonymous function, however I personally prefer Ben Alman's naming convention for this particular pattern as I believe it to be both more descriptive and more accurate.

The simplest version of an IIFE could be the following:

```
// an (anonymous) immediately-invoked function expression
(function(){ /*...*/})();
// a named immediately-invoked function expression
(function foobar(){ /*...*/})();
// this is technically a self-executing function which is quite different
function foobar(){ foobar(); }
```

whilst a slightly more expanded version of the first example might look like:

```
var namespace = namespace || {};

// here a namespace object is passed as a function
// parameter, where we assign public methods and
// properties to it
(function( o ){
    o.foo = "foo";
    o.bar = function(){
        return "bar";
    };
})(namespace);

console.log(namespace);
```

Whilst readable, this example could be significantly expanded on to address common development concerns such as defined levels of privacy (public/private

functions and variables) as well as convenient namespace extension. Let's go through some more code:

```
// namespace (our namespace name) and undefined are passed here
// to ensure 1. namespace can be modified locally and isn't
// overwritten outside of our function context
// 2. the value of undefined is guaranteed as being truly
// undefined. This is to avoid issues with undefined being
// mutable pre-ES5.

;(function ( namespace, undefined ) {
    // private properties
    var foo = "foo",
        bar = "bar";

    // public methods and properties
    namespace.foobar = "foobar";
    namespace.sayHello = function () {
        speak("hello world");
    };

    // private method
    function speak(msg) {
        console.log("You said: " + msg);
    };

    // check to evaluate whether 'namespace' exists in the
    // global namespace - if not, assign window.namespace an
    // object literal
})(window.namespace = window.namespace || {});

// we can then test our properties and methods as follows

// public
console.log(namespace.foobar); // foobar
namespace.sayHello(); // hello world

// assigning new properties
namespace.foobar2 = "foobar";
console.log(namespace.foobar2);
```

Extensibility is of course key to any scalable namespacing pattern and IIFEs can be used to achieve this quite easily. In the below example, our 'namespace' is once again passed as an argument to our anonymous function and is then extended (or decorated) with further functionality:

```
// let's extend the namespace with new functionality
(function( namespace, undefined ){
    // public method
    namespace.sayGoodbye = function(){
        console.log(namespace.foo);
        console.log(namespace.bar);
    };
});
```

```

        speak('goodbye');
    }
})( window.namespace = window.namespace || {});

```

namespace.sayGoodbye(); //goodbye

That's it for IIFEs for the time-being. If you would like to find out more about this pattern, I recommend reading both Ben's [IIFE post](#) and Elijah Manor's post on [namespace patterns from C#](#).

6. Namespace injection

Namespace injection is another variation on the IIFE where we 'inject' the methods and properties for a specific namespace from within a function wrapper using this as a namespace proxy. The benefit this pattern offers is easy application of functional behaviour to multiple objects or namespaces and can come in useful when applying a set of base methods to be built on later (e.g. getters and setters).

The disadvantages of this pattern are that there may be easier or more optimal approaches to achieving this goal (e.g. deep object extension / merging) which I cover earlier in the section.

Below we can see an example of this pattern in action, where we use it to populate the behaviour for two namespaces: one initially defined (utils) and another which we dynamically create as a part of the functionality assignment for utils (a new namespace called tools).

```

var myApp = myApp || {};
myApp.utils = {};

(function() {
    var val = 5;

    this.getValue = function() {
        return val;
    };

    this.setValue = function(newVal) {
        val = newVal;
    }

    // also introduce a new sub-namespace
    this.tools = {};

}).apply(myApp.utils);

// inject new behaviour into the tools namespace

```

```
// which we defined via the utilities module

(function(){
  this.diagnose = function(){
    return 'diagnosis';
  }
}).apply(myApp.utils.tools);

// note, this same approach to extension could be applied
// to a regular IIFE, by just passing in the context as
// an argument and modifying the context rather than just
// 'this'

// testing
console.log(myApp); //the now populated namespace
console.log(myApp.utils.getValue()); // test get
myApp.utils.setValue(25); // test set
console.log(myApp.utils.getValue());
console.log(myApp.utils.tools.diagnose());
```

Angus Croll has also [previously](#) suggested the idea of using the call API to provide a natural separation between contexts and arguments. This pattern can feel a lot more like a module creator, but as modules still offer an encapsulation solution, I'll briefly cover it for the sake of thoroughness:

```
// define a namespace we can use later
var ns = ns || {}, ns2 = ns2 || {};

// the module/namespace creator
var creator = function(val){
  var val = val || 0;

  this.next = function(){
    return val++;
  };

  this.reset = function(){
    val = 0;
  }
}

creator.call(ns);
// ns.next, ns.reset now exist
creator.call(ns2, 5000);
// ns2 contains the same methods
// but has an overridden value for val
// of 5000
```

As mentioned, this type of pattern is useful for assigning a similar base set of functionality to multiple modules or namespaces, but I'd really only suggest using it where explicitly declaring your functionality within an object/closure for direct access doesn't make sense.

Reviewing the namespace patterns above, the option that I would personally use for most larger applications is nested object namespacing with the object literal pattern.

IIFEs and single global variables may work fine for applications in the small to medium range, however, larger codebases requiring both namespaces and deep sub-namespaces require a succinct solution that promotes readability and scales. I feel this pattern achieves all of these objectives well.

I would also recommend trying out some of the suggested advanced utility methods for namespace extension as they really can save you time in the long-run.

Chapter 12. Design Patterns in jQuery Core

Now that we've taken a look at vanilla-JavaScript implementations of popular design patterns, let's switch gears and find out what of these design patterns might look like when implemented using jQuery. jQuery (as you may know) is currently the most popular JavaScript library and provides a layer of 'sugar' on top of regular JavaScript with a syntax that can be easier to understand at a glance.

Before we dive into this section, it's important to remember that many vanilla-JavaScript design patterns can be intermixed with jQuery when used correctly because jQuery is still essentially JavaScript itself.

jQuery is an interesting topic to discuss in the realm of patterns because the library actually uses a number of design patterns itself. What impresses me is just how cleanly all of the patterns it uses have been implemented so that they exist in harmony.

Let's take a look at what some of these patterns are and how they are used.

Module Pattern

We have already explored the module pattern previously, but in case you've skipped ahead: the **Module Pattern** allows us to encapsulate logic for a unit of code such that we can have both private and public methods and variables. This can be applied to writing jQuery plugins too, where a private API holds any code we don't wish to expose and a public API contains anything a user will be allowed to interact with. See below for an example:

```
!function(exports, $, undefined){  
    var Plugin = function(){  
        // Our private API  
        var priv = {},  
        // Our public API  
        Plugin = {},  
        // Plugin defaults
```

```

        defaults = {};

        // Private options and methods
        priv.options = {};
        priv.method1 = function(){};
        priv.method2 = function(){};

        // Public methods
        Plugin.method1 = function(){...};
        Plugin.method2 = function(){...};

        // Public initialization
        Plugin.init = function(options) {
            $.extend(priv.options, defaults, options);
            priv.method1();
            return Plugin;
        }

        // Return the Public API (Plugin) we want
        // to expose
        return Plugin;
    }

    exports.Plugin = Plugin;

})(this, jQuery);

```

This can then be used as follows:

```

var myPlugin = new Plugin;
myPlugin.init(/* custom options */);
myPlugin.method1();

```

Lazy Initialization

Lazy Initialization is a design pattern which allows us to delay expensive processes (e.g. the creation of objects) until the first instance they are needed. An example of this is the `.ready()` function in jQuery that only executes a function once the DOM is ready.

```
$(document).ready(function(){
    // The ajax request won't attempt to execute until
    // the DOM is ready

    var jqxhr = $.ajax({
        url: 'http://domain.com/api/',
        data: 'display=latest&order=ascending'
    })
    .done(function( data ){
        $('.status').html('content loaded');
        console.log( 'Data output:' + data );
    });
});
```

Whilst it isn't directly used in jQuery core, some developers will be familiar with the concept of LazyLoading via plugins such as [this](#). LazyLoading is effectively the same as Lazy initialization and is a technique whereby additional data on a page is loaded when needed (e.g. when a user has scrolled to the end of the page). In recent years this pattern has become quite prominent and can be currently be found in both the Twitter and Facebook UIs.

The Composite Pattern

The Composite Pattern describes a group of objects that can be treated in the same way a single instance of an object may be. Implementing this pattern allows you to treat both individual objects and compositions in a uniform manner. In jQuery, when we're accessing or performing actions on a single DOM element or a collection of elements, we can treat both sets in a uniform manner. This is demonstrated by the code sample below:

```
// Single elements
$('#singleItem').addClass('active');
$('#container').addClass('active');

// Collections of elements
$('div').addClass('active');
$('.item').addClass('active');
$('input').addClass('active');
```

The Adapter Pattern

The Adapter Pattern is a pattern which translates an interface for a class into an interface compatible with a specific system. Adapters basically allow classes to function together which normally couldn't due to their incompatible interfaces. The adapter translates calls to its interface into calls to the original interface and the code required to achieve this is usually quite minimal.

One example of an adapter you may have used is jQuery's `$(el).css()` method. Not only does it help normalize the interfaces to how styles can be applied between a number of browsers, there are plenty of good examples of this, including opacity.

```
//  
// Cross browser opacity:  
// opacity: 0.9; Chrome 4+, FF2+, Saf3.1+, Opera 9+, IE9, iOS 3.2+, Android 2.1+  
// filter: alpha(opacity=90); IE6-IE8  
//  
$('.container').css({ opacity: .5 });
```

The Facade Pattern

As we saw in earlier sections, the **Facade Pattern** is where an object provides a simpler interface to a larger (possibly more complex) body of code. Facades can be frequently found across the jQuery library and make methods both easier to use and understand, but also more readable. The following are facades for jQuery's `$.ajax()`:

```
$.get( url, data, callback, dataType );
$.post( url, data, callback, dataType );
$.getJSON( url, data, callback );
$.getScript( url, callback );
```

These are translated behind the scenes to:

```
// $.get()
$.ajax({
  url: url,
  data: data,
  dataType: dataType
}).done( callback );

// $.post
$.ajax({
  type: 'POST',
  url: url,
  data: data,
  dataType: dataType
}).done( callback );

// $.getJSON()
$.ajax({
  url: url,
  dataType: 'json',
  data: data,
}).done( callback );

// $.getScript()
$.ajax({
  url: url,
  dataType: "script",
}).done( callback );
```

What's even more interesting is that the above facades are actually facades in their own right. You see, `$.ajax` offers a much simpler interface to a complex body of code that handles cross-browser XHR (XMLHttpRequest) as well as **deferreds**.

While I could link you to the jQuery source, here's a **cross-browser XHR implementation** just so you can get an idea of how much easier this pattern makes

our lives.

The Observer Pattern

Another pattern we've look at previously is the Observer (Publish/Subscribe) pattern, where a subject (the publisher), keeps a list of its dependents (subscribers), and notifies them automatically anytime something interesting happens.

jQuery actually comes with built-in support for a publish/subscribe-like system, which it calls custom events. In earlier versions of the library, access to these custom events was possible using `.bind()` (subscribe), `.trigger()` (publish) and `.unbind()` (unsubscribe), but in recent versions this can be done using `.on()`, `.trigger()` and `.off()`.

Below we can see an example of this being used in practice:

```
// Equivalent to subscribe(topicName, callback)
$(document).on('topicName', function(){
    //..perform some behaviour
});

// Equivalent to publish(topicName)
$(document).trigger('topicName');

// Equivalent to unsubscribe(topicName)
$(document).off('topicName');
```

For those that prefer to use the conventional naming scheme for the Observer pattern, **Ben Alman** created a simple wrapper around the above methods which gives you access to `$.publish()`, `$.subscribe`, and `$.unsubscribe` methods. I've previously linked to them earlier in the book, but you can see the wrapper in full below.

```
(function($) {

    var o = $({});

    $.subscribe = function() {
        o.on.apply(o, arguments);
    };

    $.unsubscribe = function() {
        o.off.apply(o, arguments);
    };

    $.publish = function() {
        o.trigger.apply(o, arguments);
    };

});
```

```
}(jQuery));
```

Finally, in recent versions of jQuery, a multi-purpose callbacks object (\$.Callbacks) was made available to enable users to write new solutions based on callback lists. One such solution to write using this feature is another Publish/Subscribe system. An implementation of this is the following:

```
var topics = {};  
  
jQuery.Topic = function( id ) {  
    var callbacks,  
        topic = id && topics[ id ];  
    if ( !topic ) {  
        callbacks = jQuery.Callbacks();  
        topic = {  
            publish: callbacks.fire,  
            subscribe: callbacks.add,  
            unsubscribe: callbacks.remove  
        };  
        if ( id ) {  
            topics[ id ] = topic;  
        }  
    }  
    return topic;  
};
```

which can then be used as follows:

```
// Subscribers  
$.Topic( 'mailArrived' ).subscribe( fn1 );  
$.Topic( 'mailArrived' ).subscribe( fn2 );  
$.Topic( 'mailSent' ).subscribe( fn1 );  
  
// Publisher  
$.Topic( 'mailArrived' ).publish( 'hello world!' );  
$.Topic( 'mailSent' ).publish( 'woo! mail!' );  
  
// Here, 'hello world!' gets pushed to fn1 and fn2  
// when the 'mailArrived' notification is published  
// with 'woo! mail!' also being pushed to fn1 when  
// the 'mailSent' notification is published.  
  
//  
// output:  
// hello world!  
// fn2 says: hello world!  
// woo! mail!  
//
```

The Iterator Pattern

The Iterator Pattern is a design pattern where iterators (objects that allow us to traverse through all the elements of a collection) access the elements of an aggregate object sequentially without needing to expose its underlying form.

Iterators encapsulate the internal structure of how that particular iteration occurs - in the case of jQuery's `$(el).each()` iterator, you are actually able to use the underlying code behind `$.each()` to iterate through a collection, without needing to see or understand the code working behind the scenes that's providing this capability. This is a pattern similar to the facade, except it deals explicitly with iteration.

```
$.each(['john','dave','rick','julian'], function(index, value) {  
    console.log(index + ': ' + value);  
});  
  
$('li').each(function(index) {  
    console.log(index + ': ' + $(this).text());  
});
```

The Strategy Pattern

The Strategy Pattern is a pattern where a script may select a particular algorithm at runtime. The purpose of this pattern is that it's able to provide a way to clearly define families of algorithms, encapsulate each as an object and make them easily interchangeable. You could say that the biggest benefit this pattern offers is that it allows algorithms to vary independent of the clients that utilize them.

An example of this is where jQuery's `toggle()` allows you to bind two or more handlers to the matched elements, to be executed on alternate clicks. The strategy pattern allows for alternative algorithms to be used independent of the client internal to the function.

```
$('#button').toggle(function(){  
    console.log('path 1');  
},  
function(){  
    console.log('path 2');  
});
```


The Proxy Pattern

The Proxy Pattern - a proxy is basically a class that functions as an interface to something else: a file, a resource, an object in memory, something else that is difficult to duplicate, etc. jQuery's `.proxy()` method takes as input a function and returns a new one that will always have a particular context - it ensures that the value of `this` in a function is the value you desire. This is parallel to the idea of providing an interface as per the proxy pattern.

One example of where this is useful is when you're making use of a timer inside a click handler. Say we have the following handler:

```
$('#button').on('click', function(){
  // Within this function, 'this' refers to the element that was clicked
  $(this).addClass('active');
});
```

However, say we wished to add in a delay before the `active` class was added. One thought that comes to mind is using `setTimeout` to achieve this, but there's a slight problem here: whatever function is passed to `setTimeout` will have a different value for `this` inside that function (it will refer to `window` instead).

```
$('#button').on('click', function(){
  setTimeout(function(){
    // 'this' doesn't refer to our element!
    $(this).addClass('active');
  });
});
```

To solve this problem, we can use `$.proxy()`. By calling it with the function and value we would like assigned to `this` it will actually return a function that retains the value we desire. Here's how this would look:

```
$('#button').on('click', function(){
  setTimeout($.proxy(function() {
    // 'this' now refers to our element as we wanted
    $(this).addClass('active');
  }, this), 500);
  // the last 'this' we're passing tells $.proxy() that our DOM element
  // is the value we want 'this' to refer to.
});
```

The Builder Pattern

The Builder Pattern's general idea is that it abstracts the steps involved in creating objects so that different implementations of these steps have the ability to construct different representations of objects. Below are examples of how jQuery utilizes this pattern to allow you to dynamically create new elements.

```
$('<div class= "foo">bar</div>');  
  
$('<p id="test">foo <em>bar</em></p>').appendTo('body');  
  
var newParagraph = $('<p />').text("Hello world");  
  
$('<input />').attr({'type': 'text', 'id': 'sample'})  
    .appendTo('#container');
```

The Prototype Pattern

As we've seen, the **Prototype Pattern** is used when objects are created based on a template of an existing object through cloning. Essentially this pattern is used to avoid creating a new object in a more conventional manner where this process may be expensive or overly complex.

In terms of the jQuery library, your first thought when cloning is mentioned might be the `.clone()` method. Unfortunately this only clones DOM elements but if we want to clone JavaScript objects, this can be done using the `$.extend()` method as follows:

```
var myOldObject = {};  
  
// Create a shallow copy  
var myNewObject = jQuery.extend({}, myOldObject);  
  
// Create a deep copy  
var myOtherNewObject = jQuery.extend(true, {}, myOldObject);
```

This pattern has been used many times in jQuery core (as well as in jQuery plugins) quite successfully. For those wondering what deep cloning might look like in JavaScript without the use of a library, **Rick Waldron** has an implementation you can use below (and tests available [here](#)).

```
function clone( obj ) {  
    var val, length, i,  
        temp = [];  
  
    if ( Array.isArray(obj) ) {  
        for ( i = 0, length = obj.length; i < length; i++ ) {  
            // Store reference to this array item's value  
            val = obj[ i ];  
  
            // If array item is an object (including arrays), derive new value by cloning  
            if ( typeof val === "object" ) {  
                val = clone( val );  
            }  
            temp[ i ] = val;  
        }  
        return temp;  
    }  
  
    // Create a new object whose prototype is a new, empty object,  
    // Using the second properties object argument to copy the source properties  
    return Object.create({}, (function( src ) {  
        // Initialize a cache for non-inherited properties  
        var props = {};
```

```
Object.getOwnPropertyNames( src ).forEach(function( name ) {  
  // Store short reference to property descriptor  
  var descriptor = Object.getOwnPropertyDescriptor( src, name );  
  
  // Recurse on properties whose value is an object or array  
  if ( typeof src[ name ] === "object" ) {  
    descriptor.value = clone( src[ name ] );  
  }  
  props[ name ] = descriptor;  
});  
return props;  
})( obj ));  
}
```

Chapter 13. Modern Modular JavaScript Design Patterns

The Importance Of Decoupling Your Application

In the world of modern JavaScript, when we say an application is **modular**, we often mean it's composed of a set of highly decoupled, distinct pieces of functionality stored in modules. As you probably know, **loose coupling** facilitates easier maintainability of apps by removing dependencies where possible. When this is implemented efficiently, it's quite easy to see how changes to one part of a system may affect another.

Unlike some more traditional programming languages however, the current iteration of JavaScript (**ECMA-262**) doesn't provide developers with the means to import such modules of code in a clean, organized manner. It's one of the concerns with specifications that haven't required great thought until more recent years where the need for more organized JavaScript applications became apparent.

Instead, developers at present are left to fall back on variations of the **module** or **object literal** patterns, which we covered earlier in the book. With many of these, module scripts are strung together in the DOM with namespaces being described by a single global object where it's still possible to incur naming collisions in your architecture. There's also no clean way to handle dependency management without some manual effort or third party tools.

Whilst native solutions to these problems will be arriving in **ES Harmony** (the next version of JavaScript), the good news is that writing modular JavaScript has never been easier and you can start doing it today.

In this section, we're going to look at three formats for writing modular JavaScript: **AMD**, **CommonJS** and proposals for the next version of JavaScript, **Harmony**.

A Note On Script Loaders

It's difficult to discuss AMD and CommonJS modules without talking about the elephant in the room - **script loaders**. At the time of writing, script loading is a means to a goal, that goal being modular JavaScript that can be used in applications today - for this, use of a compatible script loader is unfortunately necessary. In order to get the most out of this section, I recommend gaining a **basic understanding** of how popular script loading tools work so the explanations of module formats make sense in context.

There are a number of great loaders for handling module loading in the AMD and CommonJS formats, but my personal preferences are **RequireJS** and **curl.js**. Complete tutorials on these tools are outside the scope of this book, but I can recommend reading John Hann's article about **curl.js** and James Burke's **RequireJS** API documentation for more.

From a production perspective, the use of optimization tools (like the RequireJS optimizer) to concatenate scripts is recommended for deployment when working with such modules. Interestingly, with the **Almond** AMD shim, RequireJS doesn't need to be rolled in the deployed site and what you might consider a script loader can be easily shifted outside of development.

That said, James Burke would probably say that being able to dynamically load scripts after page load still has its use cases and RequireJS can assist with this too. With these notes in mind, let's get started.

AMD

A Format For Writing Modular JavaScript In The Browser

The overall goal for the AMD (Asynchronous Module Definition) format is to provide a solution for modular JavaScript that developers can use today. It was born out of Dojo's real world experience using XHR+eval and proponents of this format wanted to avoid any future solutions suffering from the weaknesses of those in the past.

The AMD module format itself is a proposal for defining modules where both the module and dependencies can be **asynchronously** loaded. It has a number of distinct advantages including being both asynchronous and highly flexible by nature which removes the tight coupling one might commonly find between code and module identity. Many developers enjoy using it and one could consider it a reliable stepping stone towards the **module system** proposed for ES Harmony.

AMD began as a draft specification for a module format on the CommonJS list but as it wasn't able to reach full consensus, further development of the format moved to the **amdjs** group.

Today it's embraced by projects including Dojo (1.7), MooTools (2.0), Firebug (1.8) and even jQuery (1.7). Although the term CommonJS AMD format has been seen in the wild on occasion, it's best to refer to it as just AMD or Async Module support as not all participants on the CommonJS list wished to pursue it.

Note: There was a time when the proposal was referred to as Modules Transport/C, however as the spec wasn't geared for transporting existing CommonJS modules, but rather, for defining modules it made more sense to opt for the AMD naming convention.

Getting Started With Modules

The two key concepts you need to be aware of here are the idea of a **define** method for facilitating module definition and a **require** method for handling dependency loading. **define** is used to define named or unnamed modules based on

the proposal using the following signature:

```
define(  
  module_id /*optional*/,  
  [dependencies] /*optional*/,  
  definition function /*function for instantiating the module or object*/  
);
```

As you can tell by the inline comments, the `module_id` is an optional argument which is typically only required when non-AMD concatenation tools are being used (there may be some other edge cases where it's useful too). When this argument is left out, we call the module anonymous.

When working with anonymous modules, the idea of a module's identity is DRY, making it trivial to avoid duplication of filenames and code. Because the code is more portable, it can be easily moved to other locations (or around the file-system) without needing to alter the code itself or change its ID. The `module_id` is equivalent to folder paths in simple packages and when not used in packages. Developers can also run the same code on multiple environments just by using an AMD optimizer that works with a CommonJS environment such as [r.js](#).

Back to the `define` signature, the `dependencies` argument represents an array of dependencies which are required by the module you are defining and the third argument ('definition function' or 'factory function') is a function that's executed to instantiate your module. A barebone module could be defined as follows:

Understanding AMD: `define()`

```
// A module_id (myModule) is used here for demonstration purposes only  
  
define('myModule',  
  ['foo', 'bar'],  
  // module definition function  
  // dependencies (foo and bar) are mapped to function parameters  
  function ( foo, bar ) {  
    // return a value that defines the module export  
    // (i.e the functionality we want to expose for consumption)  
  
    // create your module here  
    var myModule = {  
      doStuff:function(){  
        console.log('Yay! Stuff');  
      }  
    }  
  
    return myModule;  
  });
```



```
// An alternative example could be..
define('myModule',
  ['math', 'graph'],
  function ( math, graph ) {

    // Note that this is a slightly different pattern
    // With AMD, it's possible to define modules in a few
    // different ways due to it's flexibility with
    // certain aspects of the syntax
    return {
      plot: function(x, y){
        return graph.drawPie(math.randomGrid(x,y));
      }
    };
  });
```

require on the other hand is typically used to load code in a top-level JavaScript file or within a module should you wish to dynamically fetch dependencies. An example of its usage is:

Understanding AMD: require()

```
// Consider 'foo' and 'bar' are two external modules
// In this example, the 'exports' from the two modules loaded are passed as
// function arguments to the callback (foo and bar)
// so that they can similarly be accessed

require(['foo', 'bar'], function ( foo, bar ) {
  // rest of your code here
  foo.doSomething();
});
```

Dynamically-loaded Dependencies

```
define(function ( require ) {
  var isReady = false, foobar;

  // note the inline require within our module definition
  require(['foo', 'bar'], function (foo, bar) {
    isReady = true;
    foobar = foo() + bar();
  });

  // we can still return a module
  return {
    isReady: isReady,
    foobar: foobar
  };
});
```

Understanding AMD: plugins

The following is an example of defining an AMD-compatible plugin:

```
// With AMD, it's possible to load in assets of almost any kind
// including text-files and HTML. This enables us to have template
// dependencies which can be used to skin components either on
// page-load or dynamically.

define(['./templates', 'text!./template.md', 'css!./template.css'],
function( templates, template ){
    console.log(templates);
    // do some fun template stuff here.
});
```

Note: Although `css!` is included for loading CSS dependencies in the above example, it's important to remember that this approach has some caveats such as it not being fully possible to establish when the CSS is fully loaded. Depending on how you approach your build, it may also result in CSS being included as a dependency in the optimized file, so use CSS as a loaded dependency in such cases with caution.

Loading AMD Modules Using RequireJS

```
require(['app/myModule'],
function( myModule ){
    // start the main module which in-turn
    // loads other modules
    var module = new myModule();
    module.doStuff();
});
```

This example could simply be looked at as `requirejs(['app/myModule'], function(){})` which indicates the loader's top level globals are being used. This is how to kick off top-level loading of modules with different AMD loaders however with a `define()` function, if it's passed a local `require` all `require([])` examples apply to both types of loader (`curl.js` and `RequireJS`).

Loading AMD Modules Using curl.js

```
curl(['app/myModule.js'],
function( myModule ){
    // start the main module which in-turn
    // loads other modules
    var module = new myModule();
    module.doStuff();
});
```

Modules With Deferred Dependencies

```
// This could be compatible with jQuery's Deferred implementation,
// futures.js (slightly different syntax) or any one of a number
// of other implementations
```

```
define(['lib/Deferred'], function( Deferred ){
    var defer = new Deferred();
    require(['lib/templates/?index.html','lib/data/?stats'],
        function( template, data ){
            defer.resolve({ template: template, data:data });
        }
    );
    return defer.promise();
});
```

Why Is AMD A Better Choice For Writing Modular JavaScript?

- Provides a clear proposal for how to approach defining flexible modules.
- Significantly cleaner than the present global namespace and `<script>` tag solutions many of us rely on. There's a clean way to declare stand-alone modules and dependencies they may have.
- Module definitions are encapsulated, helping us to avoid pollution of the global namespace.
- Works better than some alternative solutions (e.g. CommonJS, which we'll be looking at shortly). Doesn't have issues with cross-domain, local or debugging and doesn't have a reliance on server-side tools to be used. Most AMD loaders support loading modules in the browser without a build process.
- Provides a 'transport' approach for including multiple modules in a single file. Other approaches like CommonJS have yet to agree on a transport format.
- It's possible to lazy load scripts if this is needed.

Note: Many of the above could be said about YUI's module loading strategy as well.

Related Reading

[The RequireJS Guide To AMD](#)

[What's the fastest way to load AMD modules?](#)

[AMD vs. CommonJS, what's the better format?](#)

[AMD Is Better For The Web Than CommonJS Modules](#)

[The Future Is Modules Not Frameworks](#)

[AMD No Longer A CommonJS Specification](#)

On Inventing JavaScript Module Formats And Script Loaders

The AMD Mailing List

AMD Modules With Dojo

Defining AMD-compatible modules using Dojo is fairly straight-forward. As per above, define any module dependencies in an array as the first argument and provide a callback (factory) which will execute the module once the dependencies have been loaded. e.g:

```
define(["dijit/Tooltip"], function( Tooltip ){
    //Our dijit tooltip is now available for local use
    new Tooltip(...);
});
```

Note the anonymous nature of the module which can now be both consumed by a Dojo asynchronous loader, RequireJS or the standard `dojo.require()` module loader that you may be used to using.

For those wondering about module referencing, there are some interesting gotchas that are useful to know here. Although the AMD-advocated way of referencing modules declares them in the dependency list with a set of matching arguments, this isn't supported by the Dojo 1.6 build system - it really only works for AMD-compliant loaders. e.g:

```
define(["dojo/cookie", "dijit/Tooltip"], function( cookie, Tooltip ){
    var cookieValue = cookie("cookieName");
    new Tooltip(...);
});
```

This has many advances over nested namespacing as modules no longer need to directly reference complete namespaces every time - all we require is the 'dojo/cookie' path in dependencies, which once aliased to an argument, can be referenced by that variable. This removes the need to repeatedly type out 'dojo.' in your applications.

Note: Although Dojo 1.6 doesn't officially support user-based AMD modules (nor asynchronous loading), it's possible to get this working with Dojo using a number of different script loaders. At present, all Dojo core and Dijit modules have been transformed to the AMD syntax and improved overall AMD support will likely land between 1.7 and 2.0.

The final gotcha to be aware of is that if you wish to continue using the Dojo build system or wish to migrate older modules to this newer AMD-style, the following more verbose version enables easier migration. Notice that dojo and dijit and referenced as dependencies too:

```
define(["dojo", "dijit", "dojo/cookie", "dijit/Tooltip"], function(dojo, dijit){
    var cookieValue = dojo.cookie("cookieName");
    new dijit.Tooltip(...);
});
```

AMD Module Design Patterns (Dojo)

As we've seen in previous sections, design patterns can be highly effective in improving how we approach structuring solutions to common development problems. [John Hann](#) has given some excellent presentations about AMD module design patterns covering the Singleton, Decorator, Mediator and others and I highly recommend checking out his [slides](#) if you get a chance.

A selection of AMD design patterns can be found below.

Decorator pattern:

```
// mylib/UpdatableObservable: a decorator for dojo/store/Observable
define(['dojo', 'dojo/store/Observable'], function ( dojo, Observable ) {
    return function UpdatableObservable ( store ) {

        var observable = dojo.isFunction(store.notify) ? store :
            new Observable(store);

        observable.updated = function( object ) {
            dojo.when(object, function ( itemOrArray ) {
                dojo.forEach( [].concat(itemOrArray), this.notify, this );
            });
        };

        return observable; // makes `new` optional
    };
});

// decorator consumer
// a consumer for mylib/UpdatableObservable

define(['mylib/UpdatableObservable'], function ( makeUpdatable ) {
    var observable, updatable, someItem;
    // ... here be code to get or create `observable`

    // ... make the observable store updatable
    updatable = makeUpdatable(observable); // `new` is optional!
});
```

```

    // ... later, when a cometd message arrives with new data item
    updatable.updated(updatedItem);
  });

```

Adapter pattern

```

// 'mylib/Array' adapts `each` function to mimic jQuery's:
define(['dojo/_base/lang', 'dojo/_base/array'], function (lang, array) {
  return lang.delegate(array, {
    each: function (arr, lambda) {
      array.forEach(arr, function (item, i) {
        lambda.call(item, i, item); // like jQuery's each
      })
    }
  });
});

// adapter consumer
// 'myapp/my-module':
define(['mylib/Array'], function ( array ) {
  array.each(['uno', 'dos', 'tres'], function (i, esp) {
    // here, `this` == item
  });
});

```

AMD Modules With jQuery

The Basics

Unlike Dojo, jQuery really only comes with one file, however given the plugin-based nature of the library, we can demonstrate how straight-forward it is to define an AMD module that uses it below.

```

define(['js/jquery.js', 'js/jquery.color.js', 'js/underscore.js'],
  function($, colorPlugin, _){
    // Here we've passed in jQuery, the color plugin and Underscore
    // None of these will be accessible in the global scope, but we
    // can easily reference them below.

    // Pseudo-randomize an array of colors, selecting the first
    // item in the shuffled array
    var shuffleColor = _.first(_.shuffle(['#666', '#333', '#111']));

    // Animate the background-color of any elements with the class
    // 'item' on the page using the shuffled color
    $('item').animate({'backgroundColor': shuffleColor });

    return {};
    // What we return can be used by other modules
  });

```

There is however something missing from this example and it's the concept of registration.

Registering jQuery As An Async-compatible Module

One of the key features that landed in jQuery 1.7 was support for registering jQuery as an asynchronous module. There are a number of compatible script loaders (including RequireJS and curl) which are capable of loading modules using an asynchronous module format and this means fewer hacks are required to get things working.

If a developer wants to use AMD and does not want their jQuery version leaking into the global space, they should call `noConflict` in their top level module that uses jQuery. In addition, since multiple versions of jQuery can be on a page there are special considerations that an AMD loader must account for, and so jQuery only registers with AMD loaders that have recognized these concerns, which are indicated by the loader specifying `define.amd.jquery`. RequireJS and curl are two loaders that do so

The named AMD provides a safety blanket of being both robust and safe for most use-cases.

```
// Account for the existence of more than one global
// instances of jQuery in the document, cater for testing
// .noConflict()

var jQuery = this.jQuery || "jQuery",
    $ = this.$ || "$",
    originaljQuery = jQuery,
    original$ = $;

define(['jquery'], function ($) {
    $('items').css('background', 'green');
    return function () {};
});
```

Smarter jQuery Plugins

I've recently discussed some ideas and examples of how jQuery plugins could be written using Universal Module Definition (UMD) patterns [here](#). UMDs define modules that can work on both the client and server, as well as with all popular script loaders available at the moment. Whilst this is still a new area with a lot of concepts still being finalized, feel free to look at the code samples in the section title AMD & CommonJS below and let me know if you feel there's anything we could do better.

What Script Loaders & Frameworks Support AMD?

In-browser:

- RequireJS <http://requirejs.org>
- curl.js <http://github.com/unscriptable/curl>
- bdLoad <http://bdframework.com/bdLoad>
- Yabble <http://github.com/jbrantly/yabble>
- PINF <http://github.com/pinf/loader-js>
- (and more)

Server-side:

- RequireJS <http://requirejs.org>
- PINF <http://github.com/pinf/loader-js>

AMD Conclusions

The above are very trivial examples of just how useful AMD modules can truly be, but they hopefully provide a foundation for understanding how they work.

You may be interested to know that many visible large applications and companies currently use AMD modules as a part of their architecture. These include **IBM** and the **BBC iPlayer**, which highlight just how seriously this format is being considered by developers at an enterprise-level.

For more reasons why many developers are opting to use AMD modules in their applications, you may be interested in [this](#) post by James Burke.

CommonJS

A Module Format Optimized For The Server

CommonJS are a volunteer working group which aim to design, prototype and standardize JavaScript APIs. To date they've attempted to ratify standards for both **modules** and **packages**. The CommonJS module proposal specifies a simple API for declaring modules server-side and unlike AMD attempts to cover a broader set of concerns such as io, filesystem, promises and more.

Getting Started

From a structure perspective, a CommonJS module is a reusable piece of JavaScript which exports specific objects made available to any dependent code - there are typically no function wrappers around such modules (so you won't see `define` used here for example).

At a high-level they basically contain two primary parts: a free variable named `exports` which contains the objects a module wishes to make available to other modules and a `require` function that modules can use to import the exports of other modules.

Understanding CommonJS: `require()` and `exports`

```
// package/lib is a dependency we require
var lib = require('package/lib');

// some behaviour for our module
function foo(){
  lib.log('hello world!');
}

// export (expose) foo to other modules
exports.foo = foo;
```

Basic consumption of exports

```
// define more behaviour we would like to expose
function foobar(){
  this.foo = function(){
    console.log('Hello foo');
  }

  this.bar = function(){
    console.log('Hello bar');
  }
}
```

```

    }
}

// expose foobar to other modules
exports.foobar = foobar;

// an application consuming 'foobar'

// access the module relative to the path
// where both usage and module files exist
// in the same directory

var foobar = require('./foobar').foobar,
    test = new foobar();

test.bar(); // 'Hello bar'

```

AMD-equivalent Of The First CommonJS Example

```

define(function(require){
    var lib = require('package/lib');

    // some behaviour for our module
    function foo(){
        lib.log('hello world!');
    }

    // export (expose) foo for other modules
    return {
        foobar: foo
    };
});

```

This can be done as AMD supports a **simplified CommonJS wrapping** feature.

Consuming Multiple Dependencies

app.js:

```

var modA = require('./foo');
var modB = require('./bar');

exports.app = function(){
    console.log('Im an application!');
}

exports.foo = function(){
    return modA.helloWorld();
}

```

bar.js:

```

exports.name = 'bar';

```

foo.js:

```
require('./bar');
exports.helloWorld = function(){
  return 'Hello World!!'
}
```

What Loaders & Frameworks Support CommonJS?

In-browser:

- curl.js <http://github.com/unscriptable/curl>
- SproutCore 1.1 <http://sproutcore.com>
- PINF <http://github.com/pinf/loader-js>
- (and more)

Server-side:

- Node <http://nodejs.org>
- Narwhal <https://github.com/tlrobinson/narwhal>
- Persevere <http://www.persvr.org/>
- Wakanda <http://www.wakandasoft.com/>

Is CommonJS Suitable For The Browser?

There are developers that feel CommonJS is better suited to server-side development which is one reason there's currently a level of **disagreement** over which format should and will be used as the de facto standard in the pre-Harmony age moving forward. Some of the arguments against CommonJS include a note that many CommonJS APIs address server-oriented features which one would simply not be able to implement at a browser-level in JavaScript - for example, io, system and js could be considered unimplementable by the nature of their functionality.

That said, it's useful to know how to structure CommonJS modules regardless so that we can better appreciate how they fit in when defining modules which may be used everywhere. Modules which have applications on both the client and server include validation, conversion and templating engines. The way some developers are approaching choosing which format to use is opting for CommonJS when a module can be used in a server-side environment and using AMD if this is not the

case.

As AMD modules are capable of using plugins and can define more granular things like constructors and functions this makes sense. CommonJS modules are only able to define objects which can be tedious to work with if you're trying to obtain constructors out of them.

Although it's beyond the scope of this section, you may have also noticed that there were different types of 'require' methods mentioned when discussing AMD and CommonJS.

The concern with a similar naming convention is of course confusion and the community are currently split on the merits of a global require function. John Hann's suggestion here is that rather than calling it 'require', which would probably fail to achieve the goal of informing users about the difference between a global and inner require, it may make more sense to rename the global loader method something else (e.g. the name of the library). It's for this reason that a loader like curl.js uses `curl()` as opposed to `require`.

Related Reading

[Demystifying CommonJS Modules](#)

[JavaScript Growing Up](#)

[The RequireJS Notes On CommonJS](#)

[Taking Baby Steps With Node.js And CommonJS - Creating Custom Modules](#)

[Asynchronous CommonJS Modules for the Browser](#)

[The CommonJS Mailing List](#)

AMD & CommonJS Competing, But Equally Valid Standards

Whilst this section has placed more emphasis on using AMD over CommonJS, the reality is that both formats are valid and have a use.

AMD adopts a browser-first approach to development, opting for asynchronous behavior and simplified backwards compatibility but it doesn't have any concept of File I/O. It supports objects, functions, constructors, strings, JSON and many other types of modules, running natively in the browser. It's incredibly flexible.

CommonJS on the other hand takes a server-first approach, assuming synchronous behavior, no global baggage as John Hann would refer to it as and it attempts to cater for the future (on the server). What we mean by this is that because CommonJS supports unwrapped modules, it can feel a little more close to the ES.next/Harmony specifications, freeing you of the `define()` wrapper that AMD enforces. CommonJS modules however only support objects as modules.

Although the idea of yet another module format may be daunting, you may be interested in some samples of work on hybrid AMD/CommonJS and Universal AMD/CommonJS modules.

Basic AMD Hybrid Format (John Hann)

```
define( function (require, exports, module){  
    var shuffler = require('lib/shuffle');  
    exports.randomize = function( input ){  
        return shuffler.shuffle(input);  
    }  
});
```

Note: this is basically the 'simplified CommonJS wrapper' that is supported in the AMD spec.

AMD/CommonJS Universal Module Definition (Variation 2, UMDjs)

```
/**  
 * exports object based version, if you need to make a  
 * circular dependency or need compatibility with
```

```

* commonjs-like environments that are not Node.
*/
(function (define) {
  //The 'id' is optional, but recommended if this is
  //a popular web library that is used mostly in
  //non-AMD/Node environments. However, if want
  //to make an anonymous module, remove the 'id'
  //below, and remove the id use in the define shim.
  define('id', function (require, exports) {
    //If have dependencies, get them here
    var a = require('a');

    //Attach properties to exports.
    exports.name = value;
  });
})(typeof define === 'function' && define.amd ? define : function (id, factory) {
  if (typeof exports !== 'undefined') {
    //commonjs
    factory(require, exports);
  } else {
    //Create a global function. Only works if
    //the code does not have dependencies, or
    //dependencies fit the call pattern below.
    factory(function(value) {
      return window[value];
    }, (window[id] = {}));
  }
});

```

Extensible UMD Plugins With (Variation by myself and Thomas Davis).

core.js

```

// Module/Plugin core
// Note: the wrapper code you see around the module is what enables
// us to support multiple module formats and specifications by
// mapping the arguments defined to what a specific format expects
// to be present. Our actual module functionality is defined lower
// down, where a named module and exports are demonstrated.

;(function ( name, definition ){
  var theModule = definition(),
      // this is considered "safe":
      hasDefine = typeof define === 'function' && define.amd,
      // hasDefine = typeof define === 'function',
      hasExports = typeof module !== 'undefined' && module.exports;

  if ( hasDefine ){ // AMD Module
    define(theModule);
  } else if ( hasExports ) { // Node.js Module
    module.exports = theModule;
  } else { // Assign to common namespaces or simply the global object (window)
    (this.jquery || this.ender || this.$ || this)[name] = theModule;
  }

```

```

    }
  })( 'core', function () {
    var module = this;
    module.plugins = [];
    module.highlightColor = "yellow";
    module.errorColor = "red";

    // define the core module here and return the public API

    // this is the highlight method used by the core highlightAll()
    // method and all of the plugins highlighting elements different
    // colors
    module.highlight = function(el, strColor){
      // this module uses jQuery, however plain old JavaScript
      // or say, Dojo could be just as easily used.
      if(this.jQuery){
        jQuery(el).css('background', strColor);
      }
    }
    return {
      highlightAll:function(){
        module.highlight('div', module.highlightColor);
      }
    };
  });
});

```

myExtension.js

```

;(function ( name, definition ) {
  var theModule = definition(),
      hasDefine = typeof define === 'function',
      hasExports = typeof module !== 'undefined' && module.exports;

  if ( hasDefine ) { // AMD Module
    define(theModule);
  } else if ( hasExports ) { // Node.js Module
    module.exports = theModule;
  } else { // Assign to common namespaces or simply the global object (window)

    // account for flat-file/global module extensions
    var obj = null;
    var namespaces = name.split(".");
    var scope = (this.jQuery || this.ender || this.$ || this);
    for (var i = 0; i < namespaces.length; i++) {
      var packageName = namespaces[i];
      if (obj && i == namespaces.length - 1) {
        obj[packageName] = theModule;
      } else if (typeof scope[packageName] === "undefined") {
        scope[packageName] = {};
      }
      obj = scope[packageName];
    }
  }
});

```

```

})('core.plugin', function () {

    // define your module here and return the public API
    // this code could be easily adapted with the core to
    // allow for methods that overwrite/extend core functionality
    // to expand the highlight method to do more if you wished.
    return {
        setGreen: function ( el ) {
            highlight(el, 'green');
        },
        setRed: function ( el ) {
            highlight(el, errorColor);
        }
    };

});

```

app.js

```

$(function(){

    // the plugin 'core' is exposed under a core namespace in
    // this example which we first cache
    var core = $.core;

    // use then use some of the built-in core functionality to
    // highlight all divs in the page yellow
    core.highlightAll();

    // access the plugins (extensions) loaded into the 'plugin'
    // namespace of our core module:

    // Set the first div in the page to have a green background.
    core.plugin.setGreen("div:first");
    // Here we're making use of the core's 'highlight' method
    // under the hood from a plugin loaded in after it

    // Set the last div to the 'errorColor' property defined in
    // our core module/plugin. If you review the code further down
    // you'll see how easy it is to consume properties and methods
    // between the core and other plugins
    core.plugin.setRed('div:last');

});

```


ES Harmony

Modules Of The Future

TC39, the standards body charged with defining the syntax and semantics of ECMAScript and its future iterations is composed of a number of very intelligent developers. Some of these developers (such as [Alex Russell](#)) have been keeping a close eye on the evolution of JavaScript usage for large-scale development over the past few years and are acutely aware of the need for better language features for writing more modular JS.

For this reason, there are currently proposals for a number of exciting additions to the language including flexible **modules** that can work on both the client and server, a **module loader** and **more**. In this section, I'll be showing you some code samples of the syntax for modules in ES.next so you can get a taste of what's to come.

Note: Although Harmony is still in the proposal phases, you can already try out (partial) features of ES.next that address native support for writing modular JavaScript thanks to Google's **Traceur** compiler. To get up and running with Traceur in under a minute, read this [getting started](#) guide. There's also a JSConf [presentation](#) about it that's worth looking at if you're interested in learning more about the project.

Modules With Imports And Exports

If you've read through the sections on AMD and CommonJS modules you may be familiar with the concept of module dependencies (imports) and module exports (or, the public API/variables we allow other modules to consume). In ES.next, these concepts have been proposed in a slightly more succinct manner with dependencies being specified using an **import** keyword. **export** isn't greatly different to what we might expect and I think many developers will look at the code below and instantly 'get' it.

- **import** declarations bind a module's exports as local variables and may be renamed to avoid name collisions/conflicts.
- **export** declarations declare that a local-binding of a module is externally visible

such that other modules may read the exports but can't modify them. Interestingly, modules may export child modules however can't export modules that have been defined elsewhere. You may also rename exports so their external name differs from their local names.

```
module staff{
  // specify (public) exports that can be consumed by
  // other modules
  export var baker = {
    bake: function( item ){
      console.log('Woo! I just baked ' + item);
    }
  }
}

module skills{
  export var specialty = "baking";
  export var experience = "5 years";
}

module cakeFactory{

  // specify dependencies
  import baker from staff;

  // import everything with wildcards
  import * from skills;

  export var oven = {
    makeCupcake: function( toppings ){
      baker.bake('cupcake', toppings);
    },
    makeMuffin: function( mSize ){
      baker.bake('muffin', size);
    }
  }
}
```

Modules Loaded From Remote Sources

The module proposals also cater for modules which are remotely based (e.g. a third-party API wrapper) making it simplistic to load modules in from external locations. Here's an example of us pulling in the module we defined above and utilizing it:

```
module cakeFactory from 'http://addyosmani.com/factory/cakes.js';
cakeFactory.oven.makeCupcake('sprinkles');
cakeFactory.oven.makeMuffin('large');
```

Module Loader API

The module loader proposed describes a dynamic API for loading modules in highly controlled contexts. Signatures supported on the loader include `load(url, moduleInstance, error)` for loading modules, `createModule(object, globalModuleReferences)` and **others**. Here's another example of us dynamically loading in the module we initially defined. Note that unlike the last example where we pulled in a module from a remote source, the module loader API is better suited to dynamic contexts.

```
Loader.load('http://addyosmani.com/factory/cakes.js',
  function(cakeFactory){
    cakeFactory.oven.makeCupcake('chocolate');
  });
```

CommonJS-like Modules For The Server

For developers who are server-oriented, the module system proposed for ES.next isn't just constrained to looking at modules in the browser. Below for examples, you can see a CommonJS-like module proposed for use on the server:

```
// io/File.js
export function open(path) { ... };
export function close(hnd) { ... };

// compiler/LexicalHandler.js
module file from 'io/File';

import { open, close } from file;
export function scan(in) {
  try {
    var h = open(in) ...
  }
  finally { close(h) }
}

module lexer from 'compiler/LexicalHandler';
module stdlib from '@std';

//... scan(cmdline[0]) ...
```

Classes With Constructors, Getters & Setters

The notion of a class has always been a contentious issue with purists and we've so far got along with either falling back on JavaScript's **prototypal** nature or through using frameworks or abstractions that offer the ability to use class definitions in a form that desugars to the same prototypal behavior.

In Harmony, classes come as part of the language along with constructors and

(finally) some sense of true privacy. In the following examples, I've included some inline comments to help you understand how classes are structured, but you may also notice the lack of the word 'function' in here. This isn't a typo error: TC39 have been making a conscious effort to decrease our abuse of the `function` keyword for everything and the hope is that this will help simplify how we write code.

```
class Cake{

  // We can define the body of a class' constructor
  // function by using the keyword 'constructor' followed
  // by an argument list of public and private declarations.
  constructor( name, toppings, price, cakeSize ){
    public name = name;
    public cakeSize = cakeSize;
    public toppings = toppings;
    private price = price;
  }

  // As a part of ES.next's efforts to decrease the unnecessary
  // use of 'function' for everything, you'll notice that it's
  // dropped for cases such as the following. Here an identifier
  // followed by an argument list and a body defines a new method

  addTopping( topping ){
    public(this).toppings.push(topping);
  }

  // Getters can be defined by declaring get before
  // an identifier/method name and a curly body.
  get allToppings(){
    return public(this).toppings;
  }

  get qualifiesForDiscount(){
    return private(this).price > 5;
  }

  // Similar to getters, setters can be defined by using
  // the 'set' keyword before an identifier
  set cakeSize( cSize ){
    if( cSize < 0 ){
      throw new Error('Cake must be a valid size -
        either small, medium or large');
    }
    public(this).cakeSize = cSize;
  }
}
```

ES Harmony Conclusions

As you can see, ES.next is coming with some exciting new additions. Although Traceur can be used to an extent to try out such features in the present, remember that it may not be the best idea to plan out your system to use Harmony (just yet). There are risks here such as specifications changing and a potential failure at the cross-browser level (IE9 for example will take a while to die) so your best bets until we have both spec finalization and coverage are AMD (for in-browser modules) and CommonJS (for those on the server).

Related Reading

[A First Look At The Upcoming JavaScript Modules](#)

[David Herman On JavaScript/ES.Next \(Video\)](#)

[ES Harmony Module Proposals](#)

[ES Harmony Module Semantics/Structure Rationale](#)

[ES Harmony Class Proposals](#)

Conclusions And Further Reading A Review

In this section we reviewed several of the options available for writing modular JavaScript using modern module formats. These formats have a number of advantages over using the (classical) module pattern alone including: avoiding a need for developers to create global variables for each module they create, better support for static and dynamic dependency management, improved compatibility with script loaders, better (optional) compatibility for modules on the server and more.

In short, I recommend trying out what's been suggested today as these formats offer a lot of power and flexibility that can help when building applications based on many reusable blocks of functionality.

Chapter 14. Bonus: jQuery Plugin Design Patterns

While well-known JavaScript design patterns can be extremely useful, another side of development could benefit from its own set of design patterns are jQuery plugins. The official jQuery [plugin authoring guide](#) offers a great starting point for getting into writing plugins and widgets, but let's take it further.

Plugin development has evolved over the past few years. We no longer have just one way to write plugins, but many. In reality, certain patterns might work better for a particular problem or component than others.

Some developers may wish to use the jQuery UI [widget factory](#); it's great for complex, flexible UI components. Some may not. Some might like to structure their plugins more like modules (similar to the module pattern) or use a more formal module format such as [AMD \(asynchronous module definition\)](#). Some might want their plugins to harness the power of prototypal inheritance. Some might want to use custom events or pub/sub to communicate from plugins to the rest of their app. And so on.

I began to think about plugin patterns after noticing a number of efforts to create a one-size-fits-all jQuery plugin boilerplate. While such a boilerplate is a great idea in theory, the reality is that we rarely write plugins in one fixed way, using a single pattern all the time.

Let's assume that you've tried your hand at writing your own jQuery plugins at some point and you're comfortable putting together something that works. It's functional. It does what it needs to do, but perhaps you feel it could be structured better. Maybe it could be more flexible or could solve more issues. If this sounds familiar and you aren't sure of the differences between many of the different jQuery plugin patterns, then you might find what I have to say helpful.

My advice won't provide solutions to every possible pattern, but it will cover popular patterns that developers use in the wild.

Note: This section is targeted at intermediate to advanced developers. If you don't

feel you're ready for this just yet, I'm happy to recommend the official jQuery [Plugins/Authoring](#) guide, Ben Alman's [plugin style guide](#) and Remy Sharp's "[Signs of a Poorly Written jQuery Plugin](#)."

Patterns

jQuery plugins have very few defined rules, which one of the reasons for the incredible diversity in how they're implemented. At the most basic level, you can write a plugin simply by adding a new function property to jQuery's `$.fn` object, as follows:

```
$.fn.myPluginName = function() {  
    // your plugin logic  
};
```

This is great for compactness, but the following would be a better foundation to build on:

```
(function( $ ){  
    $.fn.myPluginName = function() {  
        // your plugin logic  
    };  
})( jQuery );
```

Here, we've wrapped our plugin logic in an anonymous function. To ensure that our use of the `$` sign as a shorthand creates no conflicts between jQuery and other JavaScript libraries, we simply pass it to this closure, which maps it to the dollar sign, thus ensuring that it can't be affected by anything outside of its scope of execution.

An alternative way to write this pattern would be to use `$.extend`, which enables you to define multiple functions at once and which sometimes make more sense semantically:

```
(function( $ ){  
    $.extend($.fn, {  
        myplugin: function(){  
            // your plugin logic  
        }  
    });  
})( jQuery );
```

We could do a lot more to improve on all of this; and the first complete pattern we'll be looking at today, the lightweight pattern, covers some best-practices that we can use for basic everyday plugin development and that takes into account

common gotchas to look out for.

Note

While most of the patterns below will be explained, I recommend reading through the comments in the code, because they will offer more insight into why certain practices are best.

I should also mention that none of this would be possible without the previous work, input and advice of other members of the jQuery community. I've listed them inline with each pattern so that you can read up on their individual work if interested.

A Lightweight Start

Let's begin our look at patterns with something basic that follows best-practices (including those in the jQuery plugin-authoring guide). This pattern is ideal for developers who are either new to plugin development or who just want to achieve something simple (such as a utility plugin). This lightweight start uses the following:

- Common best-practices, such as a semi-colon before the function's invocation; `window`, `document`, `undefined` passed in as arguments; and adherence to the jQuery core style guidelines.
- A basic defaults object.
- A simple plugin constructor for logic related to the initial creation and the assignment of the element to work with.
- Extending the options with defaults.
- A lightweight wrapper around the constructor, which helps to avoid issues such as multiple instantiations.

```
/*!
 * jQuery lightweight plugin boilerplate
 * Original author: @ajpiano
 * Further changes, comments: @addyosmani
 * Licensed under the MIT license
 */

// the semi-colon before the function invocation is a safety
// net against concatenated scripts and/or other plugins
// that are not closed properly.
;(function ( $, window, document, undefined ) {

    // undefined is used here as the undefined global
    // variable in ECMAScript 3 and is mutable (i.e. it can
    // be changed by someone else). undefined isn't really
    // being passed in so we can ensure that its value is
    // truly undefined. In ES5, undefined can no longer be
    // modified.

    // window and document are passed through as local
    // variables rather than as globals, because this (slightly)
    // quickens the resolution process and can be more
    // efficiently minified (especially when both are
    // regularly referenced in your plugin).
```

```

// Create the defaults once
var pluginName = 'defaultPluginName',
    defaults = {
        propertyName: "value"
    };

// The actual plugin constructor
function Plugin( element, options ) {
    this.element = element;

    // jQuery has an extend method that merges the
    // contents of two or more objects, storing the
    // result in the first object. The first object
    // is generally empty because we don't want to alter
    // the default options for future instances of the plugin
    this.options = $.extend( {}, defaults, options ) ;

    this._defaults = defaults;
    this._name = pluginName;

    this.init();
}

Plugin.prototype.init = function () {
    // Place initialization logic here
    // You already have access to the DOM element and
    // the options via the instance, e.g. this.element
    // and this.options
};

// A really lightweight plugin wrapper around the constructor,
// preventing against multiple instantiations
$.fn[pluginName] = function ( options ) {
    return this.each(function () {
        if (!$.data(this, 'plugin_' + pluginName)) {
            $.data(this, 'plugin_' + pluginName,
                new Plugin( this, options ));
        }
    });
}

})( jQuery, window, document );

```

Usage:

```

$('#elem').defaultPluginName({
    propertyName: 'a custom value'
});

```

Further Reading

- [Plugins/Authoring, jQuery](#)
- [“Signs of a Poorly Written jQuery Plugin,” Remy Sharp](#)

- “How to Create Your Own jQuery Plugin,” Elijah Manor
- “Style in jQuery Plugins and Why It Matters,” Ben Almon
- “Create Your First jQuery Plugin, Part 2,” Andrew Wirick

“Complete” Widget Factory

While the authoring guide is a great introduction to plugin development, it doesn't offer a great number of conveniences for obscuring away from common plumbing tasks that we have to deal with on a regular basis.

The jQuery UI Widget Factory is a solution to this problem that helps you build complex, stateful plugins based on object-oriented principles. It also eases communication with your plugin's instance, obfuscating a number of the repetitive tasks that you would have to code when working with basic plugins.

In case you haven't come across these before, stateful plugins keep track of their current state, also allowing you to change properties of the plugin after it has been initialized.

One of the great things about the Widget Factory is that the majority of the jQuery UI library actually uses it as a base for its components. This means that if you're looking for further guidance on structure beyond this template, you won't have to look beyond the jQuery UI repository.

Back to patterns. This jQuery UI boilerplate does the following:

- Covers almost all supported default methods, including triggering events.
- Includes comments for all of the methods used, so that you're never unsure of where logic should fit in your plugin.

```
/*!
 * jQuery UI Widget-factory plugin boilerplate (for 1.8/9+)
 * Author: @addyosmani
 * Further changes: @peolanh
 * Licensed under the MIT license
 */

;(function ( $, window, document, undefined ) {

    // define your widget under a namespace of your choice
    // with additional parameters e.g.
    // $.widget( "namespace.widgetname", (optional) - an
    // existing widget prototype to inherit from, an object
    // literal to become the widget's prototype );

    $.widget( "namespace.widgetname" , {

        //Options to be used as defaults
```

```

options: {
    someValue: null
},

//Setup widget (e.g. element creation, apply theming
// , bind events etc.)
_create: function () {

    // _create will automatically run the first time
    // this widget is called. Put the initial widget
    // setup code here, then you can access the element
    // on which the widget was called via this.element.
    // The options defined above can be accessed
    // via this.options this.element.addStuff();
},

// Destroy an instantiated plugin and clean up
// modifications the widget has made to the DOM
destroy: function () {

    // this.element.removeStuff();
    // For UI 1.8, destroy must be invoked from the
    // base widget
    $.Widget.prototype.destroy.call(this);
    // For UI 1.9, define _destroy instead and don't
    // worry about
    // calling the base widget
},

methodB: function ( event ) {
    //_trigger dispatches callbacks the plugin user
    // can subscribe to
    // signature: _trigger( "callbackName" , [eventObject],
    // [uiObject] )
    // e.g. this._trigger( "hover", e /*where e.type ==
    // "mouseenter"*/, { hovered: $(e.target)});
    this._trigger('methodA', event, {
        key: value
    });
},

methodA: function ( event ) {
    this._trigger('dataChanged', event, {
        key: value
    });
},

// Respond to any changes the user makes to the
// option method
_setOption: function ( key, value ) {
    switch (key) {
        case "someValue":
            //this.options.someValue = doSomethingWith( value );
            break;
        default:
            //this.options[ key ] = value;

```

```

        break;
    }

    // For UI 1.8, _setOption must be manually invoked
    // from the base widget
    $.Widget.prototype._setOption.apply( this, arguments );
    // For UI 1.9 the _super method can be used instead
    // this._super( "_setOption", key, value );
}
});

})( jQuery, window, document );

```

Usage:

```

var collection = $('#elem').widgetName({
    foo: false
});

collection.widgetName('methodB');

```

Further Reading

- [The jQuery UI Widget Factory](#)
- [“Introduction to Stateful Plugins and the Widget Factory,”](#) Doug Neiner
- [“Widget Factory”](#) (explained), Scott Gonzalez
- [“Understanding jQuery UI Widgets: A Tutorial,”](#) Hacking at 0300

Namespacing And Nested Namespacing

Namespacing your code is a way to avoid collisions with other objects and variables in the global namespace. They're important because you want to safeguard your plugin from breaking in the event that another script on the page uses the same variable or plugin names as yours. As a good citizen of the global namespace, you must also do your best not to prevent other developer's scripts from executing because of the same issues.

JavaScript doesn't really have built-in support for namespaces as other languages do, but it does have objects that can be used to achieve a similar effect. Employing a top-level object as the name of your namespace, you can easily check for the existence of another object on the page with the same name. If such an object does not exist, then we define it; if it does exist, then we simply extend it with our plugin.

Objects (or, rather, object literals) can be used to create nested namespaces, such as `namespace.subnamespace.pluginName` and so on. But to keep things simple, the namespacing boilerplate below should give you everything you need to get started with these concepts.

```
/*!
 * jQuery namespaced 'Starter' plugin boilerplate
 * Author: @dougneiner
 * Further changes: @addyosmani
 * Licensed under the MIT license
 */

;(function ( $ ) {
  if (!$.myNamespace) {
    $.myNamespace = {};
  };

  $.myNamespace.myPluginName = function ( el, myFunctionParam, options ) {
    // To avoid scope issues, use 'base' instead of 'this'
    // to reference this class from internal events and functions.
    var base = this;

    // Access to jQuery and DOM versions of element
    base.$el = $(el);
    base.el = el;

    // Add a reverse reference to the DOM object
    base.$el.data( "myNamespace.myPluginName" , base );
  };
});
```



```

base.init = function () {
    base.myFunctionParam = myFunctionParam;

    base.options = $.extend({},
        $.myNamespace.myPluginName.defaultOptions, options);

    // Put your initialization code here
};

// Sample Function, Uncomment to use
// base.functionName = function( paramaters ){
//
// };
// Run initializer
base.init();
};

$.myNamespace.myPluginName.defaultOptions = {
    myDefaultValue: ""
};

$.fn.mynamespace_myPluginName = function
( myFunctionParam, options ) {
    return this.each(function () {
        (new $.myNamespace.myPluginName(this,
            myFunctionParam, options));
    });
};

})( jQuery );

```

Usage:

```

$( '#elem' ).mynamespace_myPluginName({
    myDefaultValue: "foobar"
});

```

Further Reading

- “Namespacing in JavaScript,” Angus Croll
- “Use Your \$.fn jQuery Namespace,” Ryan Florence
- “JavaScript Namespacing,” Peter Michaux
- “Modules and namespaces in JavaScript,” Axel Rauschmayer

Custom Events For Pub/Sub (With The Widget factory)

You may have used the Observer (Pub/Sub) pattern in the past to develop asynchronous JavaScript web applications. The basic idea here is that elements will publish event notifications when something interesting occurs in your application. Other elements then subscribe to or listen for these events and respond accordingly. This results in the logic for your application being significantly more decoupled (which is always good).

In jQuery, we have this idea that custom events provide a built-in means to implement a publish and subscribe system that's quite similar to the Observer pattern. So, `bind('eventType')` is functionally equivalent to performing `subscribe('eventType')`, and `trigger('eventType')` is roughly equivalent to `publish('eventType')`.

Some developers might consider the jQuery event system as having too much overhead to be used as a publish and subscribe system, but it's been architected to be both reliable and robust for most use cases. In the following jQuery UI widget factory template, we'll implement a basic custom event-based pub/sub pattern that allows our plugin to subscribe to event notifications from the rest of our application, which publishes them.

```
/*!
 * jQuery custom-events plugin boilerplate
 * Author: DevPatch
 * Further changes: @addyosmani
 * Licensed under the MIT license
 */

// In this pattern, we use jQuery's custom events to add
// pub/sub (publish/subscribe) capabilities to widgets.
// Each widget would publish certain events and subscribe
// to others. This approach effectively helps to decouple
// the widgets and enables them to function independently.

;(function ( $, window, document, undefined ) {
    $.widget("ao.eventStatus", {
        options: {

        },

        _create : function() {
            var self = this;
```

```

        //self.element.addClass( "my-widget" );

        //subscribe to 'myEventStart'
        self.element.bind( "myEventStart", function( e ) {
            console.log("event start");
        });

        //subscribe to 'myEventEnd'
        self.element.bind( "myEventEnd", function( e ) {
            console.log("event end");
        });

        //unsubscribe to 'myEventStart'
        //self.element.unbind( "myEventStart", function(e){
            ///console.log("unsubscribed to this event");
        //});
    },

    destroy: function(){
        $.Widget.prototype.destroy.apply( this, arguments );
    },
});
})( jQuery, window , document );

// Publishing event notifications
// $(".my-widget").trigger("myEventStart");
// $(".my-widget").trigger("myEventEnd");

```

Usage:

```

var el = $('#elem');
el.eventStatus();
el.eventStatus().trigger('myEventStart');

```

Further Reading

- “Communication Between jQuery UI Widgets,” Benjamin Sternthal

Prototypal Inheritance With The DOM-To-Object Bridge Pattern

In JavaScript, we don't have the traditional notion of classes that you would find in other classical programming languages, but we do have prototypal inheritance. With prototypal inheritance, an object inherits from another object. And we can apply this concept to jQuery plugin development.

Alex Sexton and Scott Gonzalez have looked at this topic in detail. In sum, they found that for organized modular development, clearly separating the object that defines the logic for a plugin from the plugin-generation process itself can be beneficial. The benefit is that testing your plugin's code becomes easier, and you can also adjust the way things work behind the scenes without altering the way that any object APIs you've implemented are used.

In Sexton's previous post on this topic, he implements a bridge that enables you to attach your general logic to a particular plugin, which we've implemented in the template below. Another advantage of this pattern is that you don't have to constantly repeat the same plugin initialization code, thus ensuring that the concepts behind DRY development are maintained. Some developers might also find this pattern easier to read than others.

```
/*!  
 * jQuery prototypal inheritance plugin boilerplate  
 * Author: Alex Sexton, Scott Gonzalez  
 * Further changes: @addyosmani  
 * Licensed under the MIT license  
 */  
  
// myObject - an object representing a concept that you want  
// to model (e.g. a car)  
var myObject = {  
  init: function( options, elem ) {  
    // Mix in the passed-in options with the default options  
    this.options = $.extend( {}, this.options, options );  
  
    // Save the element reference, both as a jQuery  
    // reference and a normal reference  
    this.elem = elem;  
    this.$elem = $(elem);  
  
    // Build the DOM's initial structure  
    this._build();  
  }  
};
```

```

        // return this so that we can chain and use the bridge with less code.
        return this;
    },
    options: {
        name: "No name"
    },
    _build: function(){
        //this.$elem.html('<h1>'+this.options.name+'</h1>');
    },
    myMethod: function( msg ){
        // You have direct access to the associated and cached
        // jQuery element
        // this.$elem.append('<p>'+msg+'</p>');
    }
};

// Object.create support test, and fallback for browsers without it
if ( typeof Object.create !== 'function' ) {
    Object.create = function (o) {
        function F() {}
        F.prototype = o;
        return new F();
    };
}

// Create a plugin based on a defined object
$.plugin = function( name, object ) {
    $.fn[name] = function( options ) {
        return this.each(function() {
            if ( ! $.data( this, name ) ) {
                $.data( this, name, Object.create(object).init(
                    options, this ) );
            }
        });
    };
};
};

```

Usage:

```

$.plugin('myobj', myObject);

$('#elem').myobj({name: "John"});

var collection = $('#elem').data('myobj');
collection.myMethod('I am a method');

```

Further Reading

- “Using Inheritance Patterns To Organize Large jQuery Applications,” Alex Sexton

- “How to Manage Large Applications With jQuery or Whatever” (further discussion), Alex Sexton
- “Practical Example of the Need for Prototypal Inheritance,” Neeraj Singh
- “Prototypal Inheritance in JavaScript,” Douglas Crockford

jQuery UI Widget Factory Bridge

If you liked the idea of generating plugins based on objects in the last design pattern, then you might be interested in a method found in the jQuery UI Widget Factory called `$.widget.bridge`. This bridge basically serves as a middle layer between a JavaScript object that is created using `$.widget` and jQuery's API, providing a more built-in solution to achieving object-based plugin definition. Effectively, we're able to create stateful plugins using a custom constructor.

Moreover, `$.widget.bridge` provides access to a number of other capabilities, including the following:

- Both public and private methods are handled as one would expect in classical OOP (i.e. public methods are exposed, while calls to private methods are not possible);
- Automatic protection against multiple initializations;
- Automatic generation of instances of a passed object, and storage of them within the selection's internal `$.data` cache;
- Options can be altered post-initialization.

For further information on how to use this pattern, look at the comments in the boilerplate below:

```
/*!  
 * jQuery UI Widget factory "bridge" plugin boilerplate  
 * Author: @erichynds  
 * Further changes, additional comments: @addyosmani  
 * Licensed under the MIT license  
 */  
  
// a "widgetName" object constructor  
// required: this must accept two arguments,  
// options: an object of configuration options  
// element: the DOM element the instance was created on  
var widgetName = function( options, element ){  
    this.name = "myWidgetName";  
    this.options = options;  
    this.element = element;  
    this._init();  
}  
  
// the "widgetName" prototype
```

```

widgetName.prototype = {

    // _create will automatically run the first time this
    // widget is called
    _create: function(){
        // creation code
    },

    // required: initialization logic for the plugin goes into _init
    // This fires when your instance is first created and when
    // attempting to initialize the widget again (by the bridge)
    // after it has already been initialized.
    _init: function(){
        // init code
    },

    // required: objects to be used with the bridge must contain an
    // 'option'. Post-initialization, the logic for changing options
    // goes here.
    option: function( key, value ){

        // optional: get/change options post initialization
        // ignore if you don't require them.

        // signature: $('#foo').bar({ cool:false });
        if( $.isPlainObject( key ) ){
            this.options = $.extend( true, this.options, key );

            // signature: $('#foo').option('cool'); - getter
        } else if ( key && typeof value === "undefined" ){
            return this.options[ key ];

            // signature: $('#foo').bar('option', 'baz', false);
        } else {
            this.options[ key ] = value;
        }

        // required: option must return the current instance.
        // When re-initializing an instance on elements, option
        // is called first and is then chained to the _init method.
        return this;
    },

    // notice no underscore is used for public methods
    publicFunction: function(){
        console.log('public function');
    },

    // underscores are used for private methods
    _privateFunction: function(){
        console.log('private function');
    }
};

```

Usage:


```
// connect the widget obj to jQuery's API under the "foo" namespace
$.widget.bridge("foo", widgetName);

// create an instance of the widget for use
var instance = $('#foo').foo({
  baz: true
});

// your widget instance exists in the elem's data
console.log(instance.data("foo").element); // => #elem element

// bridge allows you to call public methods...
instance.foo("publicFunction"); // => "public method"

// bridge prevents calls to internal methods
instance.foo("_privateFunction"); // => #elem element
```

Further Reading

- “Using \$.widget.bridge Outside of the Widget Factory,” Eric Hynds

jQuery Mobile Widgets With The Widget factory

jQuery mobile is a framework that encourages the design of ubiquitous Web applications that work both on popular mobile devices and platforms and on the desktop. Rather than writing unique applications for each device or OS, you simply write the code once and it should ideally run on many of the A-, B- and C-grade browsers out there at the moment.

The fundamentals behind jQuery mobile can also be applied to plugin and widget development, as seen in some of the core jQuery mobile widgets used in the official library suite. What's interesting here is that even though there are very small, subtle differences in writing a "mobile"-optimized widget, if you're familiar with using the jQuery UI Widget Factory, you should be able to start writing these right away.

The mobile-optimized widget below has a number of interesting differences than the standard UI widget pattern we saw earlier:

- `$.mobile.widget` is referenced as an existing widget prototype from which to inherit. For standard widgets, passing through any such prototype is unnecessary for basic development, but using this jQuery-mobile specific widget prototype provides internal access to further "options" formatting.
- You'll notice in `_create()` a guide on how the official jQuery mobile widgets handle element selection, opting for a role-based approach that better fits the jQM mark-up. This isn't at all to say that standard selection isn't recommended, only that this approach might make more sense given the structure of jQM pages.
- Guidelines are also provided in comment form for applying your plugin methods on `pagecreate` as well as for selecting the plugin application via data roles and data attributes.

```
/*!  
 * (jQuery mobile) jQuery UI Widget-factory plugin boilerplate (for 1.8/9+)  
 * Author: @scottjehl  
 * Further changes: @addyosmani  
 * Licensed under the MIT license  
 */  
  
;(function ( $, window, document, undefined ) {
```

```

//define a widget under a namespace of your choice
//here 'mobile' has been used in the first parameter
$.widget( "mobile.widgetName", $.mobile.widget, {

    //Options to be used as defaults
    options: {
        foo: true,
        bar: false
    },

    _create: function() {
        // _create will automatically run the first time this
        // widget is called. Put the initial widget set-up code
        // here, then you can access the element on which
        // the widget was called via this.element
        // The options defined above can be accessed via
        // this.options

        //var m = this.element,
        //p = m.parents(":jqmData(role='page')"),
        //c = p.find(":jqmData(role='content')")
    },

    // Private methods/props start with underscores
    _dosomething: function(){ ... },

    // Public methods like these below can can be called
    // externally:
    // $("#myelem").foo( "enable", arguments );

    enable: function() { ... },

    // Destroy an instantiated plugin and clean up modifications
    // the widget has made to the DOM
    destroy: function () {
        //this.element.removeStuff();
        // For UI 1.8, destroy must be invoked from the
        // base widget
        $.Widget.prototype.destroy.call(this);
        // For UI 1.9, define _destroy instead and don't
        // worry about calling the base widget
    },

    methodB: function ( event ) {
        //_trigger dispatches callbacks the plugin user can
        // subscribe to
        //signature: _trigger( "callbackName" , [eventObject],
        // [uiObject] )
        // e.g. this._trigger( "hover", e /*where e.type ==
        // "mouseenter"*/, { hovered: $(e.target)});
        this._trigger('methodA', event, {
            key: value
        });
    },

```

```

methodA: function ( event ) {
    this._trigger('dataChanged', event, {
        key: value
    });
},

//Respond to any changes the user makes to the option method
_setOption: function ( key, value ) {
    switch (key) {
        case "someValue":
            //this.options.someValue = doSomethingWith( value );
            break;
        default:
            //this.options[ key ] = value;
            break;
    }

    // For UI 1.8, _setOption must be manually invoked from
    // the base widget
    $.Widget.prototype._setOption.apply(this, arguments);
    // For UI 1.9 the _super method can be used instead
    // this._super( "_setOption", key, value );
}

});

})( jQuery, window, document );

```

Usage:

```

var instance = $('#foo').widgetName({
    foo: false
});

instance.widgetName( 'methodB' );

```

We can also self-initialize this widget whenever a new page in jQuery Mobile is created. jQuery Mobile's "page" plugin dispatches a "create" event when a jQuery Mobile page (found via data-role=page attr) is first initialized. We can listen for that event (called "pagecreate") and run our plugin automatically whenever a new page is created.

```

$(document).bind("pagecreate", function (e) {
    // In here, e.target refers to the page that was created
    // (it's the target of the pagecreate event)
    // So, we can simply find elements on this page that match a
    // selector of our choosing, and call our plugin on them.
    // Here's how we'd call our "foo" plugin on any element with a
    // data-role attribute of "foo":
    $(e.target).find("[data-role='foo']").foo(options);

    // Or, better yet, let's write the selector accounting for the configurable
    // data-attribute namespace
    $(e.target).find(":jqmData(role='foo)').foo(options);
});

```

That's it. Now you can simply reference the script containing your widget and pagecreate binding in a page running jQuery Mobile site, and it will automatically run like any other jQuery Mobile plugin.

RequireJS And The jQuery UI Widget Factory

RequireJS is a script loader that provides a clean solution for encapsulating application logic inside manageable modules. It's able to load modules in the correct order (through its order plugin); it simplifies the process of combining scripts via its excellent optimizer; and it provides the means for defining module dependencies on a per-module basis.

James Burke has written a comprehensive set of tutorials on getting started with RequireJS. But what if you're already familiar with it and would like to wrap your jQuery UI widgets or plugins in a RequireJS-compatible module wrapper?.

In the boilerplate pattern below, we demonstrate how a compatible widget can be defined that does the following:

- Allows the definition of widget module dependencies, building on top of the previous jQuery UI boilerplate presented earlier;
- Demonstrates one approach to passing in HTML template assets for creating templated widgets with jQuery (in conjunction with the jQuery tmpl plugin) (View the comments in `_create()`.)
- Includes a quick tip on adjustments that you can make to your widget module if you wish to later pass it through the RequireJS optimizer

```
/*!  
 * jQuery UI Widget + RequireJS module boilerplate (for 1.8/9+)  
 * Authors: @jrburke, @addyosmani  
 * Licensed under the MIT license  
 */  
  
// Note from James:  
//  
// This assumes you are using the RequireJS+jQuery file, and  
// that the following files are all in the same directory:  
//  
// - require-jquery.js  
// - jquery-ui.custom.min.js (custom jQuery UI build with widget factory)  
// - templates/  
//   - asset.html  
// - ao.myWidget.js  
  
// Then you can construct the widget like so:
```

```

//ao.myWidget.js file:
define("ao.myWidget", ["jquery", "text!templates/asset.html", "jquery-
ui.custom.min", "jquery.tmpl"], function ($, assetHtml) {

    // define your widget under a namespace of your choice
    // 'ao' is used here as a demonstration
    $.widget( "ao.myWidget", {

        // Options to be used as defaults
        options: {},

        // Set up widget (e.g. create element, apply theming,
        // bind events, etc.)
        _create: function () {

            // _create will automatically run the first time
            // this widget is called. Put the initial widget
            // set-up code here, then you can access the element
            // on which the widget was called via this.element.
            // The options defined above can be accessed via
            // this.options

            //this.element.addStuff();
            //this.element.addStuff();
            //this.element.tmpl(assetHtml).appendTo(this.content);
        },

        // Destroy an instantiated plugin and clean up modifications
        // that the widget has made to the DOM
        destroy: function () {
            //this.element.removeStuff();
            // For UI 1.8, destroy must be invoked from the base
            // widget
            $.Widget.prototype.destroy.call( this );
            // For UI 1.9, define _destroy instead and don't worry
            // about calling the base widget
        },

        methodB: function ( event ) {
            // _trigger dispatches callbacks the plugin user can
            // subscribe to
            //signature: _trigger( "callbackName" , [eventObject],
            // [uiObject] )
            this._trigger( 'methodA', event, {
                key: value
            });
        },

        methodA: function ( event ) {
            this._trigger( 'dataChanged', event, {
                key: value
            });
        },

        //Respond to any changes the user makes to the option method

```

```

        _setOption: function ( key, value ) {
            switch (key) {
                case "someValue":
                    //this.options.someValue = doSomethingWith( value );
                    break;
                default:
                    //this.options[ key ] = value;
                    break;
            }

            // For UI 1.8, _setOption must be manually invoked from
            // the base widget
            $.Widget.prototype._setOption.apply( this, arguments );
            // For UI 1.9 the _super method can be used instead
            //this._super( "_setOption", key, value );
        }

        //somewhere assetHtml would be used for templating, depending
        // on your choice.
    });
});

```

Usage:

index.html:

```

<script data-main="scripts/main"
src="http://requirejs.org/docs/release/1.0.1/minified/require.js"></script>

```

main.js

```

require({
    paths: {
        'jquery': 'https://ajax.googleapis.com/ajax/libs/jquery/1.7.1/jquery.min',
        'jqueryui': 'https://ajax.googleapis.com/ajax/libs/jqueryui/1.8.18/jquery-
ui.min',
        'boilerplate': '../patterns/jquery.widget-factory.requirejs.boilerplate'
    }
}, ['require', 'jquery', 'jqueryui', 'boilerplate'],
function (req, $) {

    $(function () {

        var instance = $('#elem').myWidget();
        instance.myWidget('methodB');

    });
});

```

Further Reading

- Using RequireJS with jQuery, Rebecca Murphey
- “Fast Modular Code With jQuery and RequireJS,” James Burke

- “jQuery’s Best Friends,” Alex Sexton
- “Managing Dependencies With RequireJS,” Ruslan Matveev

Globally And Per-Call Overridable Options (Best Options Pattern)

For our next pattern, we'll look at an optimal approach to configuring options and defaults for your plugin. The way you're probably familiar with defining plugin options is to pass through an object literal of defaults to `$.extend`, as demonstrated in our basic plugin boilerplate.

If, however, you're working with a plugin with many customizable options that you would like users to be able to override either globally or on a per-call level, then you can structure things a little differently.

Instead, by referring to an options object defined within the plugin namespace explicitly (for example, `$.fn.pluginName.options`) and merging this with any options passed through to the plugin when it is initially invoked, users have the option of either passing options through during plugin initialization or overriding options outside of the plugin (as demonstrated here).

```
/*!
 * jQuery 'best options' plugin boilerplate
 * Author: @cowboy
 * Further changes: @addyosmani
 * Licensed under the MIT license
 */

;(function ( $, window, document, undefined ) {

    $.fn.pluginName = function ( options ) {

        // Here's a best-practice for overriding 'defaults'
        // with specified options. Note how, rather than a
        // regular defaults object being passed as the second
        // parameter, we instead refer to $.fn.pluginName.options
        // explicitly, merging it with the options passed directly
        // to the plugin. This allows us to override options both
        // globally and on a per-call level.

        options = $.extend( {}, $.fn.pluginName.options, options );

        return this.each(function () {

            var elem = $(this);

        });
    };
});
```

```

// Globally overriding options
// Here are our publicly accessible default plugin options
// that are available in case the user doesn't pass in all
// of the values expected. The user is given a default
// experience but can also override the values as necessary.
// e.g. $fn.pluginName.key ='otherval';

$.fn.pluginName.options = {

    key: "value",
    myMethod: function ( elem, param ) {

    }

};

})( jQuery, window, document );

```

Usage:

```

$( '#elem' ).pluginName({
    key: "foobar"
});

```

Further Reading

- [jQuery Pluginization](#) and the [accompanying gist](#), Ben Alman

A Highly Configurable And Mutable Plugin

Like Alex Sexton's pattern, the following logic for our plugin isn't nested in a jQuery plugin itself. We instead define our plugin's logic using a constructor and an object literal defined on its prototype, using jQuery for the actual instantiation of the plugin object.

Customization is taken to the next level by employing two little tricks, one of which you've seen in previous patterns:

- Options can be overridden both globally and per collection of elements;
- Options can be customized on a **per-element** level through HTML5 data attributes (as shown below). This facilitates plugin behavior that can be applied to a collection of elements but then customized inline without the need to instantiate each element with a different default value.

You don't see the latter option in the wild too often, but it can be a significantly cleaner solution (as long as you don't mind the inline approach). If you're wondering where this could be useful, imagine writing a draggable plugin for a large set of elements. You could go about customizing their options like this:

```
javascript
$('.item-a').draggable({'defaultPosition':'top-left'});
$('.item-b').draggable({'defaultPosition':'bottom-right'});
$('.item-c').draggable({'defaultPosition':'bottom-left'});
//etc
```

But using our patterns inline approach, the following would be possible:

```
javascript
$('.items').draggable();

html
<li class="item" data-plugin-options='{"defaultPosition":"top-left"}'></div>
<li class="item" data-plugin-options='{"defaultPosition":"bottom-left"}'></div>
```

And so on. You may well have a preference for one of these approaches, but it is another potentially useful pattern to be aware of.

```
/*
 * 'Highly configurable' mutable plugin boilerplate
 * Author: @markdalgleish
 * Further changes, comments: @addyosmani
 * Licensed under the MIT license
 */
```

```

// Note that with this pattern, as per Alex Sexton's, the plugin logic
// hasn't been nested in a jQuery plugin. Instead, we just use
// jQuery for its instantiation.

;(function( $, window, document, undefined ){

    // our plugin constructor
    var Plugin = function( elem, options ){
        this.elem = elem;
        this.$elem = $(elem);
        this.options = options;

        // This next line takes advantage of HTML5 data attributes
        // to support customization of the plugin on a per-element
        // basis. For example,
        // <div class=item' data-plugin-options='{ "message": "Goodbye World!" }'></div>
        this.metadata = this.$elem.data( 'plugin-options' );
    };

    // the plugin prototype
    Plugin.prototype = {
        defaults: {
            message: 'Hello world!'
        },

        init: function() {
            // Introduce defaults that can be extended either
            // globally or using an object literal.
            this.config = $.extend({}, this.defaults, this.options,
            this.metadata);

            // Sample usage:
            // Set the message per instance:
            // $('#elem').plugin({ message: 'Goodbye World!' });
            // or
            // var p = new Plugin(document.getElementById('elem'),
            // { message: 'Goodbye World!' }).init()
            // or, set the global default message:
            // Plugin.defaults.message = 'Goodbye World!'

            this.sampleMethod();
            return this;
        },

        sampleMethod: function() {
            // e.g. show the currently configured message
            // console.log(this.config.message);
        }
    }

    Plugin.defaults = Plugin.prototype.defaults;

    $.fn.plugin = function(options) {
        return this.each(function() {
            new Plugin(this, options).init();
        });
    }

```

```
    });  
};  
  
//optional: window.Plugin = Plugin;  
  
})( jQuery, window , document );
```

Usage:

```
$( '#elem' ).plugin({  
  message: "foobar"  
});
```

Further Reading

- “Creating Highly Configurable jQuery Plugins,” Mark Dalgleish
- “Writing Highly Configurable jQuery Plugins, Part 2,” Mark Dalgleish

UMD: AMD And CommonJS-Compatible Modules For Plugins

Whilst many of the plugin and widget patterns presented above are acceptable for general use, they aren't without their caveats. Some require jQuery or the jQuery UI Widget Factory to be present in order to function, while only a few could be easily adapted to work well as globally compatible modules in both the browser and other environments.

We've already explored both AMD and CommonJS in the last chapter, but imagine how useful it would be if we could define and load plugin modules compatible with AMD, CommonJS and other standards that are also compatible with different environments (client-side, server-side and beyond).

To provide a solution for this problem, a number of developers including James Burke, myself, Thomas Davis and Ryan Florence have been working on an effort known as UMD (or Universal Module Definition). The goal of our efforts has been to provide a set of agreed upon patterns for plugins that can work in all environments. At present, a number of such boilerplates have been completed and are available on the UMD group repo <https://github.com/umdjs/umd>.

One such pattern we've worked on for **jQuery plugins** appears below and has the following features:

- A core/base plugin is loaded into a `$.core` namespace, which can then be easily extended using plugin extensions via the namespacing pattern. Plugins loaded via script tags automatically populate a `plugin` namespace under `core` (i.e. `$.core.plugin.methodName()`).
- The pattern can be quite nice to work with because plugin extensions can access properties and methods defined in the base or, with a little tweaking, override default behavior so that it can be extended to do more.
- A loader isn't required at all to make this pattern fully function.

usage.html

```
<script type="text/javascript" src="http://code.jquery.com/jquery-1.7.1.min.js">
</script>
<script type="text/javascript" src="pluginCore.js"></script>
<script type="text/javascript" src="pluginExtension.js"></script>
```

```

<script type="text/javascript">

$(function(){

    // Our plugin 'core' is exposed under a core namespace in
    // this example, which we first cache
    var core = $.core;

    // Then use some of the built-in core functionality to
    // highlight all divs in the page yellow
    core.highlightAll();

    // Access the plugins (extensions) loaded into the 'plugin'
    // namespace of our core module:

    // Set the first div in the page to have a green background.
    core.plugin.setGreen("div:first");
    // Here we're making use of the core's 'highlight' method
    // under the hood from a plugin loaded in after it

    // Set the last div to the 'errorColor' property defined in
    // our core module/plugin. If you review the code further down,
    // you'll see how easy it is to consume properties and methods
    // between the core and other plugins
    core.plugin.setRed('div:last');
});

</script>

```

pluginCore.js

```

// Module/Plugin core
// Note: the wrapper code you see around the module is what enables
// us to support multiple module formats and specifications by
// mapping the arguments defined to what a specific format expects
// to be present. Our actual module functionality is defined lower
// down, where a named module and exports are demonstrated.
//
// Note that dependencies can just as easily be declared if required
// and should work as demonstrated earlier with the AMD module examples.

(function ( name, definition ){
    var theModule = definition(),
        // this is considered "safe":
        hasDefine = typeof define === 'function' && define.amd,
        // hasDefine = typeof define === 'function',
        hasExports = typeof module !== 'undefined' && module.exports;

    if ( hasDefine ){ // AMD Module
        define(theModule);
    } else if ( hasExports ) { // Node.js Module
        module.exports = theModule;
    } else { // Assign to common namespaces or simply the global object (window)
        (this.jquery || this.endr || this.$ || this)[name] = theModule;
    }
}

```



```

))( 'core', function () {
    var module = this;
    module.plugins = [];
    module.highlightColor = "yellow";
    module.errorColor = "red";

    // define the core module here and return the public API

    // This is the highlight method used by the core highlightAll()
    // method and all of the plugins highlighting elements different
    // colors
    module.highlight = function(el,strColor){
        if(this.jquery){
            jQuery(el).css('background', strColor);
        }
    }
    return {
        highlightAll:function(){
            module.highlight('div', module.highlightColor);
        }
    };
});

```

pluginExtension.js

```

// Extension to module core

(function ( name, definition ) {
    var theModule = definition(),
        hasDefine = typeof define === 'function',
        hasExports = typeof module !== 'undefined' && module.exports;

    if ( hasDefine ) { // AMD Module
        define(theModule);
    } else if ( hasExports ) { // Node.js Module
        module.exports = theModule;
    } else { // Assign to common namespaces or simply the global object (window)

        // account for for flat-file/global module extensions
        var obj = null;
        var namespaces = name.split(".");
        var scope = (this.jquery || this.endr || this.$ || this);
        for (var i = 0; i < namespaces.length; i++) {
            var packageName = namespaces[i];
            if (obj && i == namespaces.length - 1) {
                obj[packageName] = theModule;
            } else if (typeof scope[packageName] === "undefined") {
                scope[packageName] = {};
            }
            obj = scope[packageName];
        }
    }
})('core.plugin', function () {

```

```

// Define your module here and return the public API.
// This code could be easily adapted with the core to
// allow for methods that overwrite and extend core functionality
// in order to expand the highlight method to do more if you wish.
return {
    setGreen: function ( el ) {
        highlight(el, 'green');
    },
    setRed: function ( el ) {
        highlight(el, errorColor);
    }
};

});

```

Whilst work on improving these patterns is ongoing, please do feel free to check out the patterns suggested to date as you may find them helpful.

Further Reading

- “Using AMD Loaders to Write and Manage Modular JavaScript,” John Hann
- “Demystifying CommonJS Modules,” Alex Young
- “AMD Module Patterns: Singleton,” John Hann
- “Run-Anywhere JavaScript Modules Boilerplate Code,” Kris Zyp
- “Standards And Proposals for JavaScript Modules And jQuery,” James Burke

What Makes A Good Plugin Beyond Patterns?

At the end of the day, patterns are just one aspect of plugin development. And before we wrap up, here are my criteria for selecting third-party plugins, which will hopefully help developers write them.

Quality Do your best to adhere to best-practices with both the JavaScript and jQuery that you write. Are your solutions optimal? Do they follow the [jQuery Core Style Guidelines](#)? If not, is your code at least relatively clean and readable?

Compatibility Which versions of jQuery is your plugin compatible with? Have you tested it with the latest builds? If the plugin was written before jQuery 1.6, then it might have issues with attributes, because the way we approach them changed with that release. New versions of jQuery offer improvements and opportunities for the jQuery project to improve on what the core library offers. With this comes occasional breakages (mainly in major releases) as we move towards a better way of doing things. I'd like to see plugin authors update their code when necessary or, at a minimum, test their plugins with new versions to make sure everything works as expected.

Reliability Your plugin should come with its own set of unit tests. Not only do these prove your plugin actually works, but they can also improve the design without breaking it for end users. I consider unit tests essential for any serious jQuery plugin that is meant for a production environment, and they're not that hard to write. For an excellent guide to automated JavaScript testing with QUnit, you may be interested in "[Automating JavaScript Testing With QUnit](#)," by [Jorn Zaefferer](#).

Performance If the plugin needs to perform tasks that require a lot of computing power or that heavily manipulates the DOM, then you should follow best-practices that minimize this. Use [jsPerf.com](#) to test segments of your code so that you're aware of how well it performs in different browsers before releasing the plugin.

Documentation If you intend for other developers to use your plugin, ensure that it's well documented. Document your API. What methods and options does the plugin support? Does it have any gotchas that users need to be aware of? If users cannot figure out how to use your plugin, they'll likely look for an alternative. Also,

do your best to comment the code. This is by far the best gift you could give to other developers. If someone feels they can navigate your code base well enough to fork it or improve it, then you've done a good job.

Likelihood of maintenance When releasing a plugin, estimate how much time you'll have to devote to maintenance and support. We all love to share our plugins with the community, but you need to set expectations for your ability to answer questions, address issues and make improvements. This can be done simply by stating your intentions for maintenance in the README file, and let users decide whether to make fixes themselves.

In this section, we've explored several time-saving design patterns and best-practices that can be employed to improve your plugin development process. Some are better suited to certain use cases than others, but I hope that the code comments that discuss the ins and outs of these variations on popular plugins and widgets were useful.

Remember, when selecting a pattern, be practical. Don't use a plugin pattern just for the sake of it; rather, spend some time understanding the underlying structure, and establish how well it solves your problem or fits the component you're trying to build. Choose the pattern that best suits your needs.

Chapter 15. Conclusions

That's it for this introduction to the world of design patterns in JavaScript - I hope you've found it useful. The contents of this book should hopefully have given you sufficient information to get started using the patterns covered in your day-to-day projects.

Design patterns make it easier to reuse successful designs and architectures. It's important for every developer to be aware of design patterns but it's also essential to know how and when to use them. Implementing the right patterns intelligently can be worth the effort but the opposite is also true. A badly implemented pattern can yield little benefit to a project.

Also keep in mind that it is not the number of patterns you implement that's important but how you choose to implement them. For example, don't choose a pattern just for the sake of using 'one' but rather try understanding the pros and cons of what particular patterns have to offer and make a judgment based on it's fitness for your application.

If I've encouraged your interest in this area further and you would like to learn more about design patterns, there are a number of excellent titles on this area available for generic software development but also those that cover specific languages.

I'm happy to recommend:

1. 'Patterns Of Enterprise Application Architecture' by Martin Fowler
2. 'JavaScript Patterns' by Stoyan Stefanov
3. 'Pro JavaScript Design Patterns' by Ross Harmes and Dustin Diaz.

If you've managed to absorb most of the information in my book, I think you'll find reading these the next logical step in your learning process (beyond trying out some pattern examples for yourself of course).

Thanks for reading Essential JavaScript Design Patterns. For more educational material on learning JavaScript, please feel free to read more from me on my blog <http://addyosmani.com> or on Twitter @addyosmani.

Chapter 16. References

1. Design Principles and Design Patterns - Robert C Martin http://www.objectmentor.com/resources/articles/Principles_and_Patterns
2. Ralph Johnson - Special Issue of ACM On Patterns and Pattern Languages - <http://www.cs.wustl.edu/~schmidt/CACM-editorial.html>
3. Hillside Engineering Design Patterns Library - <http://hillside.net/patterns/>
4. Pro JavaScript Design Patterns - Ross Harmes and Dustin Diaz <http://jsdesignpatterns.com/>
5. Design Pattern Definitions - http://en.wikipedia.org/wiki/Design_Patterns
6. Patterns and Software Terminology <http://www.cmcrossroads.com/bradapp/docs/patterns-intro.html>
7. Reap the benefits of Design Patterns - Jeff Juday http://articles.techrepublic.com.com/5100-10878_11-5173591.html
8. JavaScript Design Patterns - Subramanyan Guhan <http://www.slideshare.net/rmsguhan/javascript-design-patterns>
9. What Are Design Patterns and Do I Need Them? - James Moaoriello <http://www.developer.com/design/article.php/1474561>
10. Software Design Patterns - Alex Barnett <http://alexbarnett.net/blog/archive/2007/07/20/software-design-patterns.aspx>
11. Evaluating Software Design Patterns - Gunni Rode <http://www.rode.dk/thesis/>
12. SourceMaking Design Patterns http://sourcemaking.com/design_patterns
13. The Singleton - Prototyp.ical <http://prototyp.ical.ly/index.php/2007/03/01/javascript-design-patterns-1-the-singleton/>
14. JavaScript Patterns - Stoyan Stevanov - <http://www.slideshare.net/stoyan/javascript-patterns>
15. Stack Overflow - Design Pattern Implementations in JavaScript (discussion)

<http://stackoverflow.com/questions/24642/what-are-some-examples-of-design-pattern-implementations-using-javascript>

16. The Elements of a Design Pattern - Jared Spool
http://www.uie.com/articles/elements_of_a_design_pattern/
17. Stack Overflow - Examples of Practical JS Design Patterns (discussion)
<http://stackoverflow.com/questions/3722820/examples-of-practical-javascript-object-oriented-design-patterns>
18. Design Patterns in JavaScript Part 1 - Nicholas Zakas
<http://www.webreference.com/programming/javascript/ncz/column5/>
19. Stack Overflow - Design Patterns in jQuery
<http://stackoverflow.com/questions/3631039/design-patterns-used-in-the-jquery-library>
20. Classifying Design Patterns By AntiClue - Elyse Neilson
<http://www.anticlue.net/archives/000198.htm>
21. Design Patterns, Pattern Languages and Frameworks - Douglas Schmidt
<http://www.cs.wustl.edu/~schmidt/patterns.html>
22. Show Love To The Module Pattern - Christian Heilmann <http://www.wait-till-i.com/2007/07/24/show-love-to-the-module-pattern/>
23. JavaScript Design Patterns - Mike G.
<http://www.lovemikeg.com/2010/09/29/javascript-design-patterns/>
24. Software Designs Made Simple - Anoop Mashudanan
<http://www.scribd.com/doc/16352479/Software-Design-Patterns-Made-Simple>
25. JavaScript Design Patterns - Klaus Komenda
<http://www.klauskomenda.com/code/javascript-programming-patterns/>
26. Introduction to the JavaScript Module Pattern <https://www.unleashed-technologies.com/blog/2010/12/09/introduction-javascript-module-design-pattern>
27. Design Patterns Explained - <http://c2.com/cgi/wiki?DesignPatterns>
28. Mixins explained <http://en.wikipedia.org/wiki/Mixin>

29. Working with GoF's Design Patterns In JavaScript
http://aspalliance.com/1782_Working_with_GoFs_Design_Patterns_in_JavaS
30. Using Object.create<http://stackoverflow.com/questions/2709612/using-object-create-instead-of-new>
31. t3knomanster's JavaScript Design Patterns -
<http://t3knomanser.livejournal.com/922171.html>
32. Working with GoF Design Patterns In JavaScript Programming -
http://aspalliance.com/1782_Working_with_GoFs_Design_Patterns_in_JavaS
33. JavaScript Advantages - Object Literals
<http://stackoverflow.com/questions/1600130/javascript-advantages-of-object-literal>
34. JavaScript Class Patterns - Liam McLennan
<http://geekswithblogs.net/liammclennan/archive/2011/02/06/143842.aspx>
35. Understanding proxies in jQuery -
<http://stackoverflow.com/questions/4986329/understanding-proxy-in-jquery>
36. Speaking on the Observer pattern -
<http://www.javaworld.com/javaworld/javaqa/2001-05/04-qa-0525-observer.html>
37. Singleton examples in JavaScript - Hardcode.nl -
http://www.hardcode.nl/subcategory_1/article_526-singleton-examples-in-javascript.htm
38. Design Patterns by Gamma, Helm supplement -
<http://exciton.cs.rice.edu/javaresources/DesignPatterns/>

About the Author

Addy Osmani is a writer, speaker, and a JavaScript developer. He is a member of the jQuery core [Bug Triage/Docs/Learning] teams where helps with bugs, documentation, and developer evangelism.

When not working at AOL his personal OS projects include jQuery UI Bootstrap and TodoMVC, which help developers compare JavaScript MVC frameworks. Addy is the author of the popular ebook, Essential JavaScript Design Patterns. His personal blog is <http://addyosmani.com/blog/>